

STA 2450 Computing for Math & Stat

Table of contents

Introduction	10
1 An Intro to Computing	11
1.1 Algorithms on Computers	12
1.2 Programming Languages	12
1.3 Types of Programming Languages	13
1.3.1 Low-Level vs. High-Level Languages	13
1.3.2 General-Purpose vs. Domain-Specific Languages	14
1.3.3 Interpreted vs. Compiled Languages	14
1.4 Intro to Python	16
2 Installing Python and Setting Up Your IDE	17
2.1 Installing Python	17
2.2 Understanding Integrated Development Environments (IDEs)	18
2.3 Setting Up Your IDE	19
2.4 Using Google Colab	22
2.4.1 Accessing Google Colab	22
2.4.2 Connecting Google Drive to Colab	22
2.4.3 Installing Packages in Google Colab	23
3 Basic Syntax and Expressions	24
3.1 Syntax in Programming Languages	24
3.2 Semantics in Programming Languages	26
4 Names and Variables	28
4.1 Names	28
4.1.1 Rules for Naming	28
4.2 Variables	28
4.2.1 Examples and Applications	30
4.3 Exercises	31
5 Data Types	32
5.1 Common Data Types in Python	32
5.1.1 Integer (int)	32
5.1.2 Float (float)	33
5.1.3 String (str)	33

5.1.4	Boolean (bool)	34
5.1.5	NoneType (None)	34
5.2	Dynamic Typing in Python	35
5.3	Type Conversion	35
5.4	Operators and Data Types	36
5.4.1	Arithmetic Operators	36
5.4.2	String Operators	37
5.4.3	Comparison Operators	37
5.4.4	Logical Operators	38
5.5	Exercises	39
6	Control Structures: Conditional Statements	42
6.1	The <code>if</code> Statement	42
6.2	The <code>else</code> Clause	46
6.3	The <code>elif</code> Clause	48
6.4	Conditional Expressions	51
6.5	Exercises	51
7	Control Structures: Loops	53
7.1	The <code>while</code> Loop	54
7.2	The <code>for</code> Loop	57
7.3	Exercises	62
8	Functions	64
8.1	Defining a Function	65
8.2	Calling a Function	66
8.3	Function Arguments	66
8.3.1	Positional Arguments	66
8.3.2	Keyword Arguments	68
8.3.3	Default Arguments	69
8.3.4	Variable-Length Arguments	69
8.4	Return Values	70
8.4.1	Returning Multiple Values	71
8.5	Best Practices in Function Design	73
8.6	Exercises	77
9	Introduction to Modules and Libraries	78
9.1	Understanding Modules	78
9.2	Importing Modules	79
9.3	Built-in Libraries	83
9.4	Creating Custom Modules	84
9.5	Best Practices for Modules and Libraries	85
9.6	Exercises	86

10 Data Structures: Lists and Tuples	87
10.1 Lists	87
10.1.1 Creating Lists	87
10.1.2 Accessing Elements in a List	88
10.1.3 Modifying Lists	88
10.1.4 Slicing Lists	90
10.1.5 Common List Operations	91
10.1.6 Iterating Over Lists	92
10.2 Tuples	93
10.2.1 Creating Tuples	93
10.2.2 Accessing Tuple Elements	94
10.2.3 Tuple Operations	95
10.2.4 Nested Tuples	96
10.2.5 Unpacking Tuples	96
10.2.6 Tuple Methods	97
10.3 List Comprehension	97
10.3.1 Basic Syntax of List Comprehension	98
10.3.2 List Comprehension with Conditional Logic	98
10.3.3 Nested List Comprehension	99
10.3.4 List Comprehension with Built-in Functions	100
10.3.5 List Comprehension vs. Loops	101
10.3.6 List Comprehension with External Libraries	101
10.3.7 Limitations of List Comprehension	102
10.3.8 Example: Practical Applications of List Comprehension	102
10.4 Exercises	103
11 Data Structures: Dictionaries and Sets	106
11.1 Dictionaries	106
11.1.1 Creating a Dictionary	106
11.1.2 Accessing Values	107
11.1.3 Adding and Modifying Entries	107
11.1.4 Deleting Key-Value Pairs	108
11.1.5 Dictionary Methods	109
11.1.6 Iterating Over Dictionaries	110
11.1.7 Dictionary Comprehension	111
11.1.8 Practical Applications of Dictionaries	111
11.1.9 Using **kwargs in Functions	113
11.2 Sets	116
11.2.1 Creating a Set	116
11.2.2 Adding and Removing Elements	117
11.2.3 Set Operations	118
11.2.4 Checking for Subsets and Supersets	120
11.2.5 Frozen Sets	121

11.2.6	Set Comprehensions	122
11.2.7	Set Best Practices	122
11.2.8	Combining Dictionaries and Sets	123
11.3	Exercises	123
12	JSON Files in Python	128
12.1	Introduction to JSON	128
12.1.1	Why Use JSON?	128
12.1.2	JSON Structure and Data Types	129
12.1.3	Differences Between JSON and Python Data Types	131
12.1.4	Use Cases for JSON	132
12.2	Reading and Writing JSON Data	132
12.2.1	Loading JSON Data from a File	133
12.2.2	Error Handling While Loading JSON Data	134
12.2.3	Writing JSON Data to a File	134
12.2.4	Error Handling When Writing JSON Data	135
12.2.5	Loading JSON Data from a String	136
12.2.6	Converting Python Objects to JSON Strings	136
12.2.7	Customizing JSON Serialization	137
12.3	Exercises	138
13	NumPy	142
13.1	Introduction to NumPy	142
13.1.1	Why Use NumPy?	142
13.1.2	Installing NumPy	143
13.2	Arrays in NumPy	143
13.2.1	Creating Arrays	144
13.2.2	Array Indexing and Slicing	146
13.3	Array Operations	148
13.3.1	Element-Wise Operations	148
13.3.2	Scalar Operations	148
13.3.3	Broadcasting	149
13.3.4	Array Shape and Reshaping	149
13.4	Mathematical Functions in NumPy	150
13.4.1	Common Mathematical Functions	150
13.4.2	Array Comparisons and Conditional Operations	151
13.5	NumPy Data Types	151
13.6	Exercises	152
14	Pandas	154
14.1	Introduction to Pandas	154
14.1.1	Why Use Pandas?	154
14.1.2	Key Features of Pandas	154

14.2	Series	155
14.2.1	Creating a Series	155
14.2.2	Handling Missing Data in a Series	156
14.2.3	Handling Duplicates in a Series	157
14.2.4	Performing Numeric Computations with a Series	159
14.3	DataFrame	160
14.3.1	Creating a DataFrame	160
14.3.2	Indexing DataFrames	161
14.3.3	Handling Missing Data in a DataFrame	163
14.3.4	Handling Duplicates in a DataFrame	164
14.3.5	Performing Numeric Computations with a DataFrame	165
14.4	Sorting, Filtering, and Merging DataFrames	167
14.4.1	Sorting DataFrames	167
14.4.2	Filtering DataFrames	168
14.4.3	Merging DataFrames	169
14.5	Reading and Writing Data	170
14.5.1	Reading Data from Files	171
14.5.2	Writing Data to Files	171
14.6	Exercises	171
15	Data Visualization in Python	175
15.1	Introduction to Matplotlib	175
15.1.1	Creating Basic Plots with Matplotlib	175
15.1.2	Types of Plots in Matplotlib	176
15.1.3	Customizing Plots in Matplotlib	177
15.1.4	Adding Annotations and Legends	178
15.1.5	Working with Subplots	180
15.1.6	Styling and Themes in Matplotlib	182
15.1.7	Saving Plots to Files	183
15.1.8	Common Pitfalls and Best Practices	184
15.2	Introduction to Seaborn	184
15.2.1	Getting Started with Seaborn	184
15.2.2	Creating Basic Plots with Seaborn	184
15.2.3	Customizing Visualizations with Seaborn	186
15.2.4	Visualizing Distributions	187
15.2.5	Creating Categorical Plots	188
15.2.6	Pair Plots for Multivariate Analysis	190
15.2.7	Heatmaps for Correlation Analysis	191
15.2.8	Customizing Seaborn Themes	192
15.2.9	Combining Seaborn with Matplotlib	194
15.2.10	Advantages of Using Seaborn	195
15.2.11	Limitations and Considerations	196

15.3	Introduction to ggplot (plotnine)	196
15.3.1	The Grammar of Graphics	196
15.3.2	Creating a Basic Plot with ggplot	196
15.3.3	Enhancing Visualizations with Layers	198
15.3.4	Common Plot Types in ggplot	199
15.3.5	Faceting for Multi-Plot Layouts	201
15.3.6	Customizing Aesthetics and Themes	202
15.3.7	Statistical Transformations	203
15.3.8	Advantages of Using ggplot	205
15.3.9	Limitations and Considerations	205
15.4	Exercises	205
16	Object-Oriented Programming in Python	208
16.1	Introduction to Object-Oriented Programming	208
16.1.1	Objects and Classes in Everyday Life	208
16.1.2	Attributes and Methods	209
16.1.3	Encapsulation: Keeping Attributes and Methods Together	209
16.1.4	Why Use OOP?	210
16.2	Defining a Class	211
16.2.1	Basic Class Syntax	211
16.2.2	The <code>__init__</code> Method: Class Constructor	212
16.3	Methods and Attributes in a Class	213
16.3.1	Instance Attributes vs. Class Attributes	214
16.3.2	Modifying Attributes	215
16.4	Inheritance	216
16.4.1	Inheritance in Everyday Life	216
16.4.2	Defining Inheritance in Python	216
16.4.3	Overriding Methods	218
16.4.4	The <code>super()</code> Function	218
16.4.5	Multiple Inheritance	219
16.5	Encapsulation and Data Hiding	220
16.5.1	Data Hiding: The Foundation of Encapsulation	220
16.5.2	Getter and Setter Methods	222
16.6	Polymorphism	223
16.6.1	Real-World Example: Payment Systems	224
16.6.2	Polymorphism Through Method Overriding	224
16.6.3	Polymorphism Through Method Overloading (Not Native to Python)	225
16.6.4	Polymorphism in Python with Duck Typing	226
16.6.5	Advantages of Polymorphism	227
16.7	Exercises	227

17 Writing Clear and Effective Documentation and PEP 8	229
17.1 Introduction	229
17.1.1 Types of Documentation	229
17.1.2 Characteristics of Good Documentation	229
17.1.3 Documentation vs. Self-Documenting Code	230
17.1.4 The Role of Documentation in Collaborative Projects	230
17.1.5 The Relationship Between Documentation and Code Quality	231
17.2 Best Practices for Writing Documentation	231
17.2.1 Write Meaningful Docstrings	231
17.2.2 Comment Strategically	232
17.2.3 Keep Documentation Up-to-Date	232
17.2.4 Provide Usage Examples	233
17.2.5 Use Consistent Formatting	233
17.2.6 Automate Documentation Generation	234
17.2.7 Include Error Handling Information	236
17.3 Understanding PEP 8	236
17.3.1 Why PEP 8 Matters	236
17.3.2 Key PEP 8 Guidelines	237
17.3.3 Common PEP 8 Pitfalls	239
17.3.4 Exceptions to PEP 8	240
18 Exception Handling and Debugging Techniques	241
18.1 Introduction to Exception Handling	241
18.1.1 Types of Errors in Python	241
18.1.2 Why Handle Exceptions?	242
18.1.3 The Life Cycle of an Exception	242
18.1.4 Common Exceptions in Python	243
18.1.5 The Philosophy Behind Python's Exception Handling	243
18.1.6 Exceptions in Statistical and Mathematical Computing	243
18.2 The <code>try</code> and <code>except</code> Blocks	245
18.2.1 Basic Structure of <code>try</code> and <code>except</code>	245
18.2.2 Catching Specific Exceptions	245
18.2.3 Catching Multiple Exceptions in One Block	246
18.2.4 Handling All Exceptions (Not Recommended)	246
18.2.5 Using the <code>else</code> Block	246
18.2.6 The <code>finally</code> Block	247
18.2.7 Nesting <code>try</code> and <code>except</code> Blocks	247
18.2.8 Raising Exceptions	248
18.3 Debugging Techniques	248
18.3.1 Using Print Statements for Debugging	249
18.3.2 Using Python's Built-In Debugger (<code>pdb</code>)	249
18.3.3 IDE Debugging Tools	250
18.3.4 Testing and Assertions	251

18.3.5	Logging for Debugging	253
18.3.6	Writing Clean Code to Minimize Bugs	254
18.4	Exercises	254
19	Regular Expressions	257
19.1	The Concept of Pattern Matching	257
19.1.1	Anatomy of a Regular Expression	257
19.1.2	Why Use Regular Expressions?	258
19.2	Using Regular Expressions in Python	258
19.2.1	Key Functions in Python's <code>re</code> Library	259
19.2.2	Using <code>re.search()</code>	259
19.2.3	Using <code>re.findall()</code>	260
19.2.4	<code>re.match()</code>	261
19.2.5	Using <code>re.sub()</code>	261
19.2.6	Using <code>re.split()</code>	262
19.3	Common Metacharacters	262
19.4	Building Complex Patterns	264
19.4.1	Using Grouping and Capturing	264
19.4.2	Lookahead and Lookbehind Assertions	265
19.4.3	Alternation with the <code> </code> Operator	266
19.4.4	Working with Character Classes and Ranges	266
19.4.5	Using Backreferences	267
19.4.6	Conditional Matching with Quantifiers	267
19.4.7	Combining Techniques for Complex Patterns	268
19.5	Exercises	268

Introduction

These are the lecture notes for STA 2450 - Computing for Math & Stat.

Prerequisites: None

Course Description:

Computer programming for mathematical scientists with emphasis on designing algorithms, problem solving, and coding practices. Topics include development of programs from specifications; appropriate use of data types; functions; modular program organization; documentation and style; and version control and collaborative programming.

1 An Intro to Computing

In the early 9th century AD, the Persian scholar Muhammad ibn Musa al-Khwarizmi authored a book that significantly influenced the development of European mathematics. Titled “Al-Kitāb al-mukhtaṣar fī ḥisāb al-jabr wa’l-muqābala” (“The Compendious Book on Calculation by Completion and Balancing”), Al-Khwarizmi introduced a collection of rules for solving linear and quadratic equations based on intuitive geometric arguments to the Western world. This work was translated into Latin in the 12th century, leading to the term algebra (the Latin title of the book is “Liber Algebræ et Almucabola”). Although algebra had earlier roots in Greece, it was Al-Khwarizmi’s book that made it accessible to European mathematicians, and it also laid the foundations for algorithms.

An algorithm is a finite, well-defined sequence of computational instructions or steps designed to perform a specific task or solve a particular problem. Each step in an algorithm is precise and unambiguous, and the algorithm is expected to produce a desired output within a finite amount of time. Algorithms are fundamental to computer science and mathematics, providing a systematic method for problem-solving and decision-making.

Let’s consider making a pot of coffee as an example of an algorithm. Here’s a simple, step-by-step process:

1. **Fill the Coffee Maker with Water:** Measure and pour the required amount of water into the coffee maker’s reservoir.
2. **Place the Coffee Filter:** Open the filter compartment and place a coffee filter inside.
3. **Add Coffee Grounds:** Measure the appropriate amount of coffee grounds and pour them into the filter.
4. **Close the Filter Compartment:** Securely close the filter compartment.
5. **Turn On the Coffee Maker:** Press the power button to start the brewing process.
6. **Wait for the Coffee to Brew:** Allow the coffee maker to complete its brewing cycle.
7. **Serve the Coffee:** Once the brewing is finished, pour the freshly brewed coffee into a cup and enjoy.

Each step in this process must be followed in the correct order to successfully make a pot of coffee, illustrating the fundamental nature of algorithms in organizing and executing tasks.

Algorithms are like recipes for computing. They provide instructions on how to solve problems, find answers, or move from point A to point B using only the resources available in your computer’s memory.

1.1 Algorithms on Computers

One way to illustrate the concept of a recipe in a mechanical process is by designing machines specifically intended to solve certain problems. The earliest computers were **fixed-program** computers, meaning they were designed to solve a specific mathematical problem. Some simple computers still use this approach. For example, a four-function calculator can perform basic arithmetic but is not designed for word processing or gaming. Users cannot change their programs without replacing internal components.

The first modern computer to run a program was the Manchester Mark 1 in 1949. Unlike its predecessors, it was a **stored-program** computer. These machines store sequences of instructions that can be executed by an **interpreter**. This interpreter can execute any legal set of instructions, enabling it to compute anything that can be described using those instructions. The program and the data it manipulates reside in memory.

Typically, a program **counter** points to a specific location in memory, and computation starts by executing the instruction at that point. Most often, the interpreter simply proceeds to the next instruction in the sequence. However, in some cases, it performs a test, and based on that test, execution may jump to another point in the sequence of instructions.

1.2 Programming Languages

To use execute algorithms on a computer, we need a programming language to describe the sequence of instructions. The British mathematician Alan Turing devised a theoretical device in 1936 known as the Universal Turing Machine. This hypothetical computer had unlimited memory, represented by a tape on which one could write zeroes and ones. It also consisted of simple instructions for moving, reading, and writing on the tape.

The Church-Turing thesis posits that if a function is computable, a Turing Machine can be programmed to compute it. The “if” in the Church-Turing thesis is crucial because not all problems have computational solutions.

Note

A function or problem is considered **computable** if there exists a finite sequence of well-defined steps, an algorithm, that can be implemented by a computational model (such as a Turing Machine) to produce the correct output for any valid input within a finite amount of time. In other words, a problem is computable if there is a systematic method to solve it using an algorithm.

The Church-Turing thesis leads directly to the concept of Turing completeness. A programming language is said to be **Turing complete** if it can simulate a universal Turing Machine. All modern programming languages are Turing complete, meaning that anything that can be

programmed in one language (e.g., Python) can also be programmed in another language (e.g., Java). While certain tasks may be easier to program in specific languages, all languages are fundamentally equal in terms of computational power.

1.3 Types of Programming Languages

1.3.1 Low-Level vs. High-Level Languages

Low-level languages are closer to machine language, providing little to no abstraction from a computer's hardware. They allow direct control over the hardware and memory, making them highly efficient but also more complex to write and understand. Examples of low-level languages include Assembly language and machine code.

Low-level languages operate with minimal abstraction which results in high-performance programs that execute faster and are more efficient in resource usage. However, this comes with the trade-off of complex and less intuitive code, requiring a deep understanding of computer architecture. Additionally, low-level programs are often platform-specific, meaning they are tailored to a particular type of processor or computer architecture.

Example: Assembly Language

```
MOV AX, 1    ; Move the value 1 into the AX register
ADD AX, 2    ; Add the value 2 to the AX register
MOV BX, AX   ; Move the result from AX to BX
```

High-level languages abstract away the details of computer hardware, allowing programmers to focus on the logic of the problem rather than the intricacies of the machine. They feature more intuitive syntax and semantics, making them easier to learn and use. Programs written in high-level languages are generally portable across different platforms with minimal modification. These languages also come with extensive libraries and frameworks that simplify complex tasks such as GUI development, web programming, and data manipulation. Although this abstraction can result in slower execution compared to low-level languages, the trade-off is often worth it for the ease of development and maintenance.

Example: Print a greeting to the screen in Python

```
print("Hello, World!")
```

Low-level languages offer greater control over hardware and performance optimization, whereas high-level languages prioritize ease of use and development speed. This means that while low-level languages generally result in faster and more efficient programs, they come at the cost

of greater complexity. In contrast, high-level languages, although potentially slower, provide significant productivity benefits. Their readability and maintainability make them suitable for large-scale applications and collaborative projects. Additionally, high-level languages reduce development time due to their simplicity and the availability of powerful libraries and frameworks.

Understanding the trade-offs between low-level and high-level languages is crucial for selecting the right tool for a given task. For most applications, high-level languages provide sufficient performance while dramatically simplifying development. However, for performance-critical applications, such as operating system kernels or embedded systems, low-level languages remain indispensable.

1.3.2 General-Purpose vs. Domain-Specific Languages

General-purpose programming languages are highly versatile, allowing them to be used for a wide array of programming tasks, including web development, data analysis, game development, and more. They come with extensive standard libraries and frameworks that support diverse functionalities, making it easier to implement complex features. Programs written in these languages are often portable across different operating systems and platforms. Additionally, these languages usually have large communities, extensive documentation, and abundant resources for learning and troubleshooting. Examples of well-known general-purpose programming languages include Python, Java, C++, and JavaScript.

Domain-specific Domain-specific languages (DSLs) are designed to handle specific types of tasks, offering features and abstractions directly relevant to their domain. They often allow for more concise and efficient coding within their domain, reducing the complexity and effort required to develop certain applications. However, their functionality is typically narrow, focusing on particular problem areas and lacking the versatility of general-purpose languages. Examples of well-known domain-specific languages include SQL, HTML, and MATLAB.

General-purpose languages are suitable for a wide variety of applications, providing greater flexibility and usability across many contexts. In contrast, domain-specific languages are specialized for particular tasks, making them more efficient and easier to use within their specific domains. While general-purpose languages often require learning a broader set of concepts and syntax, domain-specific languages might have a steeper but narrower learning curve focused on their particular use case.

1.3.3 Interpreted vs. Compiled Languages

Interpreted languages execute code directly through an interpreter, which translates high-level instructions into machine code line by line. This immediate execution provides several advantages, such as ease of debugging, as interpreted languages often generate more informative error messages, making it simpler to identify and correct issues by pointing to the exact

line where an error occurred. Additionally, interpreted languages offer platform independence, allowing the same source code to run on any platform with a compatible interpreter, enhancing portability. However, interpreted programs generally run slower than compiled programs because the translation occurs at runtime. Well-known interpreted languages include Python, JavaScript, Ruby, and PHP.

Compiled languages require the source code to be translated into machine code by a compiler before execution, producing an executable file. This pre-execution compilation results in faster execution, as the translation is completed beforehand, allowing the executable file to be directly executed by the hardware. Compilers also perform extensive checks during compilation, catching many types of errors before the program runs. However, compiled executables are specific to the target platform's architecture, making them less portable unless recompiled for different platforms. Well-known compiled languages include C, C++, Rust, and Go.

Example of C Code:

```
#include <stdio.h>

int main() {
    printf("Hello, World!\n");
    return 0;
}
```

This C code must be compiled into an executable file before it can be run.

Compiled languages generally produce faster-running programs since the code is translated into machine language ahead of time, whereas interpreted languages tend to be slower due to on-the-fly translation. However, interpreted languages offer a quicker development cycle, allowing code to be written and tested immediately without a separate compilation step, which is advantageous for rapid prototyping and iterative development. Interpreted languages also boast better portability, as the same code can run on any platform with the appropriate interpreter, while compiled languages require recompilation for different platforms, adding complexity.

Debugging is typically easier with interpreted languages because they provide more immediate and informative error feedback, pointing directly to issues in the source code. Although compiled languages catch many errors at compile time, their runtime error messages can be less detailed. Additionally, interpreted languages can support more dynamic programming features, such as dynamic typing and runtime evaluation of code, which can be more challenging to implement efficiently in compiled languages.

1.4 Intro to Python

In this course, we will be using Python. Python is a general-purpose programming language that can be used to build almost any kind of program that does not require direct access to the computer's hardware. However, Python is not optimal for programs with high reliability constraints or those that are built and maintained by many people over a long period.

Python has several advantages over many other languages. It is relatively simple and easy to learn. Because Python is designed to be interpreted, it provides runtime feedback that is especially helpful for novice programmers.

Python was first developed by Guido van Rossum in 1990. It went largely unnoticed during its first decade but gained popularity around 2000 with the release of Python 2.0. Python 3.0, released at the end of 2008, cleaned up many inconsistencies in Python 2's design. While Python 3.0 is not backward-compatible with earlier versions, most important public domain libraries have been ported over, eliminating the need to use outdated software.

2 Installing Python and Setting Up Your IDE

2.1 Installing Python

To start programming in Python, you'll first need to install Python on your computer. Follow the instructions below based on your operating system:

Windows:

1. Go to the official Python website: python.org.
2. Click on the “Downloads” tab.
3. Download the latest version of Python (ensure it's a stable release).
4. Run the installer and make sure to check the box that says “Add Python to PATH”.
5. Follow the installation prompts.

macOS:

1. Open Terminal.
2. Use Homebrew to install Python (if you don't have Homebrew, you can install it from brew.sh):

```
/bin/bash -c "$(curl -fsSL https://raw.githubusercontent.com/Homebrew/install/HEAD/in
```

3. Install Python:

```
brew install python
```

Linux:

1. Open Terminal.
2. Use your package manager to install Python. For Debian-based systems like Ubuntu:

```
sudo apt update
sudo apt install python3
```

For Red Hat-based systems like Fedora:

```
sudo dnf install python3
```

To verify the installation, open your terminal or command prompt and type:

```
python --version
```

You should see the installed Python version displayed.

2.2 Understanding Integrated Development Environments (IDEs)

2.2.0.1 What is an IDE?

An Integrated Development Environment (IDE) is a software application that provides comprehensive facilities to programmers for software development. An IDE typically includes:

1. **Source Code Editor:** A text editor that supports code writing with features like syntax highlighting, auto-completion, and error detection.
2. **Build Automation Tools:** Tools that automate tasks like compiling code, packaging files, and running tests.
3. **Debugger:** A tool that helps you test and debug your code by allowing you to step through your code, set breakpoints, and inspect variables.

IDEs are designed to make the process of coding easier and more efficient by combining all the tools you need into a single application.

2.2.0.2 Benefits of Using an IDE

Using an IDE offers numerous benefits and features that enhance the coding process. An IDE can significantly speed up development by providing quick access to all necessary tools, enhancing efficiency. Features like syntax highlighting and error detection help catch mistakes early, reducing the amount of debugging needed, thereby improving error reduction.

The convenience of having everything in one place makes it easier to manage projects, especially as they grow in size and complexity. Additionally, many IDEs include tools for version control, making it easier to track changes and collaborate with others, thus improving code management.

Key features of an IDE streamline the coding process further. Syntax highlighting colors keywords, variables, and syntax, making the code more readable and easier to debug. Auto-completion suggests code completions as you type, helping to write code faster and avoid typos. Code navigation allows you to quickly jump to definitions, references, and files within your

project. The integrated debugger lets you run your code step by step, inspect variables, and identify issues.

2.2.0.3 Popular Python IDEs

- **PyCharm:** A powerful IDE specifically for Python development, offering advanced features and professional tools.
- **RStudio:** Primarily used for R programming but also supports Python, providing a familiar interface for those who use both languages.
- **Visual Studio Code:** A versatile code editor with extensive Python support through extensions, combining lightweight design with powerful features.

2.2.0.4 Choosing an IDE

Choosing the right IDE depends on your needs and preferences. Some factors to consider include ease of use, which refers to how intuitive and user-friendly the IDE is, and the feature set, which involves the availability of features that match your development needs. Performance is also crucial, particularly how well the IDE performs with larger projects. Additionally, consider the community and support available, including documentation, tutorials, and community support.

Whether you use PyCharm, RStudio, or another IDE, having a good development environment can make a significant difference in your productivity and the quality of your code. A well-chosen IDE enhances your coding experience, making it easier to write, manage, and debug your projects efficiently.

2.3 Setting Up Your IDE

Using PyCharm:

PyCharm, developed by JetBrains, is a popular IDE specifically designed for Python programming. JetBrains, a software development company based in Prague, Czech Republic, was founded in 2000 and has since become renowned for creating intelligent, productivity-enhancing tools for software developers. The journey of PyCharm began in 2010, when JetBrains identified the need for a dedicated, feature-rich IDE for Python developers. At that time, Python was gaining significant traction due to its simplicity and versatility, and there was a growing demand for an IDE that could cater specifically to Python's unique requirements.

JetBrains leveraged its extensive experience in creating IDEs, such as IntelliJ IDEA for Java, to develop PyCharm. They aimed to build an IDE that would not only support the standard

features expected by developers, such as code completion and debugging, but also include advanced capabilities like scientific computing support, web development frameworks, and robust version control integration. PyCharm quickly became popular within the Python community for its powerful features, intuitive user interface, and the comprehensive support it provided for various Python frameworks and libraries. Over the years, PyCharm has continued to evolve, incorporating feedback from its user base and keeping pace with the latest advancements in Python development.

Here's how to set up PyCharm:

1. Download and Install PyCharm:

- Go to the JetBrains website: jetbrains.com/pycharm/download.
- Choose the Community edition (it's free).
- Follow the installation instructions for your operating system.

2. Configuring PyCharm:

- Open PyCharm.
- Select “New Project”.
- Choose a location for your project.
- Ensure the Python interpreter is set correctly. If Python is not listed, click on “Existing interpreter” and locate your Python installation.
- Click “Create”.

3. Writing Your First Program:

- Right-click on your project folder in the Project pane.
- Select “New” > “Python File”.
- Name your file (e.g., `hello_world.py`).
- Type the following code:

```
print("Hello, world!")
```

- Right-click on the file and select “Run” to execute your code.

Using RStudio:

RStudio, developed by RStudio, PBC, is a comprehensive IDE primarily designed for the R programming language. The inception of RStudio dates back to 2008, when its founders, including JJ Allaire, sought to create a powerful tool to support the growing community of R users. R, a language and environment for statistical computing and graphics, was gaining popularity among statisticians and data scientists due to its robust capabilities for data analysis and visualization. However, the available development tools for R at that time were limited, and there was a clear need for a more sophisticated and user-friendly IDE.

JJ Allaire, already an experienced software developer known for creating the ColdFusion web development platform, aimed to fill this gap by developing RStudio. The initial release of RStudio in 2011 was met with enthusiasm from the R community. It offered a user-friendly interface, integrating essential features like a source code editor, console, workspace, and graphics viewer within a single window. This made it significantly easier for users to write, debug, and visualize their R code.

Over the years, RStudio has expanded its functionality to support not only R but also other programming languages like Python, enhancing its utility for data scientists and statisticians. The open-source nature of RStudio and its commitment to the R community have contributed to its widespread adoption. Today, RStudio is a central tool in data science, used by individuals and organizations worldwide for data analysis, research, and reporting.

Here's how to set up RStudio:

1. Download and Install RStudio:

- Go to the RStudio website: rstudio.com.
- Download the free version for your operating system.
- Follow the installation instructions.

2. Configuring RStudio for Python:

- Open RStudio.
- Go to “Tools” > “Global Options”.
- Select “Python” from the menu on the left.
- Set the Python interpreter to your installed Python version (e.g., `/usr/bin/python3` or `C:\Python39\python.exe`).

3. Writing Your First Program:

- Go to “File” > “New File” > “Python Script”.
- Type the following code:

```
print("Hello, world!")
```

- Save the file with a `.py` extension (e.g., `hello_world.py`).
- Click the “Run” button to execute your code.

By following these instructions, you'll have Python installed and be ready to start coding with PyCharm or RStudio as your IDE.

2.4 Using Google Colab

Google Colab is a powerful tool that allows you to write and run Python code in a browser. It provides free access to computing resources, including GPU and TPU options, making it ideal for learning, experimenting, or running intensive computations without needing a high-performance computer. Colab is built on Jupyter Notebooks, which means you can easily share your code and outputs.

In this section, The following will be covered:

1. Accessing and using Google Colab
2. Connecting to your Google Drive
3. Installing additional Python packages in Colab

2.4.1 Accessing Google Colab

To access Google Colab, simply navigate to <https://colab.research.google.com> in your browser. You can start a new notebook by selecting **New Notebook** from the main dashboard.

Alternatively, if you have a Jupyter Notebook (.ipynb) saved on your local machine or in Google Drive, you can upload and work on it directly.

2.4.2 Connecting Google Drive to Colab

One of the most useful features of Google Colab is its ability to connect directly to your Google Drive, allowing you to load and save files with ease.

Here's how to connect your Google Drive to your Colab notebook:

1. **Open a new or existing Colab notebook.**
2. **Mount Google Drive** using the following code:

```
from google.colab import drive
drive.mount('/content/drive')
```

3. **Run the cell** by clicking the play button. A link will appear that directs you to authorize access to your Google Drive. Click the link and follow the instructions to grant access.
4. (If Necessary) After authorizing, **copy the authentication code** provided and paste it into the Colab input box.
5. Your Google Drive is now mounted, and you can access your files under the path `/content/drive/My Drive/`.

For example, to read a file from your Google Drive, you can use:

```
file_path = '/content/drive/My Drive/your_folder/your_file.csv'
```

To save a file back to Google Drive, specify the path similarly.

2.4.3 Installing Packages in Google Colab

Google Colab comes pre-installed with many common Python libraries (like NumPy, pandas, and Matplotlib), but you may occasionally need to install additional packages.

To install a new package, use the `!pip` or `!apt-get` commands directly within a notebook cell. For example, to install a package with pip:

```
!pip install some-package
```

For example, to install `plotly`, a Python library for creating interactive plots, you would use:

```
!pip install plotly
```

After running the command, the package will be available for immediate use in the current Colab session.

3 Basic Syntax and Expressions

In the realm of programming, understanding the basic syntax and semantics of a language is akin to grasping the fundamentals of a natural language. Syntax refers to the structure or format of the code, dictating how symbols and keywords can be combined to create valid instructions. Semantics, on the other hand, pertains to the meaning behind these instructions — what actions they perform when executed.

3.1 Syntax in Programming Languages

Syntax rules in programming languages define how the code should be written and structured. These rules ensure that the code is readable both by humans and by the computer. Just as syntax in a spoken language dictates how words and phrases should be combined to convey meaning, syntax in programming languages dictates how symbols and keywords can be combined to create valid instructions.

Key elements of syntax include keywords, operators, punctuation, and the overall structure of statements.

3.1.0.1 Keywords

Keywords are reserved words that have special meanings in a programming language. For example, in English, words like “and,” “or,” “if,” and “then” are used to connect ideas and indicate logical relationships. Similarly, in Python, keywords like `if`, `else`, `while`, `for`, and `def` are used to construct control flow and define functions.

Example in English: - “If it rains, then bring an umbrella.”

Example in Python:

```
if it_rains:
    bring_umbrella()
```


3.1.0.2 Operators

Operators are symbols that perform operations on values. In both English and programming languages, operators are used to combine or modify elements.

Example in English: - “Three plus five equals eight.”

Example in Python:

```
3 + 5
```

3.1.0.3 Punctuation

Punctuation marks, such as periods, commas, and parentheses, are used to group and organize sentences in English. Similarly, in programming languages, punctuation marks like parentheses (), brackets [], and braces {} are used to group and organize code.

Example in English: - “He said, ‘Hello!’ ”

Example in Python:

```
print("Hello!")
```

3.1.0.4 Statements and Expressions

A statement is an instruction that the Python interpreter can execute, while an expression is a combination of values, variables, and operators that can be evaluated to produce another value. Every expression is a statement, but not every statement is an expression.

Example in English: - “The sum of three and five is eight.” (Expression: “Three plus five”)

Example in Python:

```
3 + 5 # Expression  
print(3 + 5) # Statement
```

Note

Comments in Python code are denoted by the # symbol and are used to explain and document the code. They are ignored by the Python interpreter and do not affect the execution of the program, making them useful for enhancing code readability and providing context for developers.

Understanding the syntax rules in both natural and programming languages is crucial for effective communication and problem-solving. Syntax ensures that our messages, whether written in English or Python, are clear and interpretable by others, including computers.

3.2 Semantics in Programming Languages

Semantics concerns the meaning of syntactically correct strings of symbols. In other words, while syntax is about the form, semantics is about the function. Semantics ensure that a piece of code not only adheres to the rules of the language but also performs the intended operations.

For instance, the expression `3 + 4` adheres to Python's syntax rules and its semantics dictate that this expression will evaluate to 7.

To illustrate basic syntax and semantics, let's look at some simple Python expressions.

3.2.0.1 Example 1: Arithmetic Operations

Arithmetic operations in Python are straightforward and follow standard mathematical conventions.

```
# Adding two numbers
3 + 5

# Subtracting two numbers
10 - 4

# Multiplying two numbers
7 * 6

# Dividing two numbers
8 / 2
```

Each of these expressions follows the syntax rules for arithmetic operators in Python and has clear semantics: they perform addition, subtraction, multiplication, and division, respectively.

3.2.0.2 Example 2: Comparison Operations

Comparison operations are used to compare values and produce Boolean results (`True` or `False`).

```
# Checking equality
4 == 4

# Checking inequality
5 != 3

# Checking greater than
9 > 7

# Checking less than
2 < 6
```

These expressions compare two values and return a Boolean value based on the comparison.

3.2.0.3 Example 3: Logical Operations

Logical operations combine Boolean values and are useful in conditional expressions.

```
# Logical AND
True and False

# Logical OR
True or False

# Logical NOT
not True
```

In these examples, logical operators evaluate the Boolean expressions and return a Boolean result.

Understanding the basic syntax and semantics of a programming language is crucial for writing correct and efficient code. Syntax provides the structure, while semantics ensures that the code performs the intended operations.

4 Names and Variables

In programming, names and variables are fundamental concepts that form the building blocks of a program. They enable us to store, manipulate, and access data within our code, making it possible to write flexible and efficient programs. This chapter will delve into these essential concepts, exploring their significance and how they are used in Python.

4.1 Names

Names in programming are identifiers used to label and reference entities such as variables, functions, classes, and modules. They serve as a means to access stored data and functionality, making code more readable and maintainable. For instance, instead of repeatedly using a complex expression or value, we can assign it a name and refer to it using that name.

4.1.1 Rules for Naming

In Python, there are specific rules and conventions for naming identifiers:

1. **Alphabetic Characters and Underscores:** Names must begin with a letter (a-z, A-Z) or an underscore (`_`), followed by letters, digits (0-9), or underscores.
2. **Case Sensitivity:** Names are case-sensitive, meaning `Variable`, `variable`, and `VARIABLE` are distinct identifiers.
3. **Reserved Keywords:** Names cannot be Python reserved keywords, such as `if`, `else`, `for`, `while`, `class`, `def`, etc.

4.2 Variables

Variables are names assigned to data values stored in memory. They act as containers holding information that can be manipulated and accessed throughout a program. Variables allow for dynamic and flexible data handling, making it possible to perform computations, store results, and manage state.

4.2.0.1 Variable Assignment

In Python, variable assignment is straightforward. Use the assignment operator (=) to assign a value to a variable:

```
x = 10
name = "Alice"
is_student = True
```

4.2.0.2 Variable Scope

The scope of a variable determines its accessibility within different parts of a program. There are two primary types of scope in Python:

1. **Global Scope:** Variables defined outside any function (functions will be discussed in Chapter 8) or block have global scope and can be accessed anywhere in the program.
2. **Local Scope:** Variables defined within a function or block have local scope and can only be accessed within that function or block.

```
global_var = "I am global"

def my_function():
    local_var = "I am local"
    print(global_var) # Accessible
    print(local_var) # Accessible

my_function()
print(global_var) # Accessible
print(local_var) # Error: NameError
```

4.2.0.3 Constants

Constants are variables whose values are intended to remain unchanged throughout the program. While Python does not have built-in constant types, by convention, we use all uppercase letters to indicate constants:

```
PI = 3.14159
MAX_STUDENTS = 50
```

4.2.1 Examples and Applications

To solidify our understanding of names and variables, let's explore some practical examples and applications:

Example: Simple Calculator

```
# Simple Calculator Program
num1 = 10
num2 = 5

# Performing arithmetic operations
sum_result = num1 + num2
difference = num1 - num2
product = num1 * num2
quotient = num1 / num2

print("Sum:", sum_result)
print("Difference:", difference)
print("Product:", product)
print("Quotient:", quotient)
```

Example: Temperature Conversion

```
# Temperature Conversion Program (Celsius to Fahrenheit)
celsius = 25
fahrenheit = (celsius * 9/5) + 32

print(celsius, "Celsius is", fahrenheit, "Fahrenheit")
```

Understanding names and variables is crucial for effective programming. Names provide a way to reference and access data, while variables store the data we work with. By adhering to naming conventions and understanding variable scope and types, we can write clearer, more maintainable code. As we progress, these fundamental concepts will underpin more complex programming tasks and techniques.

4.3 Exercises

4.3.0.1 Exercise 1: Variable Assignment

- a. Assign the value 50 to a variable named `my_age` and use `print` to display the variable.
- b. Assign the string "Hello, Python!" to a variable named `greeting` and use `print` to display the variable.
- c. Assign the boolean value `False` to a variable named `is_raining` and use `print` to display the variable.

4.3.0.2 Exercise 2: Variable Reassignment

- a. Assign the value 10 to a variable named `number` and use `print` to display the variable.
- b. Reassign the value of `number` to 20 and use `print` to display the variable.
- c. Reassign the value of `number` to the string "twenty" and use `print` to display the variable.

4.3.0.3 Exercise 3: Basic Arithmetic

- a. Create two variables `a` and `b`, and assign them the values 8 and 3 respectively.
- b. Perform the following operations and print the results:
 - Sum of `a` and `b`
 - Difference between `a` and `b`
 - Product of `a` and `b`
 - Quotient of `a` divided by `b`

4.3.0.4 Exercise 4: Temperature Conversion

- a. Create a variable `fahrenheit` and assign it the value 86.
- b. Convert the temperature to Fahrenheit using the formula: `celsius = (fahrenheit - 32) * 5/9`.
- c. Print the Celsius value.

4.3.0.5 Exercise 5: Constants

- a. Define a constant `PI` with the value 3.14159.
- b. Define a constant `MAX_SPEED` with the value 120.
- c. Print both constants.

5 Data Types

Understanding data types is fundamental to programming in Python, as they dictate how the interpreter will manipulate and store values. Data types are the classification of data items. They tell the interpreter how the programmer intends to use the data. In Python, every value has a data type, and it is crucial to know these types to write efficient and error-free code.

5.0.0.1 The Importance of Data Types

Data types define the operations that can be performed on a particular piece of data, the structure in which the data can be stored, and the way the data can be manipulated. For example, you cannot perform arithmetic operations on a string of text, but you can concatenate it with another string. Understanding data types helps in preventing type-related errors and optimizing code performance.

5.1 Common Data Types in Python

Python supports various built-in data types, each suited for different kinds of operations and uses. The primary data types include:

1. **Integer (int)**
2. **Float (float)**
3. **String (str)**
4. **Boolean (bool)**
5. **NoneType (None)**

We will discuss each of these types in detail.

5.1.1 Integer (int)

An integer is a whole number, positive or negative, without decimals, of unlimited length. In Python, integers are of type `int`.

Example:


```
a = 10
b = -5
c = 123456789
```

To check the type of a variable, use the `type()` function:

```
print(type(a))
print(type(b))
```

```
<class 'int'>
<class 'int'>
```

5.1.2 Float (float)

A float, or floating-point number, is a number that has a decimal place. Python uses the `float` type to represent these numbers.

Example:

```
x = 10.5
y = -3.14
z = 1.0
```

Checking the type of a float variable:

```
print(type(x))
print(type(y))
```

```
<class 'float'>
<class 'float'>
```

5.1.3 String (str)

A string is a sequence of characters enclosed in quotes. Strings can be enclosed in single quotes (`'`), double quotes (`"`), or triple quotes (`'''` or `"""`).

Example:

```
name = "Alice"
greeting = 'Hello, World!'
paragraph = """This is a
```

```
multiline string."""
```

To determine if a variable is a string:

```
print(type(name))
```

```
<class 'str'>
```

5.1.4 Boolean (bool)

Booleans represent one of two values: **True** or **False**. They are commonly used in conditional statements and comparisons.

Example:

```
is_student = True  
is_raining = False
```

Checking the type of a boolean variable:

```
print(type(is_student))
```

```
<class 'bool'>
```

5.1.5 NoneType (None)

None is a special data type in Python that represents the absence of a value. It is an object of its own datatype, the **NoneType**.

Example:

```
unknown = None
```

To check if a variable is of **NoneType**:

```
print(type(unknown))
```

```
<class 'NoneType'>
```

5.2 Dynamic Typing in Python

Python is a dynamically typed language, which means you do not need to declare the type of a variable when you create one. The type is determined at runtime based on the value assigned to it.

Example:

```
variable = 5
print(type(variable))

variable = "Hello"
print(type(variable))
```

```
<class 'int'>
<class 'str'>
```

This feature provides flexibility but requires careful handling to avoid type-related errors.

5.3 Type Conversion

Sometimes, you may need to convert values from one type to another. This is known as **type casting**. Python provides several built-in functions for this purpose.

- `int()`: Converts a value to an integer.
- `float()`: Converts a value to a float.
- `str()`: Converts a value to a string.
- `bool()`: Converts a value to a boolean.

Example:

```
a = 10.5
b = int(a)  # b is now 10

c = "123"
d = int(c)  # d is now 123

e = "True"
f = bool(e)  # f is now True
```

Checking the types after conversion:

```
print(type(b))
print(type(d))
print(type(f))
```

```
<class 'int'>
<class 'int'>
<class 'bool'>
```

5.4 Operators and Data Types

Operators are used to perform operations on variables and values. Python supports various operators, and their behavior can differ based on the data types they operate on. Here, we will explore how operators work with different data types.

5.4.1 Arithmetic Operators

Arithmetic operators are used to perform common mathematical operations such as addition, subtraction, multiplication, and division.

Example with Integers:

```
a = 10
b = 3

print(a + b)
print(a - b)
print(a * b)
print(a / b)
print(a % b)
print(a ** b)
print(a // b)
```

```
13
7
30
3.3333333333333335
1
1000
3
```

Example with Floats:

```
x = 10.5
y = 3.2

print(x + y)
print(x - y)
print(x * y)
print(x / y)
```

```
13.7
7.3
33.6
3.28125
```

5.4.2 String Operators

Strings can be manipulated using the + (concatenation) and * (repetition) operators.

Example:

```
str1 = "Hello"
str2 = "World"

print(str1 + " " + str2)
print(str1 * 3)
```

```
Hello World
HelloHelloHello
```

5.4.3 Comparison Operators

Comparison operators are used to compare two values. They return a boolean value (**True** or **False**).

Example with Integers and Floats:

```
a = 10
b = 5
c = 10.0
```

```
print(a == b)
print(a != b)
print(a > b)
print(a < b)
print(a >= c)
print(a <= c)
```

False
True
True
False
True
True

Example with Strings:

```
str1 = "Hello"
str2 = "World"
str3 = "Hello"

print(str1 == str2)
print(str1 == str3)
print(str1 != str2)
print(str1 > str2)    # (Lexicographical comparison)
print(str1 < str2)
```

False
True
True
False
True

5.4.4 Logical Operators

Logical operators are used to combine conditional statements.

Example:

```
a = True
b = False
```

```
print(a and b)
print(a or b)
print(not a)
```

False

True

False

Logical operators can also be used with integers, where 0 is considered `False` and any non-zero value is considered `True`.

```
x = 0
y = 10

print(x and y)
print(x or y)
print(not x)
```

0

10

True

5.5 Exercises

5.5.0.1 Exercise 1: Identifying Data Types

For each of the following variables, use the `type()` function to determine its data type.

```
a = 42
b = 3.14
c = "Python"
d = True
e = None
```

5.5.0.2 Exercise 2: Type Conversion

Convert the following values to the specified types and print the results.

- a. Convert `x = 3.75` to an integer.
- b. Convert `y = "123"` to a float.
- c. Convert `z = 0` to a boolean.
- d. Convert `w = False` to a string.

5.5.0.3 Exercise 3: Arithmetic Operations

Perform the following arithmetic operations and print the results.

- a. Add 15 and 23.
- b. Subtract 9 from 45.
- c. Multiply 7 by 8.
- d. Divide 20 by 4.
- e. Calculate the modulus of 27 and 4.
- f. Raise 2 to the power of 5.
- g. Perform floor division of 17 by 3.

5.5.0.4 Exercise 4: String Operations

Use the appropriate operations to do the following and print the results.

- a. Concatenate the strings "Data" and "Science" with a space in between.
- b. Repeat the string "Hello" 5 times.

5.5.0.5 Exercise 5: Logical Operations

Evaluate the following logical expressions and print the results.

- a. `True and False`
- b. `True or False`
- c. `not True`
- d. `(5 > 3) and (2 < 4)`
- e. `(10 == 10) or (5 != 5)`

5.5.0.6 Exercise 6: Mixed-Type Operations

Do the following operations and print the results.

- a. Add an integer and a float: `7 + 3.14`
- b. Concatenate a string and an integer (convert the integer to a string first): `"Age: " + str(30)`

c. Multiply a string by an integer: `"Data" * 4`

6 Control Structures: Conditional Statements

In programming, the ability to make decisions is crucial. Conditional statements allow a program to take different actions based on whether certain conditions are true or false. This chapter will introduce the concept of conditional statements in Python, illustrating their significance and how they can be utilized effectively in various programming scenarios.

6.1 The `if` Statement

The `if` statement is the most fundamental building block of conditional statements in Python. It allows the program to execute a block of code only if a specified condition is true. This section will delve deeper into the mechanics of the `if` statement, its syntax, and its practical applications.

6.1.0.1 Basic Syntax

The basic syntax of an `if` statement in Python is straightforward. It consists of the keyword `if` followed by a condition, a colon, and an indented block of code that will be executed if the condition is true.

Syntax:

```
if condition:
    statement(s)
```

- **condition:** This is an expression that evaluates to either `True` or `False`.
- **statement(s):** This is the block of code that will be executed if the condition evaluates to `True`.

The condition can be any expression that returns a Boolean value (i.e., `True` or `False`). If the condition evaluates to `True`, the indented block of code following the `if` statement is executed. If the condition evaluates to `False`, the block of code is skipped.

Example:

```
x = 10
if x > 5:
    print("x is greater than 5")
```

x is greater than 5

In this example, the condition `x > 5` evaluates to `True` because 10 is greater than 5. Therefore, the code within the `if` block is executed, resulting in the output “x is greater than 5”.

6.1.0.2 Using Comparison Operators

The condition in an `if` statement often involves comparison operators. These operators compare two values and return `True` or `False` based on the comparison.

Common Comparison Operators:

- `==`: Equal to
- `!=`: Not equal to
- `>`: Greater than
- `<`: Less than
- `>=`: Greater than or equal to
- `<=`: Less than or equal to

Examples:

```
# Equal to
x = 5
if x == 5:
    print("x is equal to 5")
```

x is equal to 5

```
# Not equal to
y = 10
if y != 5:
    print("y is not equal to 5")
```

y is not equal to 5

```
# Greater than
a = 7
if a > 3:
    print("a is greater than 3")
```

a is greater than 3

```
# Less than
b = 2
if b < 5:
    print("b is less than 5")
```

b is less than 5

```
# Greater than or equal to
c = 8
if c >= 8:
    print("c is greater than or equal to 8")
```

c is greater than or equal to 8

```
# Less than or equal to
d = 4
if d <= 4:
    print("d is less than or equal to 4")
```

d is less than or equal to 4

6.1.0.3 Combining Conditions with Logical Operators

Sometimes, you need to check multiple conditions simultaneously. Python provides logical operators to combine multiple conditions.

Logical Operators:

- **and**: Returns **True** if both conditions are **True**
- **or**: Returns **True** if at least one condition is **True**
- **not**: Returns **True** if the condition is **False**

Examples:

```
# Using 'and' operator
x = 10
y = 20
if x > 5 and y > 15:
    print("Both conditions are true")
```

Both conditions are true

```
# Using 'or' operator
a = 5
b = 10
if a > 7 or b > 7:
    print("At least one condition is true")
```

At least one condition is true

```
# Using 'not' operator
c = 3
if not c > 5:
    print("c is not greater than 5")
```

c is not greater than 5

In the first example, the condition `x > 5 and y > 15` evaluates to `True` because both `10 > 5` and `20 > 15` are true. Therefore, the code within the `if` block is executed, resulting in the output “Both conditions are true”.

In the second example, the condition `a > 7 or b > 7` evaluates to `True` because `10 > 7` is true even though `5 > 7` is false. Hence, the output is “At least one condition is true”.

In the third example, the condition `not c > 5` evaluates to `True` because `c > 5` is false, and `not` operator negates it. Therefore, the output is “c is not greater than 5”.

6.1.0.4 Nested if Statements

You can nest `if` statements within other `if` statements to create more complex decision structures. This means placing one `if` statement inside another `if` statement’s block of code.

Example:

```
x = 15
if x > 10:
    print("x is greater than 10")
    if x > 20:
        print("x is also greater than 20")
```

x is greater than 10

In this example, the outer `if` statement checks if `x` is greater than 10. Since `x` is 15, the condition is true, and the code within the block is executed. Inside this block, there is another `if` statement that checks if `x` is greater than 20. Since 15 is not greater than 20, the block of the nested `if` statement is not executed.

6.2 The `else` Clause

The `else` clause in Python provides an alternative block of code that will execute if the condition in the `if` statement evaluates to `False`. This allows for a two-way decision-making process: if the condition is true, one set of statements will run, otherwise, a different set of statements will run.

6.2.0.1 Basic Syntax

The basic syntax for using the `else` clause follows directly after an `if` statement. The `else` clause must be at the same indentation level as the `if` statement, and its block of code must be indented further.

Syntax:

```
if condition:
    statement(s)
else:
    statement(s)
```

- **condition:** This is an expression that evaluates to either `True` or `False`.
- **statement(s):** This is the block of code that will be executed if the condition evaluates to `False`.

Example:

```
x = 3
if x > 5:
    print("x is greater than 5")
else:
    print("x is not greater than 5")
```

x is not greater than 5

In this example, the condition `x > 5` evaluates to **False** because 3 is not greater than 5. Therefore, the code within the **else** block is executed, resulting in the output “x is not greater than 5”.

6.2.0.2 Nested else Clauses

else clauses can be nested within other conditional blocks to create more complex decision structures. This is useful when the logic requires multiple levels of checks.

Example:

```
x = 15
if x > 10:
    print("x is greater than 10")
    if x > 20:
        print("x is also greater than 20")
    else:
        print("x is between 11 and 20")
else:
    print("x is 10 or less")
```

x is greater than 10
x is between 11 and 20

In this example, the outer **if** statement checks if `x` is greater than 10. Since `x` is 15, the condition is true, and the code within the block is executed. Inside this block, there is another **if** statement that checks if `x` is greater than 20. Since 15 is not greater than 20, the **else** block of the nested **if** statement is executed, resulting in the output “x is between 11 and 20”.

6.3 The `elif` Clause

In Python, the `elif` clause (short for “else if”) is used to check multiple conditions in a sequence. It allows you to add more than one conditional expression to an `if` statement, creating a chain of conditions that are evaluated in order. When one of these conditions evaluates to `True`, the corresponding block of code is executed, and the rest of the conditions are skipped.

6.3.0.1 Basic Syntax

The `elif` clause follows the `if` clause and is used to test additional conditions if the previous conditions were not true. You can have as many `elif` clauses as you need, and an optional `else` clause can be included at the end to handle cases where none of the `if` or `elif` conditions are true.

Syntax:

```
if condition1:
    statement(s)
elif condition2:
    statement(s)
elif condition3:
    statement(s)
```

- **condition1, condition2, condition3:** These are expressions that evaluate to either `True` or `False`.
- **statement(s):** These are the blocks of code that will be executed if the corresponding condition evaluates to `True`.

6.3.0.2 Using Multiple `elif` Clauses

The `elif` clause allows you to handle multiple potential cases in a clear and concise manner. The conditions are evaluated from top to bottom, and as soon as a `True` condition is found, the corresponding block of code is executed, and the rest of the conditions are skipped.

Example:

```
score = 85

if score >= 90:
    print("Grade: A")
elif score >= 80:
```



```
    print("Grade: B")
elif score >= 70:
    print("Grade: C")
elif score >= 60:
    print("Grade: D")
elif score < 60:
    print("Grade: F")
```

Grade: B

In this example, the program checks the score and assigns a grade based on predefined ranges. The conditions are evaluated in order: `score >= 90` is `False`, `score >= 80` is `True`, so the code within the `elif score >= 80` block is executed, resulting in the output “Grade: B”.

6.3.0.3 Combining `elif` with `else`

The `else` clause is optional but often used at the end of an `if-elif` chain to catch any cases that do not meet the previous conditions. This ensures that there is always a defined action for any possible input.

Example:

```
temperature = 75

if temperature > 85:
    print("It's hot outside.")
elif temperature > 65:
    print("The weather is nice.")
else:
    print("It's cold outside.")
```

The weather is nice.

In this example, the temperature is checked against three conditions. If the temperature is greater than 85, it prints “It’s hot outside.” If not, it checks if the temperature is greater than 65, printing “The weather is nice.” If neither condition is true, it defaults to printing “It’s cold outside.”

6.3.0.4 Practical Applications

Example: Speed Limit Checker

Let's create a program that checks a car's speed and prints a message based on the speed.

```
speed = 55

if speed > 80:
    print("You are speeding excessively.")
elif speed > 60:
    print("You are speeding.")
elif speed > 40:
    print("You are driving at a safe speed.")
else:
    print("You are driving below the speed limit.")
```

You are driving at a safe speed.

In this program, the speed is checked against multiple conditions to provide feedback on the driving speed. The conditions are evaluated in sequence, and the appropriate message is printed based on the speed.

6.3.0.5 Nested elif Clauses

Sometimes, you may need to nest `elif` clauses within other `if-elif-else` blocks to handle more complex decision-making processes.

Example: Admission Criteria

Let's create a program that checks admission criteria based on age and test scores.

```
age = 18
test_score = 85

if age >= 18:
    if test_score >= 90:
        print("Admitted with a scholarship.")
    elif test_score >= 75:
        print("Admitted.")
    else:
        print("Not admitted due to low test score.")
else:
```

```
print("Not admitted due to age requirement.")
```

Admitted.

In this example, the outer `if` statement checks if the age is 18 or older. If true, it enters a nested `if-elif-else` block that checks the test score. Depending on the test score, it prints the appropriate admission status. If the age condition is not met, it prints “Not admitted due to age requirement.”

The `elif` clause in Python is a powerful tool for handling multiple conditions in a clear and structured manner. By combining `if`, `elif`, and `else` clauses, you can create flexible decision-making processes in your programs. This allows your code to react dynamically to a wide range of inputs and conditions, making it more robust and versatile.

6.4 Conditional Expressions

Python also supports conditional expressions, which are a more concise way to write simple `if-else` statements.

Syntax:

```
value_if_true if condition else value_if_false
```

Example:

```
x = 5
result = "Positive" if x > 0 else "Non-positive"
print(result)
```

Positive

This example assigns the value “Positive” to the variable `result` if the condition `x > 0` is True, otherwise it assigns “Non-positive”. The output will be “Positive”.

6.5 Exercises

6.5.0.1 Exercise 1: Odd or Even

Write a program that checks if a number is odd or even using conditionals.

6.5.0.2 Exercise 2: Age Group

Write a program that categorizes a person's age group using only `if` and `else` statements.

6.5.0.3 Exercise 3: Positive, Negative, or Zero

Write a program that checks if a number is positive, negative, or zero. You must use at least one `elif` and one `else` statements.

7 Control Structures: Loops

Loops are one of the most powerful features in programming, enabling the repeated execution of a block of code. This repetition is often necessary for tasks such as processing elements in a collection, generating sequences of values, and executing complex algorithms. Understanding loops is essential for developing efficient and concise programs, as they reduce redundancy and simplify code that would otherwise require numerous repetitive statements.

7.0.0.1 Why Use Loops?

Loops are used to automate repetitive tasks, making code more efficient and easier to manage. Here are some key reasons to use loops:

1. **Reducing Code Duplication:** Without loops, you would need to write repetitive code multiple times. Loops allow you to write a block of code once and execute it multiple times.
2. **Dynamic Execution:** Loops can handle dynamic and varying data sizes. For instance, processing user input or reading data from a file where the number of elements is unknown beforehand.
3. **Complex Calculations:** Many algorithms, such as those used in searching and sorting, rely on loops to iterate through data and perform calculations.

7.0.0.2 Types of Loops in Python

Python supports two primary types of loops:

1. **while Loop:** Executes a block of code as long as a specified condition is `True`. The condition is checked before each iteration.
2. **for Loop:** Iterates over a sequence and executes a block of code for each item in the sequence.

7.1 The while Loop

The **while** loop is one of the fundamental control structures in Python, allowing the repeated execution of a block of code as long as a specified condition is **True**. This type of loop is particularly useful when the number of iterations is not known beforehand and depends on dynamic conditions during runtime.

7.1.0.1 Basic Syntax

The basic syntax of a **while** loop is:

```
while condition:
    statement(s)
```

- **condition**: An expression that evaluates to **True** or **False**.
- **statement(s)**: The block of code that will be executed repeatedly as long as the condition is **True**.

7.1.0.2 Example: Simple while Loop

Consider a simple example where we print numbers from 1 to 5:

```
i = 1
while i <= 5:
    print(i)
    i += 1
```

1
2
3
4
5

In this example, the variable **i** is initialized to 1. The **while** loop checks if **i** is less than or equal to 5. If the condition is **True**, it prints the value of **i** and increments **i** by 1. The loop continues until **i** becomes greater than 5.

7.1.0.3 Infinite Loops

A common pitfall with `while` loops is the creation of infinite loops, which occur when the loop's condition never becomes `False`. This can cause the program to run indefinitely, potentially causing it to become unresponsive.

Example of an infinite loop:

```
while True:
    print("This loop will run forever.")
```

To prevent infinite loops, ensure that the loop's condition will eventually become `False`.

7.1.0.4 `while` Loop with Else Clause

Python allows an optional `else` clause with `while` loops. The `else` block is executed when the loop condition becomes `False`.

Example:

```
i = 1
while i <= 5:
    print(i)
    i += 1
else:
    print("Loop ended naturally.")
```

```
1
2
3
4
5
Loop ended naturally.
```

In this example, the `else` block is executed after the `while` loop finishes executing, printing "Loop ended naturally."

7.1.0.5 Nested while Loops

`while` loops can be nested within other loops to handle more complex tasks, such as iterating over multi-dimensional data structures.

Example:

```
i = 1
while i <= 3:
    j = 1
    while j <= 3:
        print(j)
        j += 1
    i += 1
```

In this example, the outer `while` loop iterates over the variable `i` from 1 to 3, and for each iteration, the inner `while` loop iterates over the variable `j` from 1 to 3.

7.1.0.6 Controlling Loop Execution

Python provides several statements to control the execution of `while` loops:

- **break:** Terminates the loop prematurely.
- **continue:** Skips the rest of the loop body for the current iteration and proceeds to the next iteration.
- **pass:** Does nothing and is used as a placeholder in loops or functions where code will be added later.

Examples:

1. Using `break`:

```
i = 1
while i <= 10:
    if i == 5:
        break
    print(i)
    i += 1
```

1
2
3
4

2. Using continue:

```
i = 1
while i <= 10:
    if i % 2 == 0:
        i += 1
        continue
    print(i)
    i += 1
```

1
3
5
7
9

3. Using pass:

```
i = 1
while i <= 10:
    if i % 2 == 0:
        pass # Placeholder for future code
    else:
        print(i)
    i += 1
```

1
3
5
7
9

7.2 The for Loop

The for loop in Python is used for iterating over a sequence. Unlike the while loop, the for loop is preferred for its simplicity and readability when working with sequences. It provides a straightforward way to traverse elements in a collection, making it easier to write and understand.

7.2.0.1 Basic Syntax

The basic syntax of a `for` loop is:

```
for variable in sequence:
    statement(s)
```

- **variable:** A variable that takes the value of each item in the sequence during iteration.
- **sequence:** A collection of items to iterate over (e.g., list, tuple, string, range).
- **statement(s):** The block of code that will be executed for each item in the sequence.

7.2.0.2 Example: Iterating Over a List

Consider an example where we iterate over a list of numbers and print each number:

```
numbers = [1, 2, 3, 4, 5]
for num in numbers:
    print(num)
```

1
2
3
4
5

In this example, the `for` loop iterates over the list `numbers`, and the variable `num` takes the value of each item in the list during each iteration.

7.2.0.3 The `range()` Function

The `range()` function generates a sequence of numbers, which is particularly useful for creating loops that run a specific number of times. The `range()` function can take up to three arguments: `start`, `stop`, and `step`.

- **start:** The starting value of the sequence (inclusive). Default is 0.
- **stop:** The ending value of the sequence (exclusive).
- **step:** The difference between each pair of consecutive values. Default is 1.

Example using `range()`:

```
for i in range(1, 6):  
    print(i)
```

```
1  
2  
3  
4  
5
```

This loop prints numbers from 1 to 5. The `range(1, 6)` function generates the numbers 1, 2, 3, 4, and 5.

7.2.0.4 Iterating Over Strings

Strings are sequences of characters, and you can use a `for` loop to iterate over each character in a string.

Example:

```
message = "Hello, World!"  
for char in message:  
    print(char)
```

```
H  
e  
l  
l  
o  
,  
  
W  
o  
r  
l  
d  
!
```

In this example, the `for` loop iterates over the string `message`, and the variable `char` takes the value of each character in the string during each iteration.

7.2.0.5 Loop Control Statements

As with `while` loops, Python provides several statements to control the execution of `for` loops:

- **break:** Terminates the loop prematurely.
- **continue:** Skips the rest of the code inside the loop for the current iteration and moves to the next iteration.
- **pass:** Does nothing and is used as a placeholder in loops or functions where code will be added later.

Examples:

1. Using `break`:

```
for i in range(1, 10):  
    if i == 5:  
        break  
    print(i)
```

1
2
3
4

2. Using `continue`:

```
for i in range(1, 10):  
    if i % 2 == 0:  
        continue  
    print(i)
```

1
3
5
7
9

3. Using `pass`:

```
for i in range(1, 10):  
    if i % 2 == 0:  
        pass # Placeholder for future code
```

```
        else:
            print(i)

1
3
5
7
9
```

7.2.0.6 Practical Applications

for loops are used in a wide range of practical applications, from data processing to generating patterns. Here are some examples:

Example: Summing Numbers in a List

Calculate the sum of all numbers in a list:

```
numbers = [23, 45, 12, 89, 34]
total = 0

for num in numbers:
    total += num

print("Sum:", total)
```

Sum: 203

In this example, the `for` loop iterates over the list `numbers`, adding each number to the variable `total`, which holds the sum.

Example: Finding the Maximum Value

Find the maximum value in a list:

```
numbers = [23, 45, 12, 89, 34]
max_value = numbers[0]

for num in numbers:
    if num > max_value:
        max_value = num
```

```
print("Maximum value:", max_value)
```

Maximum value: 89

In this example, the `for` loop iterates over the list `numbers`, updating the variable `max_value` if a larger number is found.

7.3 Exercises

7.3.0.1 Exercise 1: Counting Down

Write a program that counts down from 10 to 1 using a `while` loop.

7.3.0.2 Exercise 2: Sum of Positive Numbers

Write a program that repeatedly asks the user for a number until they enter a negative number. The program should then print the sum of all positive numbers entered.

i `input()`

The `input()` function is used to display the given string to the console and wait for user input. The value entered by the user will be a string. Thus, the input must be converted to an `int` if an integer is needed in the rest of the code.

For this exercise, you can get the user input with the command

```
num = int(input("Enter a number (negative number to stop): "))
```

7.3.0.3 Exercise 3: Factorial Calculation

Write a program to calculate the factorial of a number using a `while` loop. Print the end result. Use an `input()` function as in Exercise 2.

7.3.0.4 Exercise 4: Guess the Number

Write a number guessing game where the program generates a random number between 1 and 100, and the user has to guess it. The program should give hints if the guess is too high or too low and keep asking until the user guesses correctly.

Use the following code to generate the random number:

```
import random
number = random.randint(1, 100)
```

7.3.0.5 Exercise 5: Sum of Even Numbers

Write a program to calculate the sum of all even numbers between 1 and 50.

7.3.0.6 Exercise 6: Prime Numbers

Write a program to print all prime numbers between 1 and 50 using a `for` loop and conditional statements.

8 Functions

Functions are integral to Python programming, and their importance cannot be overstated. At their core, functions encapsulate a specific block of code that performs a well-defined task. By isolating this task within a function, it can be easily reused, tested, and modified without affecting other parts of the program. This modular approach to programming offers several significant benefits, particularly in the context of writing efficient and maintainable code.

8.0.0.1 Reusability

One of the most compelling reasons to use functions is their ability to promote code reusability. Consider a scenario where you need to perform the same calculation at multiple points in your program. Without functions, you would have to write the same code multiple times, leading to redundancy and increased chances of errors. By defining a function, you can write the code once and reuse it wherever needed, thereby reducing duplication and ensuring consistency.

8.0.0.2 Modularity

Functions play a crucial role in breaking down complex problems into smaller, more manageable parts, a concept known as modularity. This makes it easier to develop, understand, and maintain your code. Each function is responsible for a specific task, allowing you to focus on one part of the problem at a time. This modular approach is particularly valuable in large programs, where managing the entire codebase as a single monolithic block would be overwhelming.

8.0.0.3 Abstraction

Functions enable abstraction by hiding the complexity of certain operations behind a simple interface. When you call a function, you don't need to know the intricate details of how it works—only what it does. This abstraction allows you to use complex operations without being bogged down by their implementation details.

Example: When you use Python's built-in `print()` function, you don't need to know how Python interacts with the system's output stream to display text on the screen. You just need to know that calling `print("Hello, World!")` will display the text. Similarly, when you write

your own functions, you can encapsulate complex logic and expose a simple interface to the rest of your program.

8.0.0.4 Maintainability

As software evolves, maintaining and updating code becomes increasingly important. Functions contribute to maintainability by isolating specific pieces of logic, making it easier to update or fix bugs in one part of the code without affecting others. When a function's logic needs to be updated, you only need to modify the function itself, rather than search through the entire codebase for instances where that logic was used.

8.0.0.5 Testing and Debugging

Functions simplify the testing and debugging process. Since functions are self-contained blocks of code, they can be tested independently of the rest of the program. This allows you to isolate and fix issues more efficiently. Additionally, functions with well-defined inputs and outputs are easier to verify for correctness.

8.1 Defining a Function

To define a function in Python, you use the `def` keyword, followed by the function name, parentheses `()`, and a colon `:`. The function body, which contains the code to be executed, is indented beneath the function definition.

Syntax:

```
def function_name(parameters):  
    # Function body  
    statement(s)
```

Example:

```
def greet():  
    print("Hello, World!")
```

In this example, the function `greet()` is defined to print the message “Hello, World!” when called.

8.2 Calling a Function

Once a function is defined, it can be called by using its name followed by parentheses ().

Example:

```
greet()
```

Hello, World!

Here, the `greet()` function is invoked, and the message is displayed.

8.3 Function Arguments

Functions can accept input values called arguments or parameters, allowing them to perform operations based on the input provided. These arguments are specified within the parentheses when defining the function.

8.3.1 Positional Arguments

Positional arguments are the most straightforward type of arguments. They are assigned to parameters based on their position in the function call.

Example:

```
def greet(name):  
    print(f"Hello, {name}!")  
  
greet("Alice")
```

Hello, Alice!

i f-strings

The argument in the `print()` function above is called an f-string. An f-string, introduced in Python 3.6, is a way to embed expressions inside string literals using curly braces {}. The `f` or `F` before the opening quote of the string indicates that it is an f-string. This allows you to include variables or expressions directly within the string, making string formatting more concise and readable.

Positional arguments are assigned to function parameters by their order of appearance. This means the first argument in the function call is passed to the first parameter, the second argument to the second parameter, and so on.

Example: Consider the following function definition:

```
def describe_person(name, age, city):  
    print(f"{name} is {age} years old and lives in {city}.")
```

If you call this function with the following positional arguments:

```
describe_person("Alice", 30, "New York")
```

The function execution will map the arguments as follows:

- `name` will be assigned the value "Alice"
- `age` will be assigned the value 30
- `city` will be assigned the value "New York"

The output will be:

```
Alice is 30 years old and lives in New York.
```

8.3.1.1 Importance of Order

Since positional arguments rely on the order in which they are passed, swapping the order can lead to incorrect or unintended results.

Example of Incorrect Order:

```
describe_person(30, "Alice", "New York")
```

In this case:

- `name` will be assigned 30
- `age` will be assigned "Alice"
- `city` will be assigned "New York"

This will produce the incorrect output:

```
30 is Alice years old and lives in New York.
```

8.3.2 Keyword Arguments

Keyword arguments allow you to specify the values for parameters by explicitly naming them in the function call, regardless of their order.

Example:

```
def greet(name, message):  
    print(f"{message}, {name}!")  
  
greet(name="Bob", message="Good morning")
```

Good morning, Bob!

The function can also be called with the order swapped but with the correct names.

```
greet(message="Good morning", name="Bob")
```

Good morning, Bob!

Here, the arguments are passed by specifying the parameter names, providing flexibility in the order of arguments.

8.3.2.1 Combining Positional and Keyword Arguments

Positional arguments can be combined with keyword arguments. However, when mixing them, positional arguments must always come before keyword arguments.

Example:

```
describe_person("Alice", age=30, city="New York")
```

Alice is 30 years old and lives in New York.

This call is valid and will produce the correct output, as the positional argument "Alice" is followed by keyword arguments for `age` and `city`.

8.3.3 Default Arguments

Default arguments allow you to define a function with default values for certain parameters. If no argument is provided for a parameter with a default value, the default is used.

Example:

```
def greet(name, message="Hello"):
    print(f"{message}, {name}!")

greet("Charlie")
greet("Charlie", "Goodbye")
```

```
Hello, Charlie!
Goodbye, Charlie!
```

In this example, the `message` parameter has a default value of `"Hello"`, which is used when no other value is provided.

8.3.4 Variable-Length Arguments

In Python, functions are not limited to accepting a fixed number of arguments. You can design functions to accept a variable number of arguments, allowing for greater flexibility and adaptability in different scenarios. Python provides two special types of arguments for this purpose: `*args` for positional arguments and `**kwargs` for keyword arguments. We will discuss `*args` here and come back to `**kwargs` later in [Chapter 11](#) after we discuss dictionaries.

8.3.4.1 `*args` – Variable-Length Positional Arguments

The `*args` syntax allows a function to accept any number of positional arguments. When you use `*args` in a function definition, Python collects all the positional arguments passed into the function and stores them in a tuple, which is an ordered and immutable collection of items (discussed in [Chapter 10](#)). When you define a function with `*args`, it can handle calls with any number of positional arguments—from zero to many.

Example:

```
def greet(*names):
    for name in names:
        print(f"Hello, {name}!")
```

Here's how this function works:

- If you call `greet("Alice", "Bob", "Charlie")`, the function will receive `names` as a tuple containing `("Alice", "Bob", "Charlie")`.
- The function will then iterate over the tuple and print a greeting for each name.

```
greet("Alice", "Bob", "Charlie")
```

```
Hello, Alice!
Hello, Bob!
Hello, Charlie!
```

When to Use `*args`

- **When the number of inputs is unknown:** If you're writing a function that might need to handle a varying number of inputs, `*args` is ideal.
- **For flexible APIs:** In some cases, you want to provide a flexible API that allows users to pass in any number of arguments without enforcing a strict parameter count.

Example: Imagine a function that calculates the total sum of an arbitrary number of numbers:

```
def calculate_sum(*numbers):
    total = 0
    for number in numbers:
        total += number
    return total

print(calculate_sum(1, 2, 3))
print(calculate_sum(5, 10, 15, 20))
```

```
6
50
```

This function can sum any number of integers or floats, demonstrating how `*args` enables flexible input handling.

8.4 Return Values

Functions can return values using the `return` statement. The value returned can be assigned to a variable for further use in the program.

Example:

```
def add(a, b):  
    return a + b  
  
result = add(3, 4)  
print(result)
```

7

Here, the `add` function returns the sum of `a` and `b`, which is then stored in the variable `result`.

8.4.1 Returning Multiple Values

In Python, a function can return more than one value at a time, which is a feature that adds considerable flexibility to the way functions are used. When a function returns multiple values, it does so by returning a tuple. This allows you to return several related pieces of data from a single function call, without the need to explicitly create and manage a complex data structure.

8.4.1.1 How Multiple Return Values Work

When a function is designed to return multiple values, it simply lists the values after the `return` keyword, separated by commas. Python automatically packages these values into a tuple. The caller of the function can then unpack this tuple (see [Chapter 10](#)) into separate variables, each receiving one of the returned values.

Example:

Consider the following example:

```
def add_subtract(a, b):  
    return a + b, a - b
```

In this function:

- `a + b` computes the sum of the two arguments `a` and `b`.
- `a - b` computes the difference between `a` and `b`.
- Both values are returned together as a tuple.

When this function is called:

```
sum_result, diff_result = add_subtract(10, 5)
print(sum_result)
print(diff_result)
```

15

5

Here, the tuple (15, 5) is returned, and it is immediately unpacked into the variables `sum_result` and `diff_result`. This allows the caller to easily access each result separately.

8.4.1.2 Benefits of Returning Multiple Values

Returning multiple values from a function is particularly advantageous in situations where a single calculation or process naturally produces more than one result.

Example 1: Mathematical Operations

```
def calculate_area_perimeter(length, width):
    area = length * width
    perimeter = 2 * (length + width)
    return area, perimeter

area, perimeter = calculate_area_perimeter(5, 3)
print(f"Area: {area}, Perimeter: {perimeter}")
```

Area: 15, Perimeter: 16

In this example, the function `calculate_area_perimeter` returns both the area and perimeter of a rectangle. This allows the caller to retrieve and use both pieces of information with a single function call.

Example 2: Finding Extremes

```
def find_extremes(numbers):
    return max(numbers), min(numbers)

maximum, minimum = find_extremes([10, 20, 5, 30, 15])
print(f"Maximum: {maximum}, Minimum: {minimum}")
```

Maximum: 30, Minimum: 5

Here, the function `find_extremes` computes and returns both the maximum and minimum values from a list of numbers, making it easy to handle both results simultaneously.

8.4.1.3 Unpacking Returned Values

When a function returns multiple values, the caller can unpack these values into individual variables. This is done by assigning the function call to a tuple of variables corresponding to the number of values returned.

Example:

```
sum_result, diff_result = add_subtract(10, 5)
```

In this case, the returned tuple `(15, 5)` is unpacked into `sum_result` and `diff_result`, making the individual results accessible immediately.

8.4.1.4 Single Return Value with a Tuple

If needed, the function can return a tuple directly without unpacking it in the calling code. This can be useful when the function's result is intended to be passed around or used as a single entity.

Example:

```
result = add_subtract(10, 5)
print(result)
```

`(15, 5)`

Here, the entire tuple `(15, 5)` is returned as a single object and can be used as such.

8.5 Best Practices in Function Design

Designing functions effectively is crucial for writing clean, maintainable, and efficient code. Well-designed functions not only make your code easier to understand and use but also reduce the likelihood of bugs and make it easier to extend and modify your programs. Below are some best practices to follow when designing functions in Python.

8.5.0.1 Use Descriptive Names

A function's name should clearly and concisely describe what the function does. Descriptive names make the code more readable and self-documenting, allowing others (and your future self) to understand the purpose of the function without needing extensive comments or external documentation.

Example:

```
def calculate_average(scores):  
    return sum(scores) / len(scores)
```

In this example, the function name `calculate_average` clearly indicates that the function computes the average of a list of scores. Anyone reading the code can immediately grasp the function's purpose without needing to examine its implementation.

Why This Matters:

- **Readability:** Descriptive names make your code easier to read and understand.
- **Maintainability:** When functions are clearly named, it's easier to locate and update the appropriate function when changes are needed.
- **Collaboration:** In team settings, clear function names help other developers understand and use your code correctly, reducing the potential for errors.

8.5.0.2 Keep Functions Focused

A well-designed function should perform a single, clearly defined task or a set of closely related tasks. This practice, often referred to as the “Single Responsibility Principle,” ensures that your functions are simple, modular, and reusable.

Example:

```
def read_file(file_path):  
    # Processing logic here  
    pass  
  
def process_data(data):  
    # Processing logic here  
    pass  
  
def write_file(file_path, data):  
    # Processing logic here  
    pass
```

i pass

In Python, the `pass` keyword is used as a placeholder in your code. It allows you to write syntactically correct code blocks where no action is required. Essentially, `pass` does nothing when executed. It's particularly useful in situations where you have a code structure that requires a statement, but you haven't decided what the specific code should be yet.

Think of `pass` as a placeholder in your code that let you outline the structure of your program without having to fill in the details immediately. This can be very helpful during the initial stages of writing or when planning out complex code.

In this example, each function is focused on a specific task: reading a file, processing data, and writing to a file. By keeping each function focused, the code becomes more modular and easier to maintain.

Why This Matters:

- **Simplicity:** Functions that do one thing are easier to understand, test, and debug.
- **Reusability:** Focused functions are more likely to be reusable in different parts of your program or even in other projects.
- **Maintainability:** When functions are responsible for a single task, changes to one part of the code are less likely to have unintended side effects on other parts.

8.5.0.3 Avoid Side Effects

Side effects occur when a function modifies some state or interacts with outside elements like global variables, files, or databases, which are not directly related to its inputs and outputs. While side effects are sometimes necessary, minimizing them helps ensure that functions are predictable and easier to test.

Recall that global and local variables were first discussed in [Section 4.1](#).

Example of a Function with Side Effects:

```
total = 0

def add_to_total(amount):
    global total
    total += amount
```

In this example, the function `add_to_total` modifies the global variable `total`, which is a side effect. This can lead to unpredictable behavior, especially in larger programs where the global state is modified by multiple functions.

Better Approach:

```
def calculate_new_total(current_total, amount):  
    return current_total + amount
```

In this revised example, the function `calculate_new_total` returns a new total based on the inputs without modifying any external state. The function is now pure, meaning its output depends only on its inputs and has no side effects.

Why This Matters:

- **Predictability:** Functions without side effects are easier to reason about because they produce the same output for the same input every time.
- **Testability:** Pure functions are easier to test since you don't need to set up or tear down any external state.
- **Debugging:** Functions that don't cause side effects are less likely to introduce hidden bugs related to state changes elsewhere in the program.

8.5.0.4 Document Your Functions

Even with descriptive names, adding docstrings to your functions is a good practice. A docstring provides a description of the function's purpose, parameters, and return values, making it easier for others to use your function correctly.

Example:

```
def calculate_average(scores):  
    """  
    Calculates the average of a list of scores.  
  
    Parameters:  
    scores (list of int/float): A list of numeric scores.  
  
    Returns:  
    float: The average of the scores.  
    """  
    return sum(scores) / len(scores)
```

Why This Matters:

- **Clarity:** Docstrings clarify how to use the function, what inputs it expects, and what outputs it provides.
- **Collaboration:** Docstrings make it easier for others to understand and use your code.

- **Self-Documentation:** Well-documented functions serve as a form of in-code documentation, reducing the need for external documentation.

By following these best practices in function design—using descriptive names, keeping functions focused, avoiding side effects, and documenting your functions—you can create Python code that is easier to read, maintain, and extend. These practices not only improve the quality of your code but also make it more robust and reliable, facilitating collaboration and reducing the likelihood of bugs.

8.6 Exercises

8.6.0.1 Exercise 1: Simple Greeting Function

Write a function `greet_user` that takes a user's name as input and prints a greeting message.

8.6.0.2 Exercise 2: Arithmetic Function

Write a function `calculate` that takes two numbers and returns their sum, difference, product, and quotient.

8.6.0.3 Exercise 3: Temperature Conversion

Write a function `convert_temperature` that converts a temperature from Celsius to Fahrenheit.

8.6.0.4 Exercise 4: Flexible Function

Write a function `summarize` that can take any number of numerical arguments and returns their sum and average.

9 Introduction to Modules and Libraries

In Python, modules and libraries play a crucial role in organizing code, promoting reusability, and extending the functionality of your programs. A module is a file containing Python definitions and statements that can be imported into your scripts, while a library is a collection of modules that provide related functionality. Understanding how to use both built-in and custom modules will enhance your ability to write efficient and maintainable code.

9.1 Understanding Modules

Modules allow you to break your code into separate files, making it easier to manage and understand. A module can contain functions, classes, variables, and runnable code. By organizing your code into modules, you can create reusable components that can be easily imported into other projects.

Example: Creating a Simple Module

Let's start by creating a simple module. Suppose we have a file named `greetings.py`:

```
# greetings.py

def say_hello(name):
    """Print a friendly greeting."""
    return f"Hello, {name}!"

def say_goodbye(name):
    """Print a farewell message."""
    return f"Goodbye, {name}!"
```

This module contains two functions: `say_hello` and `say_goodbye`. You can import this module into another script and use these functions.

Example: Using the Module

```
# main.py
```

```
import greetings

print(greetings.say_hello("Alice"))
print(greetings.say_goodbye("Alice"))
```

Running `main.py` would output:

```
Hello, Alice!
Goodbye, Alice!
```

i Note

The `.py` file that you are importing as a module needs to be in the same directory as the file you are importing the module into. In the example above, `greetings.py` and `main.py` needs to be in the same directory.

Here, the `greetings` module is imported, and its functions are called to print personalized messages. This simple example demonstrates the power of modules in keeping your code organized and reusable.

9.2 Importing Modules

Modules are imported to gain access to their functions, classes, and variables. Python provides several methods for importing modules, each offering different levels of control and flexibility depending on your use case. Let's explore these methods in detail.

9.2.0.1 Standard Import

The most common way to import a module is by using a standard import statement. This method imports the entire module, making all of its contents accessible. However, to use any function or class from the module, you must prefix it with the module's name. This ensures that there are no conflicts between functions or variables from different modules, as everything is neatly contained within its namespace.

Example: Using `math` Module

```
import math

# Accessing the square root function from the math module
result = math.sqrt(16)
```

```
print(result)
```

4.0

In this example, the `math` module is imported, and we use the `sqrt` function by referencing it through the module's name.

9.2.0.2 Importing Specific Functions or Classes

Sometimes, you may only need a specific function, class, or variable from a module. Python allows you to import only what you need, which can make your code cleaner and more efficient. When you import specific items, you don't need to prefix them with the module's name, as they are directly accessible.

Example: Importing `sqrt` from `math`

```
from math import sqrt

# Directly using the imported sqrt function
result = sqrt(16)
print(result)
```

4.0

This method is particularly useful when you need to use a specific function frequently and want to avoid repeatedly typing the module name. However, it's important to be cautious with this approach, as it can lead to naming conflicts if different modules contain functions or variables with the same name.

You can import multiple items from a module by separating them with commas:

```
from math import sqrt, pi

# Using both imported items directly
print(sqrt(25))
print(pi)
```

5.0

3.141592653589793

9.2.0.3 Importing All Names with *

In certain cases, you might want to import everything from a module, making all its functions and variables directly accessible without needing to prefix them with the module's name. This can be done using the `*` wildcard. However, this method is generally discouraged because it can lead to unexpected name conflicts and make the code harder to understand and debug.

Example: Importing All from `math`

```
from math import *

# Using functions and constants directly
print(sqrt(16))
print(pi)
```

```
4.0
3.141592653589793
```

While this approach can save typing and is convenient in small scripts or interactive sessions, it is not recommended for larger programs or libraries, where clarity and maintainability are critical.

9.2.0.4 Import with Alias

Modules can sometimes have long names, or you may want to avoid conflicts between modules that share the same name. In these situations, you can import a module under a different name, known as an alias. This allows you to reference the module using a shorter or more descriptive name, making your code more concise and readable.

Example: Importing `math` with an Alias

```
import math as m

# Using the alias to access functions from the math module
result = m.sqrt(16)
print(result)
```


```
4.0
```

Aliases are particularly useful when working with libraries that have long names or when you frequently use a module in your code. They help in maintaining readability while reducing the amount of typing required.

Example: Avoiding Name Conflicts with Aliases

```
import matplotlib.pyplot as plt

# Using aliases in place of longer name
plt.plot([1, 2, 3], [4, 5, 6])
```



i Installing matplotlib

Most installs of Python do not include `matplotlib` as a base library. You may need to install it. Open up a terminal. Run the following code to install `matplotlib`.

```
pip install matplotlib --user
```

Note that if you are using a virtual environment, you will need to make sure the virtual environment is activated first.

9.2.0.5 Importing Modules from a Package

Python also supports importing modules from a package, which is a collection of modules organized under a common namespace. Packages help in organizing related modules and can be imported similarly to regular modules.

Example: Importing from a Package

```
from os import path

# Using the path module from the os package
print(path.exists("example.txt"))
```

False

In this example, the `path` module is imported from the `os` package, and its `exists` function is used to check if a file exists. Packages are an essential part of Python's ecosystem, allowing you to organize and distribute your code effectively.

9.3 Built-in Libraries

Python comes with a rich set of built-in libraries that cover a wide range of functionalities, from mathematical operations to file handling and beyond. Let's explore a few common libraries that you will frequently use in your programming journey.

The math Library

As you have already seen, the `math` library provides mathematical functions and constants.

Example: Basic Usage of math Library

```
import math

# Using constants
print(math.pi)

# Using functions
print(math.factorial(5))
print(math.sqrt(25))
```

```
3.141592653589793
120
5.0
```

9.3.0.1 The random Library

The `random` library is used for generating random numbers and making random selections.

Example: Using random Library

```
import random

# Generating a random number between 1 and 10
print(random.randint(1, 10))

# Picking a random choice from a list
choices = ['apple', 'banana', 'cherry']
print(random.choice(choices))
```

```
8
banana
```

This library is particularly useful in simulations, games, and scenarios where randomness is needed.

9.4 Creating Custom Modules

Creating your own modules allows you to encapsulate code that can be reused across multiple projects. As you progress in your coding journey, you'll find this practice invaluable for maintaining clean and organized code.

Example: Building a Utility Module

Let's create a module named `utils.py` that contains some utility functions:

```
# utils.py

def reverse_string(s):
    """Reverse a given string."""
    return s[::-1]
```

```
def is_palindrome(s):  
    """Check if a string is a palindrome."""  
    return s == s[::-1]
```

You can now import and use these functions in any script:

```
# main.py  
  
import utils  
  
word = "level"  
print(utils.reverse_string(word))  
print(utils.is_palindrome(word))
```

```
level  
True
```

9.5 Best Practices for Modules and Libraries

When working with modules and libraries in Python, following best practices is essential for creating code that is both maintainable and user-friendly. One of the most important practices is to use descriptive names for your modules. The name of a module should clearly convey its purpose and functionality, making it easier for others (and yourself) to understand what the module does at a glance. For example, a module named `math_operations` is far more informative than a generic name like `utils`, as it immediately indicates that the module contains functions related to mathematical operations. Descriptive naming helps prevent confusion, especially in larger projects where multiple modules are used.

In addition to naming, keeping functions within a module focused on a single, well-defined task is crucial. Each function should do one thing and do it well. This approach not only makes your code easier to test and debug but also enhances its reusability. For instance, a function that calculates the average of a list of numbers should not also be responsible for reading the numbers from a file. By adhering to the principle of single responsibility, you ensure that each function is modular, making it easier to mix and match functions across different modules and projects.

Documentation is another critical aspect of writing good modules. Providing clear and comprehensive docstrings for your modules, functions, and classes is essential for making your code accessible to others. Docstrings should explain what the code does, how to use it, and any important details that users need to be aware of. Well-documented code not only helps others understand and use your modules but also serves as a valuable reference for yourself when you

return to the code after some time. Good documentation is a sign of professionalism and care in coding, making your work more reliable and easier to maintain.

Finally, it is important to avoid side effects in your modules. Side effects occur when a module executes code automatically upon being imported, such as modifying global variables or performing I/O operations. This can lead to unpredictable behavior and bugs, especially if the user is unaware of these side effects. To prevent this, modules should generally be passive, only providing functions and classes without executing any code unless explicitly intended.

9.6 Exercises

9.6.0.1 Exercise 1: Creating a Module

Create a module `arithmetic.py` that contains functions for addition, subtraction, multiplication, and division. Write a script that imports this module and performs these operations on user-provided inputs.

9.6.0.2 Exercise 2: Using Built-in Libraries

Write a script that uses the `random` library to generate a random integer and then uses the `math` library to find the square root of this integer.

9.6.0.3 Exercise 3: Module Composition

Create a module `geometry.py` that contains functions to calculate the area and perimeter of different shapes, such as rectangles, circles, and triangles. Write a script that imports this module and allows the user to input the dimensions of a shape, then outputs the calculated area and perimeter.

9.6.0.4 Exercise 4: Creating a Custom Math Library

Develop a custom math module `custom_math.py` that includes functions for basic arithmetic operations, factorial calculation, and prime number checking. Extend the module by adding a function to calculate the greatest common divisor (GCD) of two numbers. Write a script to demonstrate the usage of each function in the module.

10 Data Structures: Lists and Tuples

In the previous chapters, we introduced the basic elements of programming in Python, including variables, data types, control structures, and functions. Now, we turn our attention to an essential topic in computing — data structures. Data structures allow us to organize and store data in ways that enable efficient access and modification. In this chapter, we will explore two foundational data structures in Python: **Lists** and **Tuples**.

10.1 Lists

In Python, **lists** are one of the most commonly used data structures due to their flexibility and ease of use. A list is a **mutable**, **ordered** collection of elements, which means the elements in a list can be changed after the list is created, and they are stored in a specific order. This makes lists ideal for storing sequences of data that might need to be altered during the execution of a program.

10.1.1 Creating Lists

Lists in Python are defined using square brackets (`[]`), with each element separated by a comma. A list can contain elements of any data type, including integers, floats, strings, Booleans, and even other lists.

Examples:

```
# A list of integers
integer_list = [1, 2, 3, 4, 5]

# A list of mixed data types
mixed_list = [42, "apple", 3.14, True]

# A list of lists
lists_list = [integer_list, mixed_list]
```

The flexibility of lists allows you to store a wide variety of data in a single collection, making them useful in many programming contexts, such as storing records, handling input, and building dynamic datasets.

10.1.2 Accessing Elements in a List

Each element in a list is associated with an **index**—an integer representing the element’s position in the list. In Python, indices start at 0, meaning the first element is accessed with index 0, the second element with index 1, and so on. Negative indices can also be used to access elements from the end of the list, with -1 referring to the last element, -2 to the second-to-last, and so on.

Examples:

```
# Accessing elements by positive index
my_list = [10, 20, 30, 40]
print(my_list[0]) # Output: 10
print(my_list[2]) # Output: 30

# Accessing elements by negative index
print(my_list[-1]) # Output: 40
print(my_list[-2]) # Output: 30
```

```
10
30
40
30
```

10.1.3 Modifying Lists

One of the most powerful features of lists is their mutability. This means that after a list is created, its elements can be changed, added, or removed without creating a new list. There are several ways to modify lists in Python.

10.1.3.1 Changing Elements

You can change individual elements in a list by assigning a new value to a specific index:

```
my_list = [1, 2, 3]
my_list[1] = 99
print(my_list)
```

```
[1, 99, 3]
```


10.1.3.2 Adding Elements

You can add elements to a list using the `append()` method (to add a single element) or the `extend()` method (to add multiple elements):

```
# Adding a single element
my_list = [1, 2, 3]
my_list.append(4)
print(my_list)

# Adding multiple elements
my_list.extend([5, 6])
print(my_list)
```

```
[1, 2, 3, 4]
```

```
[1, 2, 3, 4, 5, 6]
```

You can also use the `insert()` method to add an element at a specific index:

```
my_list = [1, 2, 3]
my_list.insert(1, "new")
print(my_list)
```

```
[1, 'new', 2, 3]
```

Note

In Python, a **method** is a function that is associated with an object. Every method is a function, but not every function is a method. The key difference is that methods are called on objects, and they are designed to perform actions related to those objects. The syntax for calling a method involves the dot notation, where you first specify the object and then call the method.

Methods are tightly bound to the object they belong to and often manipulate or interact with the data stored in the object. Python has built-in methods for its standard data types like lists, strings, dictionaries, etc.

10.1.3.3 Removing Elements

Elements can be removed from a list using the `remove()` method (to remove the first occurrence of a specific element) or the `pop()` method (to remove an element by its index):

```
# Removing an element by value
my_list = [1, 2, 3, 2]
my_list.remove(2)
print(my_list)

# Removing an element by index
my_list.pop(1) #note that this returns the value being removed
print(my_list)
```

[1, 3, 2]

[1, 2]

To clear all elements from a list, use the `clear()` method:

```
my_list = [1, 2, 3]
my_list.clear()
print(my_list)
```

[]

10.1.4 Slicing Lists

In addition to accessing individual elements, Python allows you to **slice** lists, which means extracting a portion of the list to create a new list. Slicing is done using the colon (`:`) operator, with the format `list[start:end]`. The **start** index is inclusive, while the **end** index is exclusive.

Examples:

```
my_list = [10, 20, 30, 40, 50]

# Extract elements from index 1 to 3 (2nd to 4th elements)
print(my_list[1:4]) # Output: [20, 30, 40]

# Extract elements from the start to index 3
print(my_list[:4]) # Output: [10, 20, 30, 40]

# Extract elements from index 2 to the end
print(my_list[2:]) # Output: [30, 40, 50]
```

```
[20, 30, 40]
[10, 20, 30, 40]
[30, 40, 50]
```

Slicing can also be used with a step value, which specifies how many elements to skip between items:

```
# Extract every other element
print(my_list[::2])
```

```
[10, 30, 50]
```

10.1.5 Common List Operations

Python provides several built-in functions and operators that can be applied to lists. Below are some of the most commonly used list operations:

1. **Checking Length:** You can find the number of elements in a list using the `len()` function:

```
my_list = [10, 20, 30]
print(len(my_list))
```

```
3
```

2. **Membership Testing:** You can check whether an element is in a list using the `in` keyword:

```
my_list = [10, 20, 30]
print(20 in my_list)
```

```
True
```

3. **Sorting Lists:** Lists can be sorted in place using the `sort()` method, or a sorted copy of the list can be returned using the `sorted()` function:

```
my_list = [3, 1, 4, 1, 5, 9]
my_list.sort() # Sort the list in place
print(my_list)
```

```
[1, 1, 3, 4, 5, 9]
```

4. **Reversing a List:** You can reverse the order of elements in a list using the `reverse()` method:

```
my_list = [1, 2, 3]
my_list.reverse()
print(my_list)
```

```
[3, 2, 1]
```

10.1.6 Iterating Over Lists

Lists are iterable, which means you can loop through the elements using a `for` loop:

```
my_list = ["apple", "banana", "cherry"]
for fruit in my_list:
    print(fruit)
```

```
apple
banana
cherry
```

You can also iterate over both the index and the element using the `enumerate()` function:

```
my_list = ["apple", "banana", "cherry"]
for index, fruit in enumerate(my_list):
    print(f"Index {index}: {fruit}")
```

10.1.6.1 Nesting Lists

Lists can contain other lists as elements, allowing you to create more complex data structures like matrices or grids. This is known as **nesting**.

Example:

```
matrix = [[1, 2, 3], [4, 5, 6], [7, 8, 9]]

# Accessing the second element of the first list
print(matrix[0][1]) # Output: 2

# Iterating over nested lists
```

```
for row in matrix:
    print(row)
```

```
2
[1, 2, 3]
[4, 5, 6]
[7, 8, 9]
```

10.1.6.2 Lists and Built-in Libraries

Lists can be used effectively with common Python libraries. For example, the `random` module allows for random selection of elements from a list:

Example:

```
import random

fruits = ["apple", "banana", "cherry", "date"]
print(random.choice(fruits)) # Randomly selects and prints one fruit
```

date

10.2 Tuples

While **lists** offer flexibility through their mutability, Python also provides **tuples**, which are similar in structure but **immutable**. Once a tuple is created, its elements cannot be changed. This immutability makes tuples useful for representing fixed collections of items that should not or cannot change throughout the execution of a program. Tuples are ideal when you need to ensure data integrity, such as coordinates, configuration settings, or when passing multiple values from functions where immutability is expected.

10.2.1 Creating Tuples

Tuples are defined using parentheses `()` and can hold elements of any data type, much like lists. However, since they are immutable, you cannot modify the elements of a tuple after it is created.

Syntax:

```
# Creating a tuple
my_tuple = (10, 20, 30)
```

It's also possible to create tuples without using parentheses, though this is less common:

```
my_tuple = 10, 20, 30
```

Tuples can contain elements of different types:

```
mixed_tuple = (1, "apple", 3.14, True)
```

Tuples with a single element need to have a comma after the element to avoid confusion with parentheses used in expressions:

```
single_element_tuple = (5,)
print(type(single_element_tuple))
```

```
<class 'tuple'>
```

10.2.2 Accessing Tuple Elements

Like lists, tuples are indexed starting at 0, and individual elements can be accessed using their index. However, since tuples are immutable, you cannot modify their elements.

Examples:

```
my_tuple = (10, 20, 30, 40)

# Accessing the first element
print(my_tuple[0])

# Accessing the last element using negative indexing
print(my_tuple[-1])
```

```
10
40
```

You can slice tuples just like lists, returning a new tuple containing the specified range of elements:

```
# Slicing a tuple
print(my_tuple[1:3])
```

(20, 30)

10.2.3 Tuple Operations

Although tuples are immutable, there are still several operations you can perform on them:

1. **Length of a Tuple:** You can find the number of elements in a tuple using the `len()` function:

```
my_tuple = (10, 20, 30)
print(len(my_tuple))
```

3

2. **Membership Testing:** Use the `in` keyword to check whether an element exists in a tuple:

```
my_tuple = (10, 20, 30)
print(20 in my_tuple)
```

True

3. **Iterating Over a Tuple:** Tuples are iterable, so you can loop through their elements using a `for` loop:

```
for item in my_tuple:
    print(item)
```

10

20

30

4. **Indexing and Slicing:** Like lists, tuples support indexing and slicing:

```
my_tuple = (10, 20, 30, 40)
print(my_tuple[1:3])
```

(20, 30)

10.2.4 Nested Tuples

Tuples can be **nested** inside other tuples, allowing you to create multi-level data structures. This is useful when organizing complex data, such as in multidimensional arrays or coordinate systems.

Example:

```
nested_tuple = ((1, 2), (3, 4), (5, 6))

# Accessing the first tuple
print(nested_tuple[0])

# Accessing an element from a nested tuple
print(nested_tuple[0][1])
```

```
(1, 2)
2
```

10.2.5 Unpacking Tuples

One of the most powerful features of tuples is **tuple unpacking**, which allows you to assign the elements of a tuple to individual variables in a single statement.

Example:

```
# Unpacking a tuple into variables
coordinates = (10, 20)
x, y = coordinates
print(x)
print(y)
```

```
10
20
```

Tuple unpacking is especially useful when functions return multiple values in a tuple, allowing you to capture and work with these values directly.

Example:


```
def rectangle_properties(length, width):
    area = length * width
    perimeter = 2 * (length + width)
    return (area, perimeter)

# Unpacking the returned tuple into variables
area, perimeter = rectangle_properties(5, 10)
print(f"Area: {area}, Perimeter: {perimeter}")
```

Area: 50, Perimeter: 30

10.2.6 Tuple Methods

Since tuples are immutable, they have fewer methods compared to lists. However, they do provide two useful methods:

1. **count()**: Returns the number of times a specified value appears in the tuple.

```
my_tuple = (1, 2, 2, 3, 2)
print(my_tuple.count(2))
```

3

2. **index()**: Returns the index of the first occurrence of a specified value.

```
my_tuple = (1, 2, 3)
print(my_tuple.index(2))
```

1

10.3 List Comprehension

In Python, **list comprehension** provides a concise way to generate lists. It offers a more readable and often more efficient alternative to using loops and **append()** to build lists. List comprehension allows you to apply an expression to each element in a sequence and optionally include conditional statements to filter elements. This compact syntax makes it easy to create lists based on existing sequences or from operations.

10.3.1 Basic Syntax of List Comprehension

The basic syntax of list comprehension is:

```
[expression for item in iterable]
```

- **expression:** This is the operation or value you want to apply to each item in the iterable.
- **item:** Each element from the iterable (e.g., a list, tuple, or string).
- **iterable:** The sequence of elements to loop over (e.g., a list, range, or other iterable object).

This simple form of list comprehension generates a new list by evaluating the expression for each element in the iterable.

Example: Creating a list of squares

```
squares = [x**2 for x in range(5)]  
print(squares)
```

```
[0, 1, 4, 9, 16]
```

In this example, for each value `x` in `range(5)`, Python evaluates `x**2` and adds the result to the new list. The result is a list of the squares of the numbers from 0 to 4.

10.3.2 List Comprehension with Conditional Logic

List comprehension can include **conditional logic**, allowing you to filter elements or apply an operation only when a condition is met. The syntax for adding a condition is as follows:

```
[expression for item in iterable if condition]
```

The **condition** is a logical statement that is evaluated for each item in the iterable. Only items for which the condition evaluates to `True` are included in the resulting list.

Example: Filtering even numbers

```
even_numbers = [x for x in range(10) if x % 2 == 0]  
print(even_numbers)
```

```
[0, 2, 4, 6, 8]
```

Here, only numbers that satisfy the condition `x % 2 == 0` (i.e., the even numbers) are included in the resulting list.

You can also use an `if-else` expression within list comprehension for more complex logic:

```
[expression_if_true if condition else expression_if_false for item in iterable]
```

Example: Labeling even and odd numbers

```
labels = ["even" if x % 2 == 0 else "odd" for x in range(5)]  
print(labels)
```

```
['even', 'odd', 'even', 'odd', 'even']
```

In this case, the comprehension adds the string `"even"` to the list if `x` is divisible by 2 and `"odd"` otherwise.

10.3.3 Nested List Comprehension

List comprehension can also be nested to create lists from multidimensional structures, such as matrices. Nested list comprehensions are powerful but can become harder to read if not used carefully.

The basic syntax for nested list comprehension is:

```
[expression for item1 in iterable1 for item2 in iterable2]
```

Example: Flattening a matrix (a list of lists)

```
matrix = [[1, 2, 3], [4, 5, 6], [7, 8, 9]]  
flat_list = [num for row in matrix for num in row]  
print(flat_list)
```

```
[1, 2, 3, 4, 5, 6, 7, 8, 9]
```

In this example, we loop over each row in the matrix, and for each row, we loop over each number, adding it to a single list (`flat_list`).

Example: Multiplying elements in a matrix by 2

```
matrix = [[1, 2, 3], [4, 5, 6], [7, 8, 9]]
doubled_matrix = [[num * 2 for num in row] for row in matrix]
print(doubled_matrix)
```

```
[[2, 4, 6], [8, 10, 12], [14, 16, 18]]
```

In this case, the list comprehension multiplies each element in the matrix by 2, resulting in a new matrix where all values are doubled.

10.3.4 List Comprehension with Built-in Functions

You can combine list comprehensions with built-in Python functions to perform more complex transformations and calculations. This is a common pattern when dealing with operations such as string manipulation, mathematical calculations, or other functional transformations.

Example: Using the `len()` function

```
words = ["apple", "banana", "cherry"]
word_lengths = [len(word) for word in words]
print(word_lengths)
```

```
[5, 6, 6]
```

Here, the list comprehension applies the `len()` function to each word in the list `words`, resulting in a list of the word lengths.

Example: Using `sum()` with list comprehension Suppose we have a list of lists representing test scores for different students. We can calculate the sum of each student's test scores using list comprehension:

```
scores = [[75, 80, 85], [60, 70, 75], [90, 95, 100]]
total_scores = [sum(student_scores) for student_scores in scores]
print(total_scores)
```

```
[240, 205, 285]
```

In this example, `sum()` calculates the total score for each student, and list comprehension gathers these sums into a new list `total_scores`.

10.3.5 List Comprehension vs. Loops

List comprehension is often favored over traditional loops because it provides a more **compact** and **readable** syntax. However, there are cases where a loop might be more appropriate, especially when the logic is complex, or when you need to modify elements in place.

Example: Traditional loop

```
squares = []
for x in range(5):
    squares.append(x**2)
print(squares)
```

[0, 1, 4, 9, 16]

Equivalent list comprehension:

```
squares = [x**2 for x in range(5)]
print(squares)
```

[0, 1, 4, 9, 16]

List comprehension is faster for simple operations because it avoids the overhead of repeatedly calling `append()` and performing function calls. However, for more complex logic, such as multiple conditional statements or nested loops, the clarity of traditional loops might outweigh the brevity of list comprehension.

10.3.6 List Comprehension with External Libraries

List comprehension can be used effectively with other basic libraries like `math` or `random`.

Example: Using `math.sqrt()` with list comprehension

```
import math
numbers = [1, 4, 9, 16, 25]
square_roots = [math.sqrt(num) for num in numbers]
print(square_roots)
```

[1.0, 2.0, 3.0, 4.0, 5.0]

In this example, `math.sqrt()` is applied to each number in the list `numbers`, resulting in a list of square roots.

Example: Using `random.randint()` You can also use list comprehension with the `random` module to generate random numbers:

```
import random
random_numbers = [random.randint(1, 100) for _ in range(5)]
print(random_numbers)
```

```
[43, 72, 17, 4, 3]
```

Here, the `_` is a placeholder variable (indicating that the value is not important), and `random.randint()` generates a random integer for each iteration.

10.3.7 Limitations of List Comprehension

While list comprehension is a powerful tool, there are some cases where it may not be the best choice:

- **Readability:** Overusing or nesting list comprehensions can make code difficult to read and maintain, especially when working with complex logic. In such cases, using traditional loops or helper functions might result in clearer and more maintainable code.
- **Complex operations:** When performing complex operations with multiple steps, it's often better to use traditional loops to avoid confusion and improve code clarity.
- **Memory efficiency:** Since list comprehensions create a new list in memory, they might not be suitable for extremely large datasets. In such cases, consider using generator expressions (which we will cover in later sections) to optimize memory usage.

10.3.8 Example: Practical Applications of List Comprehension

To conclude, let's explore a practical application of list comprehension in data processing.

Example: Filtering and transforming data Suppose we are processing a list of student scores, and we want to filter out scores below 50 and increase the remaining scores by 10%. We can accomplish this efficiently with list comprehension.

```
scores = [45, 67, 85, 30, 78, 92, 40]
adjusted_scores = [score * 1.1 for score in scores if score >= 50]
print(adjusted_scores)
```

```
[73.7, 93.50000000000001, 85.80000000000001, 101.2]
```

In this example, only scores greater than or equal to 50 are included, and each selected score is increased by 10%. List comprehension simplifies the process of filtering and transforming the data in a single readable line.

10.4 Exercises

10.4.0.1 Exercise 1: Basic List Operations

- a. Create a list named `fruits` with the values: “apple”, “banana”, “cherry”.
- b. Add the fruit “orange” to the list.
- c. Remove “banana” from the list.
- d. Print the second fruit in the list.
- e. Print the length of the list.

10.4.0.2 Exercise 2: Modifying Lists

- a. Create a list `numbers` containing the values 1, 2, 3, 4, and 5.
- b. Replace the third element in the list with the value 10.
- c. Add the number 6 to the end of the list.
- d. Insert the number 0 at the beginning of the list.
- e. Print the updated list.

10.4.0.3 Exercise 3: Slicing Lists

Given the list `colors = ["red", "green", "blue", "yellow", "purple"]`,

- a. Slice and print the first three colors.
- b. Slice and print the last two colors.
- c. Slice and print every second color in the list.

10.4.0.4 Excercise 4: Nested Lists

Create a nested list `matrix = [[1, 2, 3], [4, 5, 6], [7, 8, 9]]`.

- Print the element at the first row and second column.
- Change the element at the second row and third column to 10.
- Print the entire second row.

10.4.0.5 Excercise 5: Basic Tuple Operations

- Create a tuple `my_tuple` with values: 10, 20, 30, 40, 50.
- Print the value at index 2.
- Try to change the value at index 1 to 15 (what happens?).
- Print the length of the tuple.

10.4.0.6 Excercise 6: Tuple Unpacking

- Create a tuple `dimensions = (1920, 1080)`.
- Unpack the tuple into variables `width` and `height`.
- Print the values of `width` and `height`.

10.4.0.7 Excercise 7: Returning Tuples from Functions

- Write a function `calculate_stats(numbers)` that takes a list of numbers and returns a tuple containing the sum and the average of the list.
- Call the function with the list `[10, 20, 30, 40, 50]` and unpack the returned tuple into variables `total_sum` and `average`.
- Print the values of `total_sum` and `average`.

10.4.0.8 Excercise 8: Nested Tuples

Given the nested tuple `nested = ((1, 2), (3, 4), (5, 6))`,

- Print the first element of the second tuple.
- Try to change the second element of the third tuple to 7 (what happens?).

10.4.0.9 Excercise 9: Basic List Comprehension

- a. Create a list comprehension that generates a list of squares of numbers from 1 to 10.
- b. Create a list comprehension that generates a list of even numbers between 1 and 20.

10.4.0.10 Excercise 10: Basic List Comprehension

- a. Create a list comprehension that generates a list of all numbers between 1 and 50 that are divisible by 3.
- b. Create a list comprehension that generates a list of all numbers between 1 and 100 that are divisible by both 2 and 5.

10.4.0.11 Excercise 11: Nested List Comprehension

Using nested list comprehension, create a list of all possible combinations of two numbers, where the first number is from the list [1, 2, 3] and the second number is from the list [4, 5, 6].

10.4.0.12 Excercise 12: String Manipulation with List Comprehension

Given the list `words = ["apple", "banana", "cherry", "date"]`,

- a. Create a list comprehension that returns the lengths of each word in the list.
- b. Create a list comprehension that converts each word in the list to uppercase.

10.4.0.13 Excercise 13: Tuples and List Comprehension

Given a list of tuples representing students and their scores:

```
students = [("Alice", 85), ("Bob", 60), ("Charlie", 95), ("David", 70)],
```

- a. Use a list comprehension to create a list of names of students who scored 70 or above.
- b. Use a list comprehension to create a list of tuples where each student's score is increased by 5 points.

10.4.0.14 Excercise 14: Matrix Flattening

Given a nested list (matrix): `matrix = [[1, 2, 3], [4, 5, 6], [7, 8, 9]]`, use list comprehension to flatten the matrix into a single list.

11 Data Structures: Dictionaries and Sets

In the previous chapter, we explored lists and tuples, which are foundational data structures in Python. Now, we turn our attention to two additional data structures: dictionaries and sets. These structures provide efficient ways to store and manipulate data, particularly when managing large or complex datasets.

11.1 Dictionaries

A **dictionary** in Python is a versatile and powerful data structure that stores data in key-value pairs. Unlike lists, which use numerical indices, dictionaries use keys that can be of any immutable data type (e.g., strings, numbers, tuples). This makes dictionaries an ideal structure for tasks that require fast lookups, updates, and association between related pieces of data.

11.1.1 Creating a Dictionary

Dictionaries are created using curly braces `{}` or the `dict()` constructor, and key-value pairs are defined using the syntax `key: value`. Let's explore various ways to create dictionaries:

11.1.1.1 Literal Notation

The most straightforward way to create a dictionary is by using curly braces:

```
# Creating a dictionary with student grades
grades = {
    "John": 85,
    "Alice": 92,
    "Bob": 78
}
```

11.1.1.2 Using the dict() Constructor

Alternatively, dictionaries can be created using the `dict()` constructor, which allows for the creation of dictionaries using keyword arguments or iterables of key-value pairs:

```
# Creating a dictionary using keyword arguments
grades = dict(John=85, Alice=92, Bob=78)

# Creating a dictionary from a list of tuples
grades = dict([("John", 85), ("Alice", 92), ("Bob", 78)])
```

In this example, both methods create the same dictionary as the one using literal notation.

11.1.2 Accessing Values

To access values in a dictionary, you use the key as the index. This allows for quick lookup time, which is one of the key advantages of dictionaries over lists or tuples.

```
# Accessing the value associated with the key "Alice"
print(grades["Alice"])
```

92

If the key is not found, Python raises a `KeyError`. To avoid this, you can use the `get()` method, which returns `None` or a specified default value if the key does not exist:

```
# Safely accessing a key
print(grades.get("David", "Not Found"))
```

Not Found

Note the second argument for the `get` method is a value that will be returned if the specified key does not exist. The default is `None`.

11.1.3 Adding and Modifying Entries

Dictionaries are mutable, meaning you can add or modify key-value pairs after the dictionary has been created.

11.1.3.1 Modifying Values

To modify the value associated with an existing key, simply assign a new value to that key:

```
# Modifying the value associated with "John"
grades["John"] = 88
print(grades)
```

```
{'John': 88, 'Alice': 92, 'Bob': 78}
```

11.1.3.2 Adding New Key-Value Pairs

To add a new key-value pair, use the same syntax as modifying an existing pair:

```
# Adding a new key-value pair
grades["David"] = 90
print(grades)
```

```
{'John': 88, 'Alice': 92, 'Bob': 78, 'David': 90}
```

11.1.4 Deleting Key-Value Pairs

There are several ways to remove key-value pairs from a dictionary:

11.1.4.1 Using the del Statement

The `del` statement removes the key-value pair from the dictionary:

```
# Removing the key-value pair for "Bob"
del grades["Bob"]
print(grades)
```

```
{'John': 88, 'Alice': 92, 'David': 90}
```

11.1.4.2 Using the pop() Method

The `pop()` method removes a key-value pair and returns the value. If the key is not found, it raises a `KeyError`, unless a default value is provided.

```
# Removing and returning the value for "Alice"
alice_grade = grades.pop("Alice")
print(alice_grade)
print(grades)
```

92

```
{'John': 88, 'David': 90}
```

11.1.5 Dictionary Methods

Dictionaries come with a variety of built-in methods that simplify common tasks, such as adding, removing, and checking for keys and values.

11.1.5.1 keys(), values(), and items()

- `keys()` returns a view of all the keys in the dictionary.
- `values()` returns a view of all the values in the dictionary.
- `items()` returns a view of all key-value pairs as tuples.

```
print(grades.keys())
print(grades.values())
print(grades.items())
```

```
dict_keys(['John', 'David'])
dict_values([88, 90])
dict_items([('John', 88), ('David', 90)])
```

11.1.5.2 update()

The `update()` method allows you to merge two dictionaries or add key-value pairs from another iterable:

```
# Merging dictionaries
extra_grades = {"Eve": 85, "Charlie": 79}
grades.update(extra_grades)
print(grades)
```

```
{'John': 88, 'David': 90, 'Eve': 85, 'Charlie': 79}
```

11.1.5.3 clear()

The `clear()` method removes all key-value pairs from the dictionary:

```
grades.clear()
print(grades)
```

```
{}
```

11.1.6 Iterating Over Dictionaries

There are several ways to iterate over dictionaries, depending on whether you need keys, values, or both:

11.1.6.1 Iterating Over Keys

By default, iterating over a dictionary yields its keys:

```
for student in grades:
    print(student)
```

11.1.6.2 Iterating Over Values

You can iterate over the values by using the `values()` method:

```
for grade in grades.values():
    print(grade)
```

11.1.6.3 Iterating Over Key-Value Pairs

The `items()` method allows you to iterate over both keys and values simultaneously:

```
for student, grade in grades.items():
    print(f"{student}: {grade}")
```

11.1.7 Dictionary Comprehension

Like list comprehensions, Python also supports **dictionary comprehensions**, which provide a concise way to create dictionaries from iterables.

```
# Creating a dictionary of squares
squares = {x: x**2 for x in range(1, 6)}
print(squares)
```

```
{1: 1, 2: 4, 3: 9, 4: 16, 5: 25}
```

11.1.8 Practical Applications of Dictionaries

Dictionaries are used in a variety of real-world applications, particularly where fast lookups or associations between pieces of data are needed.

11.1.8.1 Frequency Count

One common use of dictionaries is counting the frequency of elements in a collection. Here is an example that counts the occurrence of each character in a string:

```
def char_frequency(text):
    freq = {}
    for char in text:
        if char in freq:
            freq[char] += 1
        else:
            freq[char] = 1
    return freq

text = "data science"
print(char_frequency(text))
```

```
{'d': 1, 'a': 2, 't': 1, ' ': 1, 's': 1, 'c': 2, 'i': 1, 'e': 2, 'n': 1}
```

11.1.8.2 Storing Configuration Settings

Dictionaries are often used to store configuration settings because they allow easy lookups by key:

```

config = {
    "host": "localhost",
    "port": 8080,
    "debug": True
}
print(config["host"])

```

localhost

11.1.8.3 Caching Computations

Dictionaries can be used to cache results of expensive computations to avoid recalculating them:

i cache

A cache (pronounced “cash”) is memory used to store something, usually data, temporarily in a computing environment.

```

factorial_cache = {}

def factorial(n):
    if n in factorial_cache:
        return factorial_cache[n]
    if n == 0:
        result = 1
    else:
        result = n * factorial(n-1)
    factorial_cache[n] = result
    return result

print(factorial(5))
print(factorial_cache)

```

120
{0: 1, 1: 1, 2: 2, 3: 6, 4: 24, 5: 120}

By caching the results, subsequent calls to `factorial(n)` for previously computed values are faster, as they avoid redundant calculations.

11.1.9 Using ****kwargs** in Functions

In Python, dictionaries are often used to pass and manage named arguments to functions. One powerful feature that leverages dictionaries is the ****kwargs** mechanism, which allows functions to accept an arbitrary number of keyword arguments (recall Chapter 8). These keyword arguments are collected into a dictionary, which provides flexibility when you do not know in advance what arguments might be passed to a function.

The ****kwargs** construct is particularly useful when writing functions that need to accept a variable number of named parameters or when extending existing functions with new optional arguments without changing their function signature.

11.1.9.1 How ****kwargs** Works

The term **kwargs** stands for “keyword arguments,” and when used with ******, it allows you to pass a variable number of named arguments to a function. Inside the function, these keyword arguments are captured as a dictionary.

```
def print_student_scores(**kwargs):
    for student, score in kwargs.items():
        print(f"{student}: {score}")

# Calling the function with multiple keyword arguments
print_student_scores(John=85, Alice=92, Bob=78)
```

```
John: 85
Alice: 92
Bob: 78
```

In the example above, the function `print_student_scores()` accepts any number of keyword arguments and prints them. The ****kwargs** parameter collects the keyword arguments as a dictionary, where the keys are the argument names (`John`, `Alice`, `Bob`), and the values are the respective scores.

11.1.9.2 Accessing and Using ****kwargs**

Once inside the function, ****kwargs** behaves like a normal dictionary. You can access, iterate over, and modify its elements just as you would with any other dictionary.

```
def get_student_grade(**kwargs):
    student = kwargs.get("student")
    grade = kwargs.get("grade")
    if student and grade:
        print(f"{student}'s grade is {grade}")
    else:
        print("Missing student or grade information")

# Providing student and grade as keyword arguments
get_student_grade(student="John", grade=85)

# Missing one argument
get_student_grade(student="Alice")
```

```
John's grade is 85
Missing student or grade information
```

In this example, the `kwargs.get()` method is used to safely retrieve values from the `kwargs` dictionary. If the key does not exist, `get()` returns `None`, which prevents the function from throwing a `KeyError`.

11.1.9.3 Combining `**kwargs` with Regular and Positional Arguments

You can combine `**kwargs` with regular and positional arguments. However, `**kwargs` must always be placed after regular arguments in the function signature:

```
def student_info(course, **kwargs):
    print(f"Course: {course}")
    for key, value in kwargs.items():
        print(f"{key}: {value}")

# Calling the function with both positional and keyword arguments
student_info("Mathematics", name="John", grade=90, age=20)
```

```
Course: Mathematics
name: John
grade: 90
age: 20
```

Here, `course` is a regular argument, and the remaining keyword arguments (e.g., `name`, `grade`, and `age`) are captured into the `kwargs` dictionary.

11.1.9.4 Passing a Dictionary as ****kwargs**

If you already have a dictionary of key-value pairs, you can pass it to a function using ****** to unpack the dictionary into keyword arguments.

```
def print_details(name, age, occupation):
    print(f"Name: {name}, Age: {age}, Occupation: {occupation}")

# Creating a dictionary of arguments
person = {"name": "Alice", "age": 30, "occupation": "Data Scientist"}

# Passing the dictionary as keyword arguments
print_details(**person)
```

Name: Alice, Age: 30, Occupation: Data Scientist

In this case, the ****person** syntax unpacks the dictionary into keyword arguments, allowing you to pass the dictionary directly into the function.

11.1.9.5 ****kwargs** and Default Arguments

While ****kwargs** allows for flexible keyword arguments, you can also combine it with default arguments to give function parameters some predefined behavior.

```
def log_message(level="INFO", **kwargs):
    message = kwargs.get("message", "No message provided")
    timestamp = kwargs.get("timestamp", "No timestamp")
    print(f"[{level}] {timestamp}: {message}")

# Logging a message with a default level
log_message(message="System started", timestamp="2023-09-12 10:00:00")

# Overriding the default level
log_message(level="ERROR", message="System failure", timestamp="2023-09-12 10:01:00")
```

```
[INFO] 2023-09-12 10:00:00: System started
[ERROR] 2023-09-12 10:01:00: System failure
```

In this example, the `log_message()` function uses a default level of “INFO” and then utilizes ****kwargs** to collect additional information like the message and timestamp.

11.1.9.6 Summary

The ****kwargs** construct provides a flexible way to pass and handle keyword arguments in Python. By collecting all keyword arguments into a dictionary, you gain the ability to write dynamic and adaptable functions. Whether used for configuration settings, logging, or passing optional parameters, ****kwargs** is a powerful tool that makes functions more reusable and extensible.

11.2 Sets

A **set** in Python is an unordered collection of unique elements. This data structure is useful when you need to store distinct items and perform operations such as union, intersection, difference, or membership testing efficiently. Sets are particularly powerful when handling large datasets where duplication is unnecessary or undesirable.

11.2.1 Creating a Set

Sets are created by placing items inside curly braces `{}` or by using the `set()` function. Unlike lists or dictionaries, sets do not maintain any order, and duplicate values are automatically removed. Sets can hold items of any immutable data type, such as numbers, strings, or tuples.

11.2.1.1 Literal Notation

You can create a set directly by enclosing a sequence of values in curly braces:

```
# Creating a set of integers
numbers = {1, 2, 3, 4, 5}
print(numbers)

# Creating a set of strings
fruits = {"apple", "banana", "cherry"}
print(fruits)
```

```
{1, 2, 3, 4, 5}
{'apple', 'banana', 'cherry'}
```

11.2.1.2 Using the set() Function

The `set()` function is particularly useful when creating a set from an iterable, such as a list or a string. It automatically removes duplicates:

```
# Creating a set from a list with duplicate values
numbers_list = [1, 2, 2, 3, 4, 4, 5]
unique_numbers = set(numbers_list)
print(unique_numbers)

# Creating a set from a string
letters = set("hello")
print(letters)
```

```
{1, 2, 3, 4, 5}
{'o', 'l', 'e', 'h'}
```

11.2.2 Adding and Removing Elements

11.2.2.1 Adding Elements

You can add elements to a set using the `add()` method. However, since sets do not allow duplicate values, adding an existing element has no effect:

```
# Adding elements to a set
fruits = {"apple", "banana"}
fruits.add("cherry")
print(fruits)

# Attempting to add a duplicate element
fruits.add("apple")
print(fruits)
```

```
{'cherry', 'banana', 'apple'}
{'cherry', 'banana', 'apple'}
```

11.2.2.2 Removing Elements

There are multiple ways to remove elements from a set, including the `remove()`, `discard()`, and `pop()` methods:

- `remove()` raises a `KeyError` if the element does not exist.
- `discard()` does not raise an error if the element is not found.
- `pop()` removes and returns an arbitrary element, as sets are unordered.

```
# Removing elements using remove()
fruits.remove("banana")
print(fruits)

# Using discard() to remove an element safely
fruits.discard("apple")
print(fruits)

# Removing a random element with pop()
random_fruit = fruits.pop()
print(random_fruit)
print(fruits)
```

```
{'cherry', 'apple'}
{'cherry'}
cherry
set()
```

11.2.3 Set Operations

One of the most powerful features of sets is their support for mathematical operations such as union, intersection, difference, and symmetric difference. These operations are efficient and allow for concise and readable code.

11.2.3.1 Union (`|` or `union()`)

The union operation combines all elements from two sets, excluding duplicates. This operation can be performed using the `|` operator or the `union()` method.

```
set1 = {1, 2, 3}
set2 = {3, 4, 5}

# Using the | operator
union_set = set1 | set2
print(union_set)
```

```
# Using the union() method
union_set = set1.union(set2)
print(union_set)
```

{1, 2, 3, 4, 5}

{1, 2, 3, 4, 5}

11.2.3.2 Intersection (& or intersection())

The intersection operation returns only the elements that are present in both sets. It can be performed using the & operator or the `intersection()` method.

```
set1 = {1, 2, 3}
set2 = {2, 3, 4}

# Using the & operator
intersection_set = set1 & set2
print(intersection_set)

# Using the intersection() method
intersection_set = set1.intersection(set2)
print(intersection_set)
```

{2, 3}

{2, 3}

11.2.3.3 Difference (- or difference())

The difference operation returns the elements that are in the first set but not in the second. This can be done using the - operator or the `difference()` method.

```
set1 = {1, 2, 3}
set2 = {2, 3, 4}

# Using the - operator
difference_set = set1 - set2
print(difference_set)

# Using the difference() method
```

```
difference_set = set1.difference(set2)
print(difference_set)
```

{1}

{1}

11.2.3.4 Symmetric Difference (^ or symmetric_difference())

The symmetric difference operation returns elements that are in either of the sets but not in both. This operation can be performed using the ^ operator or the `symmetric_difference()` method.

```
set1 = {1, 2, 3}
set2 = {2, 3, 4}

# Using the ^ operator
sym_diff_set = set1 ^ set2
print(sym_diff_set)

# Using the symmetric_difference() method
sym_diff_set = set1.symmetric_difference(set2)
print(sym_diff_set)
```

{1, 4}

{1, 4}

11.2.4 Checking for Subsets and Supersets

Sets also support operations that allow you to check whether one set is a subset or superset of another. These operations are particularly useful in scenarios where you need to compare sets.

11.2.4.1 Subset (<= or issubset())

A set A is a subset of set B if all elements of A are also in B. You can check for subsets using the <= operator or the `issubset()` method.


```

set1 = {1, 2}
set2 = {1, 2, 3, 4}

# Using the <= operator
print(set1 <= set2)

# Using the issubset() method
print(set1.issubset(set2))

```

True

True

11.2.4.2 Superset (>= or issuperset())

A set A is a superset of set B if all elements of B are also in A. You can check for supersets using the >= operator or the issuperset() method.

```

set1 = {1, 2, 3, 4}
set2 = {1, 2}

# Using the >= operator
print(set1 >= set2)

# Using the issuperset() method
print(set1.issuperset(set2))

```

True

True

11.2.5 Frozen Sets

A **frozen set** is an immutable version of a set, meaning that once a frozen set is created, its elements cannot be modified (i.e., you cannot add or remove elements). Frozen sets are useful when you need a collection of unique elements that should remain constant throughout the program. You create a frozen set using the `frozenset()` function:

```

# Creating a frozen set
immutable_set = frozenset([1, 2, 3, 4])
print(immutable_set)

```

```
# Attempting to add an element to a frozen set will raise an error
# immutable_set.add(5) # Raises AttributeError
```

```
frozenset({1, 2, 3, 4})
```

11.2.6 Set Comprehensions

Python provides a concise way to create sets using **set comprehensions**, similar to list comprehensions. A set comprehension is written with curly braces `{}` and allows you to define sets based on existing iterables, often including a filtering condition.

```
# Creating a set of squares for numbers 1 to 5
squares = {x**2 for x in range(1, 6)}
print(squares)

# Set comprehension with a condition
even_squares = {x**2 for x in range(1, 11) if x % 2 == 0}
print(even_squares)
```

```
{1, 4, 9, 16, 25}
{64, 100, 4, 36, 16}
```

Set comprehensions are a powerful tool when you need to transform or filter data while maintaining unique elements.

11.2.7 Set Best Practices

While sets are efficient for handling unique elements, there are a few best practices to keep in mind when working with sets in Python:

1. **Avoid Unnecessary Duplicates:** Since sets automatically remove duplicates, there's no need to check for duplicates before adding elements.
2. **Use Sets for Membership Testing:** When you need to check if an item exists in a collection and the collection does not need to maintain order or allow duplicates, sets are the best choice due to their $O(1)$ membership testing time complexity.
3. **Choose the Right Operation:** Use set operations such as union, intersection, and difference to simplify complex data comparison tasks. These operations are more efficient than writing custom loops to achieve the same results.

Sets are an invaluable data structure for handling collections of unique items. Their efficiency in membership testing, combined with their ability to perform set operations such as union, intersection, and difference, makes them ideal for a wide variety of tasks, from data processing to mathematical computations. With their unordered nature and automatic deduplication, sets help simplify code and ensure efficient performance, especially when working with large datasets.

11.2.8 Combining Dictionaries and Sets

You can often combine dictionaries and sets in practical applications. For example, to find unique words and their counts from a list of sentences:

```
def unique_words(sentences):
    word_dict = {}
    for sentence in sentences:
        words = set(sentence.split()) # Use a set to find unique words
        for word in words:
            if word in word_dict:
                word_dict[word] += 1
            else:
                word_dict[word] = 1
    return word_dict

sentences = ["data science is great", "data science is evolving"]
result = unique_words(sentences)
print(result) # Output: {'data': 2, 'science': 2, 'is': 2, 'great': 1, 'evolving': 1}
```

```
{'data': 2, 'is': 2, 'science': 2, 'great': 1, 'evolving': 1}
```

This function uses sets to ensure that each word in a sentence is only counted once per sentence and then stores the results in a dictionary.

11.3 Exercises

11.3.0.1 Exercise 1: Student Grades

Write a function called `update_grades()` that accepts a dictionary of student names and their grades. The function should accept new student names and grades and update the dictionary. Finally, it should return the updated dictionary.

Example:

```
grades = {"John": 85, "Alice": 92}
new_grades = {"Bob": 78, "Alice": 95}

updated_grades = update_grades(grades, **new_grades)
print(updated_grades)
```

Expected Output:

```
{"John": 85, "Alice": 95, "Bob": 78}
```

11.3.0.2 Excercise 2: Word Frequency Counter

Write a function `word_frequency(text)` that takes a string and returns a dictionary with the frequency count of each word in the string.

Example:

```
text = "data science is data fun science is fun"
result = word_frequency(text)
print(result)
```

Expected Output:

```
{'data': 2, 'science': 2, 'is': 2, 'fun': 2}
```

11.3.0.3 Excercise 3: Dictionary of Squares

Create a function `squares_dict(n)` that generates a dictionary where the keys are numbers from 1 to `n` and the values are their corresponding squares.

Example:

```
print(squares_dict(5))
```

Expected Output:

```
{1: 1, 2: 4, 3: 9, 4: 16, 5: 25}
```

11.3.0.4 Exercise 4: Merge Dictionaries

Write a function `merge_dictionaries(*args)` that accepts any number of dictionaries and merges them into a single dictionary. If a key is repeated, the value from the last dictionary should be retained.

Example:

```
dict1 = {"a": 1, "b": 2}
dict2 = {"b": 3, "c": 4}
dict3 = {"d": 5}

result = merge_dictionaries(dict1, dict2, dict3)
print(result)
```

Expected Output:

```
{'a': 1, 'b': 3, 'c': 4, 'd': 5}
```

11.3.0.5 Exercise 5: Remove Duplicates

Write a function `remove_duplicates(lst)` that takes a list and returns a new list with all the duplicates removed using a set.

Example:

```
numbers = [1, 2, 2, 3, 4, 4, 5]
print(remove_duplicates(numbers))
```

Expected Output:

```
[1, 2, 3, 4, 5]
```

11.3.0.6 Exercise 6: Set Operations

Given two sets of students enrolled in two different courses, write functions to:

1. Find students enrolled in both courses.
2. Find students enrolled only in the first course.
3. Find students enrolled in either course but not both.

Example:

```
course_A = {"Alice", "Bob", "Charlie", "David"}
course_B = {"Charlie", "David", "Eve", "Frank"}

# Students in both courses
print(students_in_both(course_A, course_B))

# Students only in course A
print(only_in_first(course_A, course_B))

# Students in either course but not both
print(either_but_not_both(course_A, course_B))
```

Expected Output:

```
{'Charlie', 'David'}
{'Alice', 'Bob'}
{'Alice', 'Bob', 'Eve', 'Frank'}
```

11.3.0.7 Exercise 7: Flexible Function with **kwargs

Write a function `student_profile(**kwargs)` that accepts student information (name, age, grade, etc.) as keyword arguments and prints the information in a readable format.

Example:

```
student_profile(name="Alice", age=20, grade="A", major="Mathematics")
```

Expected Output:

```
Name: Alice
Age: 20
Grade: A
Major: Mathematics
```

11.3.0.8 Exercise 8: Summing Keyword Arguments

Write a function `sum_values(**kwargs)` that takes any number of keyword arguments where the values are integers and returns their sum.

Example:

```
result = sum_values(a=5, b=10, c=3)
print(result)
```

Expected Output:

```
18
```

12 JSON Files in Python

JSON (JavaScript Object Notation) is a lightweight format used for storing and exchanging data. It is often used to transmit data between a server and a web application, serving as a common data-interchange format. In Python, working with JSON is straightforward thanks to the built-in `json` module, which provides functionality for parsing, serializing, and deserializing JSON data. This chapter introduces JSON, how to handle it using standard libraries, and how to create and manage custom modules for JSON processing.

12.1 Introduction to JSON

JSON (JavaScript Object Notation) is a widely-used format for data interchange, particularly in web development and API communication. JSON is designed to be both human-readable and machine-readable, making it an ideal choice for data exchange across platforms and programming languages. JSON represents structured data as a series of key-value pairs, much like Python dictionaries, but with a more constrained and universal syntax. It can easily handle data types such as objects, arrays, strings, numbers, booleans, and null values.

12.1.1 Why Use JSON?

JSON (JavaScript Object Notation) has become the dominant format for data interchange due to its simplicity, flexibility, and widespread support. Here are several key reasons why JSON is preferred in many applications:

1. **Lightweight and Efficient:** JSON is a minimalistic format that uses a concise structure to represent data. Unlike XML, JSON eliminates the need for heavy markup tags, making it more compact and faster to process. This lightweight nature is especially beneficial in network communication, where reducing data size can significantly improve performance.
2. **Cross-Platform Compatibility:** Although JSON originated from JavaScript, it is now a language-independent standard. Most modern programming languages, including Python, Java, C++, and Ruby, offer built-in support for parsing and generating JSON. This makes JSON ideal for systems where data needs to be transferred between different technologies.

3. **Human-Readable:** JSON's clear and straightforward syntax makes it easy for humans to read and write. The structure, based on key-value pairs and arrays, is intuitive and similar to Python's dictionaries and lists, which helps developers quickly understand the data.
4. **Common in Web Development:** JSON is the default data format for most web APIs. RESTful services, in particular, rely heavily on JSON to structure the data exchanged between clients and servers. Its popularity in web applications makes it a critical skill for developers working with modern web technologies.
5. **Easy to Parse:** JSON is simple to parse in most programming environments. Libraries like Python's `json` module provide straightforward methods for converting JSON data into native data structures and vice versa, making JSON a practical choice for data interchange.

Overall, JSON is widely used because it strikes an effective balance between being machine-friendly and human-friendly, making it an optimal choice for a variety of applications.

12.1.2 JSON Structure and Data Types

JSON's structure is based on two universal data structures:

- **Objects:** Collections of key-value pairs, where keys are strings and values can be any valid JSON type.
- **Arrays:** Ordered lists of values, where each value can be any valid JSON type.

In addition, JSON supports the following primitive types:

- **String:** A sequence of characters, enclosed in double quotes.
- **Number:** Integers or floating-point numbers.
- **Boolean:** `true` or `false`.
- **Null:** A special value representing the absence of data (`null` in JSON, `None` in Python).

12.1.2.1 Example JSON Object

A typical JSON object that describes a student might look like this:

```
{  
  "name": "Alice",  
  "age": 21,  
  "major": "Statistics",  
  "graduated": false,  
  "courses": ["Calculus", "Linear Algebra", "Statistics"],  
}
```

```
    "details": {
      "GPA": 3.8,
      "credits_completed": 95
    }
  }
```

Here, the JSON object contains several fields:

- **name:** A string representing the student's name.
- **age:** A number representing the student's age.
- **major:** A string representing the student's field of study.
- **graduated:** A boolean indicating whether the student has graduated.
- **courses:** An array of strings representing the courses the student has taken.
- **details:** A nested object with more specific information about the student's GPA and credits completed.

12.1.2.2 JSON Arrays

In JSON, arrays are used to store lists of data. An array can contain any type of value: numbers, strings, booleans, objects, or even other arrays. This makes JSON very flexible for representing complex data structures.

Example of an array of student objects:

```
[
  {
    "name": "Alice",
    "age": 21,
    "major": "Statistics"
  },
  {
    "name": "Bob",
    "age": 23,
    "major": "Mathematics"
  },
  {
    "name": "Charlie",
    "age": 22,
    "major": "Computer Science"
  }
]
```

This JSON array contains three objects, each representing a student with their name, age, and major.

12.1.3 Differences Between JSON and Python Data Types

Although JSON and Python share many similarities, there are important differences to keep in mind when converting between the two:

- **Python dictionaries** (`dict`) correspond to **JSON objects**.
- **Python lists** (`list`) correspond to **JSON arrays**.
- **Python strings** (`str`) map directly to **JSON strings**.
- **Python integers** and **floats** map to **JSON numbers**.
- **Python None** is equivalent to **JSON null**.
- **Python True** and **False** are equivalent to **JSON true** and **false**, respectively.

These mappings allow for seamless conversions between Python data structures and JSON, but developers must be aware of slight differences in how Python and JSON handle certain data. For example, in JSON, only double quotes (") are allowed for strings, while Python allows both single and double quotes.

Example: Converting Python Data to JSON

Suppose you have the following Python dictionary:

```
student_data = {
    "name": "Alice",
    "age": 21,
    "major": "Statistics",
    "graduated": False,
    "courses": ["Calculus", "Linear Algebra", "Statistics"],
    "details": {
        "GPA": 3.8,
        "credits_completed": 95
    }
}
```

This Python dictionary can be converted to JSON using the `json` module:

```
import json

json_data = json.dumps(student_data)
print(json_data)
```

```
{"name": "Alice", "age": 21, "major": "Statistics", "graduated": false, "courses": ["Calculus"]}
```

Note the subtle differences, such as the use of lowercase `false` instead of Python's `False`, and the use of double quotes around strings.

12.1.4 Use Cases for JSON

JSON is widely used in various applications, some of the most common being:

- **Web APIs:** JSON is the standard format for exchanging data between client-side applications (e.g., web browsers) and server-side applications.
- **Configuration Files:** JSON is often used for configuration files in modern software applications because it is easy to read and write.
- **Data Serialization:** JSON is commonly used to serialize and deserialize data in a format that can be easily exchanged across different programming languages and platforms.
- **Data Storage:** JSON can be used as a lightweight alternative to databases for small-scale data storage, particularly for configuration settings or user preferences.

By understanding the structure of JSON and how it relates to Python's data types, we can efficiently use it to handle data in real-world scenarios, particularly when working with web applications, APIs, or data storage systems.

12.2 Reading and Writing JSON Data

Serialization is the process of converting an object or data structure into a format that can be easily stored or transmitted and then reconstructed later. This process allows data to be saved to a file, sent over a network, or stored in a database, and later deserialized (reconstructed) back into its original form. In Python, serialization often refers to converting Python objects into formats like JSON, XML, or binary formats.

For example, when you serialize a Python dictionary into a JSON string, you are converting the dictionary into a format that can be written to a file or transmitted over a network. The reverse process—converting a serialized format back into a Python object—is called **deserialization**.

One of the key benefits of JSON in Python is the ease with which it can be read from and written to files using the built-in `json` module. This section explores the core methods provided by this module, including reading (parsing) JSON data from files, writing (serializing) Python objects into JSON, and handling JSON data as strings. Understanding these operations is essential for working with APIs, configurations, or any structured data exchange.

12.2.1 Loading JSON Data from a File

To read (or deserialize) JSON data from a file, the `json.load()` method is used. This method reads the entire content of a file and converts it into a Python object (such as a dictionary or list). Here's a simple example:

12.2.1.1 Example: Reading from a JSON File

Assume you have a file `student.json` that contains the following JSON data:

```
{
    "name": "Alice",
    "age": 21,
    "major": "Statistics",
    "graduated": false
}
```

You can load this data into a Python dictionary using the `json.load()` method as follows:

```
import json

# Open the JSON file for reading
with open("student.json", "r") as file:
    student_data = json.load(file)

# Accessing the data
print(student_data["name"]) # Output: Alice
```

Alice

In this example:

- We open the `student.json` file in read mode.
- The `json.load()` function parses the JSON data and converts it into a Python dictionary.
- You can then access the values in the dictionary as you would with any Python dictionary.

12.2.2 Error Handling While Loading JSON Data

When reading JSON data from a file, errors can occur if the file is improperly formatted or does not exist. Python's `json` module raises a `json.JSONDecodeError` if the content is not valid JSON, and a `FileNotFoundError` if the file is missing. To handle these potential errors, you can use `try-except` blocks.

```
import json

# Safely loading JSON data from a file
try:
    with open("student.json", "r") as file:
        student_data = json.load(file)
except FileNotFoundError:
    print("Error: The file was not found.")
except json.JSONDecodeError:
    print("Error: The file contains invalid JSON.")
```

This ensures that the program gracefully handles common file and parsing errors instead of crashing unexpectedly.

12.2.3 Writing JSON Data to a File

Writing (or serializing) Python objects into JSON format is done using the `json.dump()` method as shown in the previous section. This method takes a Python object and writes it to a file in JSON format.

12.2.3.1 Example: Writing to a JSON File

Let's say we want to save a dictionary representing a student's data to a JSON file:

```
import json

# Python dictionary
student_data = {
    "name": "Bob",
    "age": 23,
    "major": "Mathematics",
    "graduated": True
}
```

```
# Writing to a JSON file
with open("student.json", "w") as file:
    json.dump(student_data, file, indent=4)
```

In this example:

- We open a file `student.json` in write mode.
- The `json.dump()` function writes the `student_data` dictionary to the file in JSON format.
- The `indent=4` argument is used to format the output with indentation, making the JSON more readable.

The resulting `student.json` file will look like this:

```
{
  "name": "Bob",
  "age": 23,
  "major": "Mathematics",
  "graduated": true
}
```

12.2.4 Error Handling When Writing JSON Data

Just like reading JSON files, writing to JSON files can also result in errors, such as `IOError` if the file cannot be opened or written to. To handle these cases, wrap the `json.dump()` operation in a `try-except` block.

```
import json

# Safely writing JSON data to a file
try:
    with open("student.json", "w") as file:
        json.dump(student_data, file, indent=4)
except IOError as e:
    print(f"Error writing to file: {e}")
```

This approach ensures that your program responds appropriately if file operations fail.

12.2.5 Loading JSON Data from a String

In some cases, JSON data might not come from a file but from a string, such as when receiving data from a web API. The `json.loads()` function is used to parse JSON data from a string and convert it into a Python object.

12.2.5.1 Example: Parsing JSON from a String

Here's how you can parse a JSON string into a Python dictionary:

```
import json

# JSON string
json_string = '{"name": "Charlie", "age": 22, "major": "Computer Science"}'

# Parsing the JSON string
student_data = json.loads(json_string)

print(student_data) # Output: {'name': 'Charlie', 'age': 22, 'major': 'Computer Science'}
```

```
{'name': 'Charlie', 'age': 22, 'major': 'Computer Science'}
```

The `json.loads()` method converts the JSON string into a Python dictionary, which can be used like any other dictionary.

12.2.6 Converting Python Objects to JSON Strings

In addition to writing JSON to files, you might need to generate JSON-formatted strings for data exchange, such as sending data over a network or printing it to the console. The `json.dumps()` function allows you to convert Python objects to JSON strings.

12.2.6.1 Example: Converting Python Dictionary to JSON String

```
import json

# Python dictionary
student_data = {
    "name": "Diana",
    "age": 20,
```



```

        "major": "Engineering"
    }

    # Converting to JSON string
    json_string = json.dumps(student_data, indent=4)
    print(json_string)

{
    "name": "Diana",
    "age": 20,
    "major": "Engineering"
}

```

The `indent` parameter is optional, but it improves readability by formatting the JSON with proper indentation.

12.2.7 Customizing JSON Serialization

Sometimes, Python objects may contain data types that are not directly serializable by the `json` module, such as datetime objects. In these cases, you can provide a custom function to handle the serialization of these complex types.

12.2.7.1 Example: Custom Serialization

```

import json
from datetime import datetime

# Python dictionary with a datetime object
student_data = {
    "name": "Emily",
    "graduation_date": datetime(2023, 5, 15)
}

# Custom serialization function
def custom_serializer(obj):
    if isinstance(obj, datetime):
        return obj.strftime('%Y-%m-%d')
    raise TypeError(f"Type {type(obj)} is not serializable")

```

```
# Converting to JSON string with custom serialization
json_string = json.dumps(student_data, default=custom_serializer, indent=4)
print(json_string)
```

```
{
  "name": "Emily",
  "graduation_date": "2023-05-15"
}
```

In this example, the `default` parameter is used to specify a custom serialization function for handling the `datetime` object. The resulting JSON string will look like this:

Without the custom function, attempting to serialize a `datetime` object would raise a `TypeError`.

12.3 Exercises

12.3.0.1 Exercise 1: Loading JSON from a File

You have a file named `book.json` that contains the following JSON data:

```
{
  "title": "Python Programming",
  "author": "John Doe",
  "year": 2020,
  "genres": ["Programming", "Technology"],
  "available": true
}
```

A. Write a Python program to read the contents of the file `book.json` and print the title of the book. B. Extend your program to print the author and the list of genres as well.

12.3.0.2 Exercise 2: Writing JSON to a File

Create a Python dictionary that represents the following student data:

- Name: “Sarah”
- Age: 24
- Major: “Data Science”
- Courses: [“Machine Learning”, “Statistics”, “Python Programming”]

- Graduated: False

A. Write a Python script that saves this dictionary to a file named `student_data.json` in JSON format with indentation. B. Open the file and verify that the content is properly formatted as JSON.

12.3.0.3 Exercise 3: Parsing JSON from a String

You receive the following JSON string from an API response:

```
{
  "city": "Austin",
  "temperature": 30,
  "conditions": "Sunny",
  "forecast": ["Sunny", "Partly Cloudy", "Rain"]
}
```

A. Write a Python program to parse this JSON string and convert it into a Python dictionary. B. Print the current weather condition ("conditions") and the second item in the forecast list.

12.3.0.4 Exercise 4: Serializing Python Data to JSON

You are given the following Python dictionary:

```
employee_data = {
    "name": "Alice",
    "id": 12345,
    "position": "Software Engineer",
    "start_date": "2021-09-01",
    "salary": 85000,
    "active": True
}
```

A. Write a Python script to convert this dictionary into a JSON-formatted string. B. Ensure that the resulting JSON string is printed with an indentation of 4 spaces for better readability. C. Save this JSON string to a file called `employee.json`.

12.3.0.5 Exercise 5: Handling Errors in JSON Files

You are working with JSON files, and sometimes they may not be formatted correctly or may be missing. Write a Python program that:

A. Attempts to load JSON data from a file named `config.json`. B. If the file is missing or contains invalid JSON, catch and handle the exceptions appropriately, printing an error message like:

- "Error: `config.json` not found." for missing files.
- "Error: Invalid JSON format." for JSON decoding errors.

12.3.0.6 Exercise 6: Custom Serialization of Python Objects

Consider a Python dictionary that contains a `datetime` object:

```
from datetime import datetime

event = {
    "name": "Conference",
    "location": "New York",
    "date": datetime(2024, 5, 15, 10, 30)
}
```

A. Write a Python program to serialize this dictionary into a JSON string. Use a custom serialization function to convert the `datetime` object into a string formatted as YYYY-MM-DD HH:MM. B. Save the resulting JSON string to a file named `event.json`.

12.3.0.7 Exercise 7: Converting a List of Dictionaries to JSON

You have the following list of dictionaries, each representing a book in a library:

```
books = [
    {"title": "Python Basics", "author": "Alice", "year": 2019},
    {"title": "Data Science Handbook", "author": "Bob", "year": 2021},
    {"title": "Machine Learning 101", "author": "Charlie", "year": 2020}
]
```

A. Write a Python script to convert this list into a JSON-formatted string. B. Save the JSON data to a file called `books.json`.

12.3.0.8 Exercise 8: Modifying and Writing JSON Data

You are given a file named `users.json` with the following data:

```
[
  {"username": "john_doe", "email": "john@example.com", "active": true},
  {"username": "jane_doe", "email": "jane@example.com", "active": false}
]
```

A. Write a Python script that loads this data into a Python list. B. Modify the script to activate all users by setting the "active" field to `true` for all entries. C. Save the modified data back to `users.json`.

12.3.0.9 Exercise 9: Nested JSON Parsing

You are given a JSON string representing nested product data:

```
{
  "product": {
    "id": 101,
    "name": "Laptop",
    "price": 1200,
    "specifications": {
      "processor": "Intel i7",
      "ram": "16GB",
      "storage": "512GB SSD"
    }
  }
}
```

A. Write a Python program to parse this JSON string into a Python dictionary. B. Extract and print the product name, price, and the processor specification from the nested dictionary.

13 NumPy

In computational mathematics and statistics, arrays are fundamental structures for storing and manipulating large datasets. The Python library NumPy (Numerical Python) offers an efficient and powerful way to handle multi-dimensional arrays and matrices, enabling fast numerical computations. This chapter will introduce the concept of arrays using NumPy and explore basic operations that can be performed on these arrays.

13.1 Introduction to NumPy

In computational tasks, particularly in statistics and mathematics, efficiency and speed are crucial when working with large datasets or performing complex numerical computations. The Python library **NumPy** (short for **Numerical Python**) addresses these needs by providing support for large, multi-dimensional arrays and matrices, along with a collection of high-level mathematical functions to operate on these arrays.

13.1.1 Why Use NumPy?

While Python's built-in data structures like lists are flexible and easy to use, they are not optimized for numerical computation. NumPy offers a significant performance boost due to its underlying implementation in C, allowing for more efficient memory use and faster execution of mathematical operations. This makes NumPy a cornerstone for scientific computing, as it can handle large-scale numerical data with a much higher efficiency than native Python structures.

Here are several key reasons why NumPy is indispensable in computational mathematics and statistics:

1. **Performance:** NumPy arrays are faster than Python lists, especially when performing repeated or large-scale computations. This speed advantage arises from the fact that NumPy arrays are implemented in C, and operations on arrays are performed in compiled code, reducing the overhead associated with Python's dynamic typing and interpreted nature.

2. **Convenience:** NumPy provides a wide range of functions that allow users to manipulate arrays in various ways, such as reshaping, indexing, slicing, and performing mathematical operations, with far less code than is required with Python's built-in data structures.
3. **Mathematical Functionality:** NumPy includes built-in functions for linear algebra, random number generation, statistical operations, and much more, all optimized for efficiency. These functions can be applied element-wise to arrays, which allows for vectorized operations. This contrasts with Python lists, where applying mathematical operations requires explicit looping, which is slower and less intuitive.
4. **Interoperability:** NumPy is widely used in conjunction with other scientific libraries such as **SciPy** (for additional mathematical functions), **Matplotlib** (for plotting), and **TensorFlow** (for machine learning). Its use is ubiquitous in data science, making it a key foundation for many other libraries.
5. **Memory Efficiency:** Arrays in NumPy are homogeneous, meaning that all elements in an array are of the same type. This contrasts with Python lists, which can contain elements of different types. The homogeneity of arrays allows NumPy to allocate memory more efficiently and perform operations much faster, particularly for large datasets.

13.1.2 Installing NumPy

Before you can use NumPy, you need to install it. If you are using Google Colab, Jupyter notebooks, or a typical Python development environment, NumPy is often pre-installed. If not, you can install it using the following command in your terminal or command prompt:

```
pip install numpy
```

After installation, you can import the library in your Python script:

```
import numpy as np
```

Using the alias `np` is a widely adopted convention in the Python community and keeps your code clean and concise when calling NumPy functions.

13.2 Arrays in NumPy

Arrays are the fundamental building blocks in NumPy, providing a way to organize and manipulate large amounts of data efficiently. NumPy arrays, or **ndarrays** (short for N-dimensional arrays), can store data in multiple dimensions, allowing for efficient data processing in a format that's highly optimized for both performance and memory usage. This section explores the creation, manipulation, and essential operations with NumPy arrays.

13.2.1 Creating Arrays

Creating arrays in NumPy can be done in several ways depending on the need. Arrays can be initialized from lists, created using predefined shapes like zeros or ones, or generated using sequences of values.

13.2.1.1 Creating Arrays from Python Lists

One of the simplest ways to create an array is by converting a Python list into a NumPy array using `np.array()`.

```
import numpy as np

# Creating a 1D array from a list
arr = np.array([1, 2, 3, 4, 5])
print(arr)
```

```
[1 2 3 4 5]
```

This creates a one-dimensional array. Similarly, you can create multi-dimensional arrays by passing in lists of lists.

```
# Creating a 2D array (a matrix) from nested lists
arr_2d = np.array([[1, 2, 3], [4, 5, 6]])
print(arr_2d)
```

```
[[1 2 3]
 [4 5 6]]
```

In this case, `arr_2d` is a 2x3 matrix with two rows and three columns.

13.2.1.2 Creating Arrays with NumPy Functions

NumPy also provides a variety of built-in functions to generate arrays with specific patterns or values.

- `np.zeros()`: Creates an array filled with zeros.


```
# 3x4 array filled with zeros
zeros_array = np.zeros((3, 4))
print(zeros_array)
```

```
[[0. 0. 0. 0.]
 [0. 0. 0. 0.]
 [0. 0. 0. 0.]]
```

- **np.ones()**: Creates an array filled with ones.

```
# 2x5 array filled with ones
ones_array = np.ones((2, 5))
print(ones_array)
```

```
[[1. 1. 1. 1. 1.]
 [1. 1. 1. 1. 1.]]
```

- **np.full()**: Creates an array filled with a specified value.

```
# 3x3 array filled with the value 7
full_array = np.full((3, 3), 7)
print(full_array)
```

```
[[7 7 7]
 [7 7 7]
 [7 7 7]]
```

- **np.arange()**: Creates an array with a sequence of values, similar to Python's `range()` function, but returns an array.

```
# Array with values [0, 2, 4, 6, 8]
range_array = np.arange(0, 10, 2)
print(range_array)
```

```
[0 2 4 6 8]
```

- **np.linspace()**: Creates an array with a specified number of evenly spaced values between a start and end point.

```
# 5 values between 0 and 1, inclusive
linspace_array = np.linspace(0, 1, 5)
print(linspace_array)
```

```
[0.    0.25 0.5  0.75 1.   ]
```

These functions are incredibly useful when you need to quickly generate arrays for mathematical computations, simulations, or testing algorithms.

13.2.2 Array Indexing and Slicing

NumPy arrays allow you to access and manipulate specific elements, subarrays, or slices of the array efficiently using indexing and slicing techniques. This functionality mirrors Python's list slicing but extends it to multiple dimensions.

13.2.2.1 Indexing in 1D Arrays

In a one-dimensional array, you can access elements by specifying the index of the element you want:

```
arr = np.array([10, 20, 30, 40, 50])
print(arr[0])
print(arr[-1])
```

```
10
```

```
50
```

Just like Python lists, NumPy arrays support negative indexing, where `-1` refers to the last element, `-2` to the second-to-last, and so on.

13.2.2.2 Indexing in Multi-Dimensional Arrays

In multi-dimensional arrays, you access elements by specifying indices for each dimension:

```
arr_2d = np.array([[1, 2, 3], [4, 5, 6], [7, 8, 9]])
print(arr_2d[1, 2])
```

```
6
```

To access entire rows or columns, you can use slicing. For example:

```
# Access the second row
print(arr_2d[1, :])

# Access the third column
print(arr_2d[:, 2])
```

```
[4 5 6]
[3 6 9]
```

13.2.2.3 Slicing Arrays

Slicing allows you to access subarrays or parts of an array. The syntax is similar to Python lists: `start:stop:step`, where:

- `start` is the starting index (inclusive),
- `stop` is the stopping index (exclusive),
- `step` is the interval between indices.

```
arr = np.array([10, 20, 30, 40, 50, 60])

# Slice from index 1 to 4 (not including 4)
print(arr[1:4])

# Slice with a step of 2
print(arr[::2])
```

```
[20 30 40]
[10 30 50]
```

Slicing works similarly with multi-dimensional arrays, allowing you to select rows and columns:

```
arr_2d = np.array([[1, 2, 3], [4, 5, 6], [7, 8, 9]])

# Slice rows 0 to 2, columns 1 to 3
print(arr_2d[0:2, 1:3])
```

```
[[2 3]
 [5 6]]
```

13.3 Array Operations

One of the most powerful features of NumPy arrays is the ability to perform element-wise operations without needing loops. This feature is called **vectorization**, and it makes mathematical operations on arrays much more efficient.

13.3.1 Element-Wise Operations

You can perform arithmetic operations on entire arrays, and the operation is applied element-wise. This means each element in one array is paired with the corresponding element in another array (if the shapes match) to compute the result.

```
arr1 = np.array([1, 2, 3])
arr2 = np.array([4, 5, 6])

# Element-wise addition
print(arr1 + arr2)

# Element-wise multiplication
print(arr1 * arr2)

# Element-wise subtraction and division
print(arr2 - arr1)
print(arr2 / arr1)
```

```
[5 7 9]
[ 4 10 18]
[3 3 3]
[4.  2.5 2. ]
```

13.3.2 Scalar Operations

NumPy arrays also support operations with scalars, where the scalar is applied to every element in the array:

```
arr = np.array([10, 20, 30])

# Multiply each element by 2
print(arr * 2)
```

```
# Add 5 to each element
print(arr + 5)
```

```
[20 40 60]
[15 25 35]
```

13.3.3 Broadcasting

When performing operations between arrays of different shapes, NumPy automatically expands the smaller array to match the larger one. This process is called **broadcasting**.

For example:

```
arr = np.array([1, 2, 3])
scalar = 5

# Broadcasting allows the scalar to be added to each element of the array
print(arr + scalar)
```

```
[6 7 8]
```

Broadcasting is particularly useful when performing operations between arrays and scalars or between arrays with compatible shapes.

13.3.4 Array Shape and Reshaping

The shape of a NumPy array refers to the number of elements along each dimension. You can check the shape of an array using the `.shape` attribute:

```
arr_2d = np.array([[1, 2, 3], [4, 5, 6]])
print(arr_2d.shape)
```

```
(2, 3)
```

If necessary, you can reshape arrays using the `reshape()` function to change their dimensions while keeping the total number of elements constant:

```
arr = np.array([1, 2, 3, 4, 5, 6])
reshaped_arr = arr.reshape((2, 3))
print(reshaped_arr)
```

```
[[1 2 3]
 [4 5 6]]
```

The above code reshapes a 1D array into a 2D array with two rows and three columns.

13.4 Mathematical Functions in NumPy

NumPy comes with a wide array of built-in mathematical functions that operate efficiently on arrays. These include functions for statistical calculations, linear algebra, and more.

13.4.1 Common Mathematical Functions

- **Sum:** Computes the sum of all elements in the array.

```
arr = np.array([1, 2, 3, 4, 5])
print(np.sum(arr))
```

```
15
```

- **Mean:** Computes the mean (average) of the array elements.

```
print(np.mean(arr))
```

```
3.0
```

- **Standard Deviation:** Computes the standard deviation of the array elements.

```
print(np.std(arr))
```

```
1.4142135623730951
```

- **Min/Max:** Finds the minimum and maximum values in the array.

```
print(np.min(arr))
print(np.max(arr))
```

1
5

These functions are crucial when working with large datasets, as they allow for quick and efficient analysis of the data.

13.4.2 Array Comparisons and Conditional Operations

NumPy arrays allow for efficient comparison operations, resulting in boolean arrays. This is particularly useful in applications such as filtering data or applying conditions.

```
arr = np.array([10, 20, 30, 40, 50])

# Compare each element to 30
comparison = arr > 30
print(comparison)

# Use the comparison to filter the array
filtered_arr = arr[comparison]
print(filtered_arr)
```

```
[False False False  True  True]
[40 50]
```

13.5 NumPy Data Types

One of the strengths of NumPy is its ability to handle a wide range of data types. These types are much more memory-efficient compared to Python's built-in data types, particularly when dealing with large datasets.

Some common NumPy data types include:

- **int**: Signed integers (e.g., `np.int8`, `np.int16`, `np.int32`, `np.int64`), where the number indicates the bit-length.
- **float**: Floating-point numbers (e.g., `np.float16`, `np.float32`, `np.float64`).
- **bool**: Boolean type, where values can be `True` or `False`.
- **complex**: Complex numbers (e.g., `np.complex64`, `np.complex128`).

You can specify the data type of an array when creating it, or NumPy will infer the type based on the data you provide:

```
# Creating an array of integers
arr_int = np.array([1, 2, 3], dtype=np.int32)

# Creating an array of floats
arr_float = np.array([1.0, 2.0, 3.0], dtype=np.float64)

print(arr_int.dtype)
print(arr_float.dtype)
```

```
int32
float64
```

NumPy arrays are **homogeneous**; every element of the array has to be of the same data type. This is what allows NumPy to be so efficient: the array data is stored in contiguous blocks of memory, which makes it easier and faster to access and manipulate.

13.6 Exercises

13.6.0.1 Exercise 1: Creating Arrays

- A. Create a 1D NumPy array containing the values 10, 20, 30, 40, and 50. Print the array.
- B. Create a 2D NumPy array from the following list of lists: `[[1, 2, 3], [4, 5, 6], [7, 8, 9]]`. Print the array.
- C. Use `np.arange()` to create an array containing the numbers from 0 to 15, with a step size of 3. Print the array.

13.6.0.2 Exercise 2: Indexing and Slicing

- A. Given the array `arr = np.array([100, 200, 300, 400, 500])`, print the element at index 2.
- B. For the array `arr_2d = np.array([[10, 20, 30], [40, 50, 60], [70, 80, 90]])`, slice and print the subarray containing the first two rows and the first two columns.
- C. Given the array `arr = np.array([5, 10, 15, 20, 25, 30, 35, 40])`, slice and print every second element starting from the first.

13.6.0.3 Exercise 3: Element-Wise Operations

- A. Create two 1D arrays `arr1 = np.array([1, 2, 3])` and `arr2 = np.array([4, 5, 6])`. Perform element-wise addition, subtraction, and multiplication of these arrays and print the results.
- B. Multiply the array `arr = np.array([2, 4, 6, 8, 10])` by 3 and print the result.
- C. Create an array with values `[2, 4, 6]` and square each element (i.e., multiply each element by itself). Print the resulting array.

13.6.0.4 Exercise 4: Array Shape and Reshaping

- A. Create a 1D array with 12 elements using `np.arange()` and reshape it into a 3x4 array. Print the reshaped array.
- B. Given the array `arr = np.array([[1, 2, 3], [4, 5, 6], [7, 8, 9]])`, print the shape of the array.
- C. Reshape the array `arr = np.array([1, 2, 3, 4, 5, 6])` into a 2x3 matrix and print it.

13.6.0.5 Exercise 5: Mathematical Operations

- A. Given the array `arr = np.array([5, 10, 15, 20, 25])`, compute and print the sum of all elements in the array.
- B. Create an array of numbers from 1 to 5. Compute and print the mean and standard deviation of the array.
- C. Given the array `arr = np.array([3, 7, 2, 8, 1, 9, 6, 5, 4])`, find and print the minimum and maximum values.

13.6.0.6 Exercise 6: Array Comparisons and Conditional Operations

- A. Create an array `arr = np.array([10, 20, 30, 40, 50])`. Use a comparison operator to create a boolean array that checks which elements are greater than 25. Print the result.
- B. Use the boolean array from part (A) to filter the original array and print only the elements greater than 25.
- C. Given the array `arr = np.array([1, 2, 3, 4, 5, 6, 7, 8, 9, 10])`, filter and print only the even numbers.

14 Pandas

14.1 Introduction to Pandas

Data analysis often involves handling large volumes of structured data and performing operations such as data cleaning, transformation, and statistical analysis. The ability to efficiently manipulate data is crucial for extracting meaningful insights. Pandas is a powerful open-source library in Python specifically designed to make these tasks easier and more intuitive.

14.1.1 Why Use Pandas?

Pandas, short for “Panel Data,” is built on top of NumPy and provides high-level data manipulation tools that are both fast and versatile. It is particularly well-suited for handling and analyzing data in the form of tables, such as those found in spreadsheets or SQL databases. The two primary data structures in Pandas, the **Series** and the **DataFrame**, offer powerful ways to organize and manipulate data, making it possible to perform complex data analysis with just a few lines of code.

The development of Pandas by Wes McKinney in 2008 was driven by the need for a flexible and efficient tool that could handle the demands of modern data analysis. Since its creation, Pandas has become the de facto standard for data manipulation in Python, widely used in fields such as finance, statistics, machine learning, and scientific research.

14.1.2 Key Features of Pandas

Pandas offers several key features that make it an essential tool for data analysis:

1. **Data Alignment and Indexing:** Pandas automatically aligns data during arithmetic operations, handling missing data with ease. This feature ensures that data manipulations are efficient and that operations can be performed across different datasets without losing track of the relationships between data points.
2. **Handling Missing Data:** In real-world datasets, missing data is a common occurrence. Pandas provides robust methods for detecting, handling, and imputing missing values, allowing analysts to maintain data integrity and make informed decisions.

3. **Data Cleaning and Preparation:** Pandas simplifies the process of data cleaning, including handling duplicates, filtering data, and transforming variables into the desired formats. Clean data is essential for producing accurate and reliable analysis results.
4. **Flexible Data Structures:** The **Series** and **DataFrame** objects in Pandas are highly flexible. They can handle different data types, including integers, floats, strings, and even objects, all within the same data structure. This flexibility is crucial when dealing with heterogeneous datasets.
5. **Integration with Other Libraries:** Pandas integrates seamlessly with other Python libraries such as NumPy, Matplotlib, and SciPy, making it easy to combine data analysis with numerical operations, visualization, and statistical techniques.
6. **High Performance:** Pandas is optimized for performance with large datasets. Its operations are built on top of the highly efficient NumPy library, allowing for quick execution of complex data manipulations.

14.2 Series

A **Series** in Pandas is a one-dimensional array-like object that can hold data of any type (integers, strings, floating-point numbers, etc.). A Series is similar to a column in a spreadsheet or a database. Each element in a Series has an associated label called an **index**, which allows for intuitive data handling and alignment.

14.2.1 Creating a Series

You can create a Series in Pandas using the `pd.Series()` function, where `pd` is the standard alias for the Pandas module.

```
import pandas as pd

# Creating a Series from a list
data = [10, 20, 30, 40, 50]
series = pd.Series(data)
print(series)
```

```
0    10
1    20
2    30
3    40
4    50
dtype: int64
```

Here, the numbers on the left represent the **index** of the Series, and the numbers on the right are the data values.

14.2.1.1 Accessing Elements in a Series

You can access individual elements in a Series using the index. For example, to access the third element:

```
value = series[2]
print(value)
```

30

14.2.2 Handling Missing Data in a Series

Missing data is a common issue when dealing with real-world datasets. Pandas represents missing values as **NaN** (Not a Number), and it provides built-in methods to handle these missing values.

14.2.2.1 Creating a Series with Missing Data

```
# Creating a Series with missing data
data = [10, 20, None, 40, 50]
series_with_nan = pd.Series(data)
print(series_with_nan)
```

```
0    10.0
1    20.0
2     NaN
3    40.0
4    50.0
dtype: float64
```

In this example, the third value is **None**, which Pandas automatically converts to **NaN** to signify missing data.

14.2.2.2 Handling Missing Data

Pandas provides several methods to handle missing data in a Series:

- **Removing Missing Values:** You can use the `dropna()` method to remove any missing values from the Series.

```
# Removing missing values
clean_series = series_with_nan.dropna()
print(clean_series)
```

```
0    10.0
1    20.0
3    40.0
4    50.0
dtype: float64
```

- **Filling Missing Values:** The `fillna()` method allows you to replace missing values with a specified value.

```
# Filling missing values with a specific number
filled_series = series_with_nan.fillna(0)
print(filled_series)
```

```
0    10.0
1    20.0
2     0.0
3    40.0
4    50.0
dtype: float64
```

These methods provide flexibility in handling missing data, depending on the specific requirements of your analysis.

14.2.3 Handling Duplicates in a Series

Duplicates can often occur in data, and it is essential to handle them to ensure the accuracy of your analysis. Pandas makes it easy to identify and remove duplicate values in a Series.

14.2.3.1 Identifying Duplicates

You can use the `uplicated()` method to identify duplicate values in a Series. This method returns a boolean Series indicating whether each value is a duplicate.

```
# Creating a Series with duplicate values
data = [10, 20, 20, 30, 40, 40, 50]
series_with_duplicates = pd.Series(data)

# Identifying duplicates
duplicates = series_with_duplicates.duplicated()
print(duplicates)
```

```
0    False
1    False
2     True
3    False
4    False
5     True
6    False
dtype: bool
```

14.2.3.2 Removing Duplicates

To remove duplicate values, you can use the `drop_duplicates()` method:

```
# Removing duplicate values
unique_series = series_with_duplicates.drop_duplicates()
print(unique_series)
```

```
0    10
1    20
3    30
4    40
6    50
dtype: int64
```

This method retains only the first occurrence of each value, removing all subsequent duplicates.

14.2.4 Performing Numeric Computations with a Series

One of the strengths of a Pandas Series is its ability to perform vectorized operations, making numeric computations fast and intuitive. You can perform arithmetic operations on a Series just as you would with individual numbers.

14.2.4.1 Basic Arithmetic Operations

Pandas supports element-wise operations, allowing you to perform calculations directly on a Series.

```
# Creating a Series
data = [10, 20, 30, 40, 50]
numeric_series = pd.Series(data)

# Performing arithmetic operations
sum_series = numeric_series + 5
product_series = numeric_series * 2
squared_series = numeric_series ** 2

print("Sum Series:\n", sum_series)
print("Product Series:\n", product_series)
print("Squared Series:\n", squared_series)
```

```
Sum Series:
0    15
1    25
2    35
3    45
4    55
dtype: int64
Product Series:
0    20
1    40
2    60
3    80
4   100
dtype: int64
Squared Series:
0    100
1    400
```

```
2    900
3   1600
4   2500
dtype: int64
```

14.2.4.2 Summary Statistics

Pandas also provides convenient methods for calculating summary statistics on a Series, such as the mean, median, and standard deviation.

```
# Calculating summary statistics
mean_value = numeric_series.mean()
median_value = numeric_series.median()
std_dev = numeric_series.std()

print("Mean:", mean_value)
print("Median:", median_value)
print("Standard Deviation:", std_dev)
```

```
Mean: 30.0
Median: 30.0
Standard Deviation: 15.811388300841896
```

These built-in methods allow you to quickly analyze the statistical properties of your data, providing insights with minimal code.

14.3 DataFrame

A **DataFrame** is a two-dimensional, size-mutable, and potentially heterogeneous tabular data structure with labeled axes (rows and columns). It is similar to a table in a database or an Excel spreadsheet. DataFrames are one of the most commonly used data structures in data analysis and provide a powerful tool for working with structured data in Python.

14.3.1 Creating a DataFrame

DataFrames can be created in various ways, such as from dictionaries, lists, or external data sources like CSV files. Here's an example of creating a DataFrame from a dictionary:


```
import pandas as pd

data = {
    'Name': ['Alice', 'Bob', 'Charlie', 'David'],
    'Age': [24, 27, 22, 32],
    'City': ['New York', 'Los Angeles', 'Chicago', 'Houston']
}

df = pd.DataFrame(data)
print(df)
```

	Name	Age	City
0	Alice	24	New York
1	Bob	27	Los Angeles
2	Charlie	22	Chicago
3	David	32	Houston

The DataFrame structure consists of rows and columns, where each column represents a Series.

14.3.2 Indexing DataFrames

Indexing in Pandas DataFrames is a way to access, filter, and manipulate data efficiently. DataFrames support multiple methods of indexing, which can be used to select specific rows, columns, or even individual values. Understanding how to use indexing properly is essential for effective data analysis.

14.3.2.1 Types of Indexing in DataFrames

Pandas provides several ways to index a DataFrame:

1. **Label-based indexing** using `.loc[]`
2. **Integer-based indexing** using `.iloc[]`
3. **Boolean indexing**

14.3.2.2 Label-based Indexing with `.loc[]`

The `.loc[]` method is used to select data based on the labels of rows and columns. It allows you to access a group of rows and columns by labels or a boolean array.

```
import pandas as pd

# Creating a DataFrame
data = {
    'Name': ['Alice', 'Bob', 'Charlie', 'David'],
    'Age': [24, 27, 22, 32],
    'City': ['New York', 'Los Angeles', 'Chicago', 'Houston']
}

df = pd.DataFrame(data)

# Using .loc[] to select data by labels
# Selecting the row with label 1 (Bob) and specific columns
selected_data = df.loc[1, ['Name', 'City']]
print(selected_data)
```

```
Name      Bob
City  Los Angeles
Name: 1, dtype: object
```

In this example, the `.loc[]` method selects the data from the row with index label 1 and returns the values from the `Name` and `City` columns.

14.3.2.3 Integer-based Indexing with `.iloc[]`

The `.iloc[]` method is used to select data by position, similar to how you would index arrays in Python. It uses integer-based positions to access data.

```
# Using .iloc[] to select data by position
# Selecting the first two rows and the first two columns
selected_data = df.iloc[0:2, 0:2]
print(selected_data)
```

```
   Name  Age
0  Alice  24
1   Bob   27
```

Here, `.iloc[]` selects the data from the first two rows (0:2) and the first two columns (0:2), providing a quick way to subset the DataFrame using integer positions.

14.3.2.4 Boolean Indexing

Boolean indexing allows you to filter data in a DataFrame based on conditions. This is useful for selecting rows that meet specific criteria.

```
# Using Boolean indexing to filter rows where Age is greater than 25
filtered_data = df[df['Age'] > 25]
print(filtered_data)
```

	Name	Age	City
1	Bob	27	Los Angeles
3	David	32	Houston

In this example, we used a condition (`df['Age'] > 25`) to filter the DataFrame, returning only the rows where the `Age` column has values greater than 25.

14.3.3 Handling Missing Data in a DataFrame

Missing data is a frequent issue in datasets. Pandas provides several methods to handle missing data, allowing you to clean and prepare your data efficiently.

14.3.3.1 Creating a DataFrame with Missing Data

```
# Creating a DataFrame with missing data
data = {
    'Name': ['Alice', 'Bob', 'Charlie', 'David'],
    'Age': [24, None, 22, 32],
    'City': ['New York', 'Los Angeles', None, 'Houston']
}

df_with_nan = pd.DataFrame(data)
print(df_with_nan)
```

	Name	Age	City
0	Alice	24.0	New York
1	Bob	NaN	Los Angeles
2	Charlie	22.0	None
3	David	32.0	Houston

14.3.3.2 Handling Missing Data

Pandas provides several ways to handle missing data in a DataFrame:

- **Removing Missing Values:** Use the `dropna()` method to remove any rows or columns with missing values.

```
# Dropping rows with missing values
cleaned_df = df_with_nan.dropna()
print(cleaned_df)
```

	Name	Age	City
0	Alice	24.0	New York
3	David	32.0	Houston

In this example, only the rows without missing values are retained.

- **Filling Missing Values:** Use the `fillna()` method to fill missing values with a specified value or a method like the mean.

```
# Filling missing values in the Age column with the average age
df_filled = df_with_nan.fillna({'Age': df_with_nan['Age'].mean(), 'City': 'Unknown'})
print(df_filled)
```

	Name	Age	City
0	Alice	24.0	New York
1	Bob	26.0	Los Angeles
2	Charlie	22.0	Unknown
3	David	32.0	Houston

In this example, missing values in the **Age** column are filled with the mean of the existing ages, while missing values in the **City** column are replaced with “Unknown.”

14.3.4 Handling Duplicates in a DataFrame

Duplicate data can often cause inaccuracies in analysis. Pandas makes it easy to identify and remove duplicates.

14.3.4.1 Identifying Duplicates

You can use the `duplicated()` method to find duplicate rows in the DataFrame.

```
# Creating a DataFrame with duplicate rows
data = {
    'Name': ['Alice', 'Bob', 'Charlie', 'David', 'Alice'],
    'Age': [24, 27, 22, 32, 24],
    'City': ['New York', 'Los Angeles', 'Chicago', 'Houston', 'New York']
}

df_with_duplicates = pd.DataFrame(data)

# Identifying duplicate rows
duplicates = df_with_duplicates.duplicated()
print(duplicates)
```

```
0    False
1    False
2    False
3    False
4     True
dtype: bool
```

14.3.4.2 Removing Duplicates

To remove duplicate rows, use the `drop_duplicates()` method:

```
# Removing duplicate rows
df_unique = df_with_duplicates.drop_duplicates()
print(df_unique)
```

This removes the duplicate occurrence of Alice from the DataFrame.

14.3.5 Performing Numeric Computations with a DataFrame

DataFrames support a wide range of arithmetic operations that can be performed on numerical data, making it easy to compute and analyze your data.

14.3.5.1 Basic Arithmetic Operations

You can perform operations on individual columns or the entire DataFrame.

```
# Creating a DataFrame
data = {
    'Product': ['A', 'B', 'C'],
    'Price': [100, 150, 200],
    'Quantity': [3, 4, 5]
}

df = pd.DataFrame(data)

# Calculating the total cost for each product
df['Total Cost'] = df['Price'] * df['Quantity']
print(df)
```

	Product	Price	Quantity	Total Cost
0	A	100	3	300
1	B	150	4	600
2	C	200	5	1000

This example demonstrates how to create a new column (Total Cost) by multiplying the Price and Quantity columns.

14.3.5.2 Summary Statistics

Pandas makes it easy to calculate summary statistics for a DataFrame using built-in methods like `mean()`, `sum()`, and `max()`.

```
# Calculating summary statistics
average_price = df['Price'].mean()
total_quantity = df['Quantity'].sum()

print("Average Price:", average_price)
print("Total Quantity Sold:", total_quantity)
```

```
Average Price: 150.0
Total Quantity Sold: 12
```

These operations allow you to quickly analyze data trends and obtain insights with minimal code.

14.4 Sorting, Filtering, and Merging DataFrames

When working with data, it is often necessary to sort, filter, and merge datasets to analyze information efficiently. Pandas provides intuitive methods for these operations, allowing you to organize and manipulate data easily.

14.4.1 Sorting DataFrames

Sorting data is essential when you want to analyze your DataFrame in a specific order. Pandas allows you to sort DataFrames by their row indices or column values using the `sort_values()` and `sort_index()` methods.

14.4.1.1 Sorting by Column Values

You can sort a DataFrame by the values in one or more columns using the `sort_values()` method.

```
import pandas as pd

# Creating a DataFrame
data = {
    'Name': ['Alice', 'Bob', 'Charlie', 'David'],
    'Age': [24, 27, 22, 32],
    'City': ['New York', 'Los Angeles', 'Chicago', 'Houston']
}

df = pd.DataFrame(data)

# Sorting the DataFrame by the 'Age' column in ascending order
sorted_df = df.sort_values(by='Age')
print(sorted_df)
```

	Name	Age	City
2	Charlie	22	Chicago
0	Alice	24	New York
1	Bob	27	Los Angeles
3	David	32	Houston

This example sorts the DataFrame by the **Age** column, arranging the rows in ascending order. You can sort in descending order by setting the parameter `ascending=False`.

14.4.1.2 Sorting by Index

To sort the DataFrame by its index, use the `sort_index()` method.

```
# Sorting the DataFrame by its index in descending order
sorted_by_index = df.sort_index(ascending=False)
print(sorted_by_index)
```

	Name	Age	City
3	David	32	Houston
2	Charlie	22	Chicago
1	Bob	27	Los Angeles
0	Alice	24	New York

This arranges the DataFrame in descending order based on the row index.

14.4.2 Filtering DataFrames

Filtering data allows you to extract specific rows that meet particular conditions. This is useful for analyzing subsets of your data.

14.4.2.1 Filtering Rows Based on a Condition

You can filter a DataFrame to include only the rows that meet a specific condition using boolean indexing.

```
# Filtering rows where the Age is greater than 25
filtered_df = df[df['Age'] > 25]
print(filtered_df)
```

	Name	Age	City
1	Bob	27	Los Angeles
3	David	32	Houston

In this example, only the rows where the **Age** column has a value greater than 25 are included.

14.4.2.2 Filtering with Multiple Conditions

You can also filter rows based on multiple conditions using the logical operators `&` (and), `|` (or), and `~` (not).

```
# Filtering rows where Age is greater than 25 and City is 'Los Angeles'
filtered_df = df[(df['Age'] > 25) & (df['City'] == 'Los Angeles')]
print(filtered_df)
```

	Name	Age	City
1	Bob	27	Los Angeles

This filters the DataFrame to include only rows where the **Age** is greater than 25 and the **City** is 'Los Angeles'.

14.4.3 Merging DataFrames

Merging is the process of combining multiple DataFrames into a single DataFrame based on a common key or index. Pandas provides several methods for merging, including `merge()`, `concat()`, and `join()`.

14.4.3.1 Merging DataFrames with `merge()`

The `merge()` function is similar to SQL joins and is used to combine DataFrames based on a key column.

```
# Creating two DataFrames to merge
data1 = {'Name': ['Alice', 'Bob', 'Charlie'],
         'Age': [24, 27, 22]}
df1 = pd.DataFrame(data1)

data2 = {'Name': ['Alice', 'Bob', 'David'],
         'City': ['New York', 'Los Angeles', 'Houston']}
df2 = pd.DataFrame(data2)

# Merging the DataFrames on the 'Name' column
merged_df = pd.merge(df1, df2, on='Name', how='inner')
print(merged_df)
```

	Name	Age	City
0	Alice	24	New York
1	Bob	27	Los Angeles

In this example, we performed an **inner join** on the **Name** column, returning only the rows with matching values in both DataFrames. You can change the **how** parameter to perform different types of joins:

- **how='inner'** (default) returns only the intersection of the keys.
- **how='outer'** returns all keys from both DataFrames.
- **how='left'** returns all keys from the left DataFrame and matches them with the right.
- **how='right'** returns all keys from the right DataFrame and matches them with the left.

14.4.3.2 Concatenating DataFrames with `concat()`

The `concat()` function is used to concatenate DataFrames along a particular axis (rows or columns).

```
# Concatenating two DataFrames along the rows (axis=0)
df3 = pd.DataFrame({'Name': ['Eve', 'Frank'],
                    'Age': [29, 34],
                    'City': ['Seattle', 'Boston']})

concatenated_df = pd.concat([df, df3], ignore_index=True)
print(concatenated_df)
```

	Name	Age	City
0	Alice	24	New York
1	Bob	27	Los Angeles
2	Charlie	22	Chicago
3	David	32	Houston
4	Eve	29	Seattle
5	Frank	34	Boston

This example shows how to concatenate two DataFrames by stacking them on top of each other.

14.5 Reading and Writing Data

One of the most common tasks in data analysis is reading data from files and writing data to them. Pandas makes it easy to handle these operations with its built-in functions.

14.5.1 Reading Data from Files

Pandas supports reading data from various file formats, such as CSV, Excel, JSON, and more. The most commonly used format is CSV (Comma-Separated Values). You can read a CSV file into a DataFrame using the `pd.read_csv()` function:

```
df = pd.read_csv('data.csv')
```

This command reads the contents of the `data.csv` file and stores it as a DataFrame in the variable `df`. Pandas automatically infers the data types and the structure of the file.

14.5.2 Writing Data to Files

Similarly, you can write a DataFrame to a file using the `to_csv()` method. This is useful for saving your data after performing analysis or making modifications.

```
df.to_csv('output.csv', index=False)
```

The parameter `index=False` ensures that the index is not included in the saved file.

14.6 Exercises

14.6.0.1 Exercise 1: Creating and Manipulating Series

Create a Pandas Series from the following list of temperatures (in Celsius): `[22, 25, 23, 20, 27, None, 28, 24]`.

- A. Print the Series.
- B. Identify any missing data.
- C. Replace the missing data with the average of the other temperatures.

14.6.0.2 Exercise 2: Creating and Manipulating Series

Create a Series of temperatures using the list `[22, 25, 23, 22, 25, 28, 24, 22]`.

- A. Identify and remove any duplicate values.
- B. Sort the Series in descending order.

14.6.0.3 Exercise 3: Creating and Manipulating Series

Perform the following calculations on the Series you created in Exercise 2:

- A. Multiply each temperature by $9/5$ and then add 32 to convert the values from Celsius to Fahrenheit.
- B. Find the mean, median, and standard deviation of the converted temperatures.

14.6.0.4 Exercise 4: Working with DataFrames

Create a DataFrame from the following dictionary:

```
data = {
    'Name': ['Alice', 'Bob', 'Charlie', 'David'],
    'Age': [24, None, 22, 32],
    'City': ['New York', 'Los Angeles', 'Chicago', 'Houston']
}
```

- A. Print the DataFrame.
- B. Identify and remove any rows with missing values.
- C. Fill in the missing values in the **Age** column with the average age.

14.6.0.5 Exercise 5: Working with DataFrames

Using the same DataFrame from Exercise 4, perform the following tasks:

- A. Sort the DataFrame by the **Age** column in ascending order.
- B. Filter the DataFrame to include only rows where the **Age** is greater than 23.

14.6.0.6 Exercise 6: Working with DataFrames

Create two new DataFrames:

- DataFrame 1:

```
data1 = {'Name': ['Alice', 'Bob', 'Charlie'],
         'Age': [24, 27, 22]}
```

- DataFrame 2:

```
data2 = {'Name': ['Alice', 'Charlie', 'David'],
         'City': ['New York', 'Chicago', 'Houston']}
```

- A. Merge these two DataFrames on the `Name` column using an inner join.
- B. Perform an outer join on the same DataFrames and observe the differences.

14.6.0.7 Exercise 7: Sorting, Filtering, and Merging

Create a DataFrame using the following data:

```
data = {
    'Product': ['Apples', 'Bananas', 'Cherries', 'Dates'],
    'Price': [3, 1, 4, 2],
    'Quantity': [5, 7, 6, 3]
}
```

- A. Sort the DataFrame by `Price` in ascending order.
- B. Filter the DataFrame to include only products with a `Quantity` greater than 4.
- C. Add a new column to the DataFrame called `Total Cost` that calculates the total cost as `Price * Quantity`.

14.6.0.8 Exercise 8: Sorting, Filtering, and Merging

Create a new DataFrame with the following data:

```
data = {
    'Product': ['Apples', 'Cherries', 'Dates'],
    'Supplier': ['Supplier A', 'Supplier B', 'Supplier C']
}
```

- A. Merge this DataFrame with the DataFrame from Exercise 7 using the `Product` column as the key.
- B. Experiment with different types of joins (inner, outer, left, and right) and note the changes in the resulting DataFrame.

14.6.0.9 Exercise 9: Reading and Writing Data

Save the following DataFrame to a CSV file named `students.csv`:

```
data = {  
    'Name': ['Anna', 'Ben', 'Cathy', 'Dan'],  
    'Score': [88, 92, 85, 90]  
}
```

- A. Write the DataFrame to the CSV file, ensuring that the index is not included in the file.
- B. Read the data back into a DataFrame and display its contents.

15 Data Visualization in Python

Data visualization is a crucial aspect of data analysis that enables the understanding and interpretation of data through graphical representation. It allows us to observe trends, patterns, and outliers in data that might not be evident in raw numerical forms. In Python, there are several powerful libraries for creating visualizations, with Matplotlib, ggplot, and Seaborn being among the most popular. This section will introduce these libraries and demonstrate how to create a variety of plots to explore and present data effectively.

15.1 Introduction to Matplotlib

Matplotlib is a comprehensive library for creating static, animated, and interactive visualizations in Python. It is often the first choice for many data scientists and analysts due to its flexibility and the detailed control it offers over plots. Matplotlib is organized into several components, with the `pyplot` module being the most commonly used. This module provides a MATLAB-like interface that simplifies the process of creating and customizing plots.

15.1.1 Creating Basic Plots with Matplotlib

The most fundamental plot in Matplotlib is the line plot, which is typically used to visualize trends over time or continuous data. To get started with plotting in Matplotlib, you need to import the `pyplot` module:

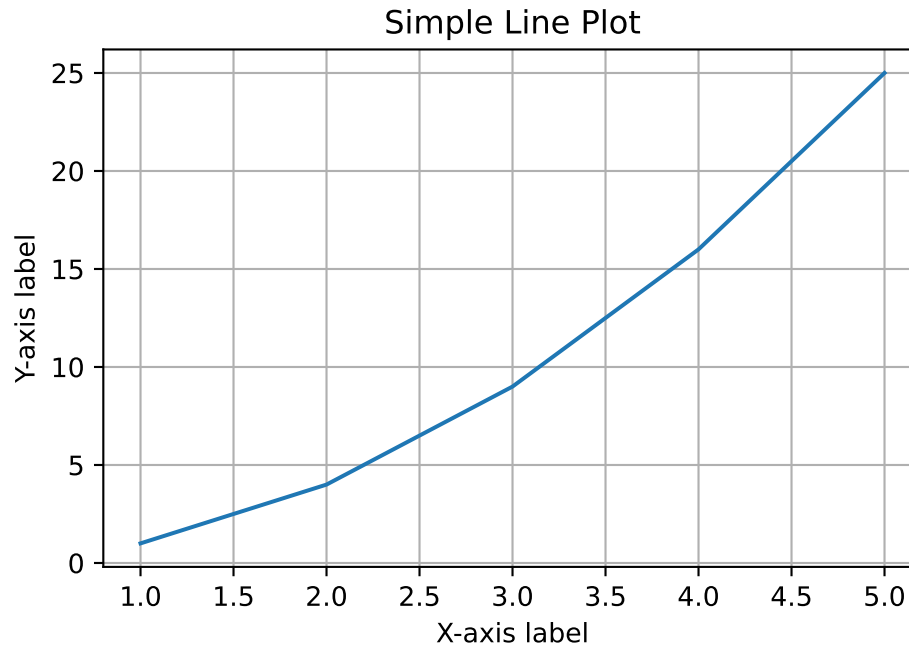
```
import matplotlib.pyplot as plt
```

Example: Plotting a Simple Line Graph

```
# Data for plotting
x = [1, 2, 3, 4, 5]
y = [1, 4, 9, 16, 25]

# Creating a line plot
plt.plot(x, y)
plt.xlabel('X-axis label')
```

```
plt.ylabel('Y-axis label')
plt.title('Simple Line Plot')
plt.grid(True) # Adds a grid to the plot
plt.show()
```



This example creates a simple line plot with labeled axes, a title, and a grid. The method `plt.show()` is used to render the plot.

15.1.2 Types of Plots in Matplotlib


Matplotlib supports a wide variety of plot types beyond simple line graphs. Some of the most commonly used plot types include:

1. **Scatter Plots:** Useful for exploring relationships between two variables.
2. **Bar Charts:** Ideal for comparing discrete categories.
3. **Histograms:** Commonly used to display the distribution of data.
4. **Pie Charts:** Used to show proportions of a whole.
5. **Box Plots:** Effective for visualizing the spread and outliers of data.

Example: Creating a Scatter Plot


```
# Data for scatter plot
x = [5, 7, 8, 5, 9, 7]
y = [3, 8, 4, 7, 2, 5]

# Creating the scatter plot
plt.scatter(x, y, color='blue', marker='^')
plt.xlabel('X-axis label')
plt.ylabel('Y-axis label')
plt.title('Scatter Plot Example')
plt.show()
```



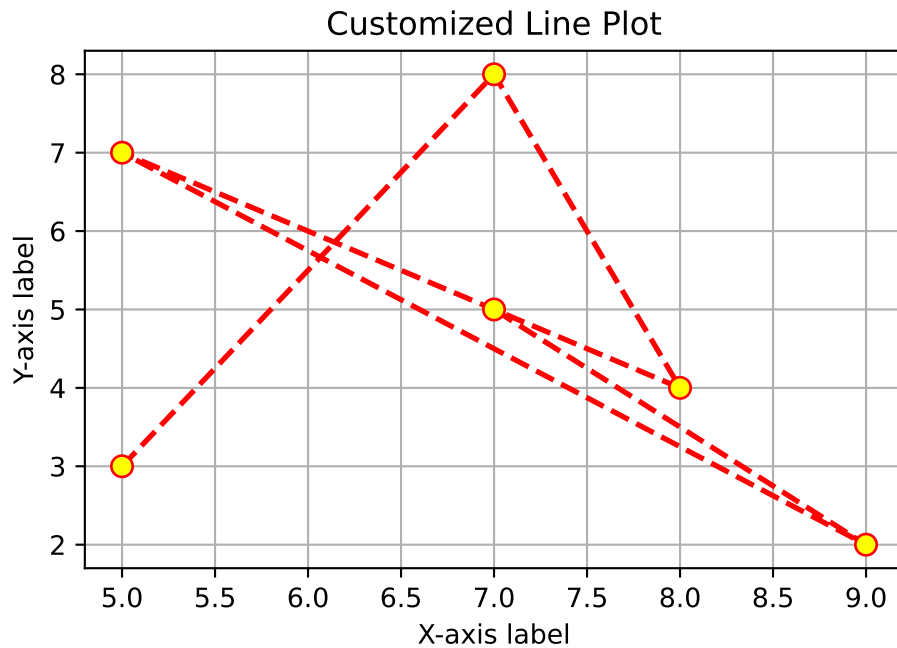
In this example, we use `plt.scatter()` to create a scatter plot, where the color and marker style of the points can be easily customized.

15.1.3 Customizing Plots in Matplotlib

Customization is one of Matplotlib's greatest strengths. You can modify almost every aspect of a plot, including colors, line styles, markers, axis scales, legends, and more. This flexibility allows you to create publication-quality visualizations.

Example: Customizing a Line Plot

```
# Customizing the line plot
plt.plot(x, y, color='red', linestyle='--', linewidth=2, marker='o', markersize=8, markerfacecolor='yellow')
plt.xlabel('X-axis label')
plt.ylabel('Y-axis label')
plt.title('Customized Line Plot')
plt.grid(True)
plt.show()
```



In this example, we modify the line style to be dashed (--), set the line color to red, increase the line width, and customize the marker appearance with size and color adjustments.

15.1.4 Adding Annotations and Legends

Annotations and legends can significantly enhance the interpretability of your plots. Annotations allow you to highlight specific data points or trends, while legends help distinguish different datasets in a multi-line plot.

Example: Adding a Legend and Annotations

```


# Data for multiple lines
x = [1, 2, 3, 4, 5]
y1 = [1, 4, 9, 16, 25]
y2 = [2, 3, 4, 5, 6]

# Plotting multiple lines
plt.plot(x, y1, label='Dataset 1', color='blue')
plt.plot(x, y2, label='Dataset 2', color='green')

# Adding an annotation
plt.annotate('Intersection', xy=(3, 9), xytext=(4, 15),
            arrowprops=dict(facecolor='black', shrink=0.05))

# Customizing plot
plt.xlabel('X-axis label')
plt.ylabel('Y-axis label')
plt.title('Annotated Plot with Legend')
plt.legend(loc='upper left')
plt.grid(True)
plt.show()

```



This example demonstrates how to use the `annotate()` function to add a text label to a specific

point on the plot, along with an arrow pointing to that location. The `legend()` function is used to add a legend that describes the different lines in the plot.

15.1.5 Working with Subplots

Subplots allow you to create multiple plots in a single figure, making it easier to compare different visualizations side by side.



Example: Creating Subplots

```
# Creating a figure with 2 rows and 1 column of subplots
fig, axs = plt.subplots(2, 1, figsize=(6, 8))

# First subplot
axs[0].plot(x, y1, color='blue')
axs[0].set_title('Line Plot 1')
axs[0].set_xlabel('X-axis')
axs[0].set_ylabel('Y-axis')

# Second subplot
axs[1].scatter(x, y2, color='green')
axs[1].set_title('Scatter Plot 2')
axs[1].set_xlabel('X-axis')
axs[1].set_ylabel('Y-axis')

# Display the plots
plt.tight_layout() # Adjusts spacing between subplots
plt.show()
```



This example uses the `subplots()` function to create a figure with two subplots arranged in a column. Each subplot is customized with its own title, labels, and data.

15.1.6 Styling and Themes in Matplotlib


Matplotlib offers several built-in styles that can be applied to your plots to give them a consistent and visually appealing look. Styles can be easily switched using the `plt.style.use()` function.

Example: Using Different Styles

```
# Applying a style to the plot
plt.style.use('ggplot')

# Replotting the data with the new style
plt.plot(x, y1, color='red', linestyle='-', marker='o')
plt.title('Styled Plot with ggplot Theme')
plt.xlabel('X-axis')
plt.ylabel('Y-axis')
plt.show()

# Reset to the default style
plt.style.use('default')
```



In this example, we apply the `ggplot` style to the plot, which gives it a different color scheme and grid design inspired by the aesthetics of the `ggplot2` package in R.

15.1.7 Saving Plots to Files

Once you've created your visualizations, you might want to save them as image files for use in presentations or publications.

Example: Saving a Plot

```
# Saving the plot to a file
plt.plot(x, y1, color='blue', linestyle='--')
plt.title('Plot to be Saved')
plt.savefig('my_plot.png', dpi=300, bbox_inches='tight')
```

The `savefig()` function saves the plot as a PNG file with high resolution (specified by `dpi=300`) and tight bounding boxes to remove excess whitespace.

15.1.8 Common Pitfalls and Best Practices

- **Aspect Ratio:** Always check the aspect ratio of your plots to ensure they are not distorted.
- **Labels and Titles:** Make sure to label all axes and include a title to provide context to the visualization.
- **Color Blindness:** Use color palettes that are accessible to those with color vision deficiencies.
- **Legend Placement:** Place legends in a position that does not overlap with important data points.

Understanding how to leverage the power of Matplotlib will enable you to create informative and professional-quality data visualizations, enhancing your ability to communicate data-driven insights.

15.2 Introduction to Seaborn

Seaborn is a powerful Python visualization library built on top of Matplotlib, designed specifically for creating informative and attractive statistical graphics. It provides a high-level interface for drawing a variety of statistical plots and integrates seamlessly with data structures like Pandas DataFrames. Seaborn's design philosophy emphasizes the importance of aesthetics, making it easy to generate complex visualizations with concise code.

15.2.1 Getting Started with Seaborn

To use Seaborn, you first need to import the library. If you haven't already installed it, you can do so using the following command:

```
!pip install seaborn
```

Once installed, import Seaborn into your Python environment:

```
import seaborn as sns
import matplotlib.pyplot as plt
```


15.2.2 Creating Basic Plots with Seaborn

Seaborn simplifies the process of creating plots by offering specific functions for different types of statistical visualizations. Some of the most commonly used plot types include scatter plots, line plots, histograms, and bar charts.

Example: Creating a Scatter Plot

```
# Sample data
tips = sns.load_dataset('tips')

# Creating a scatter plot
sns.scatterplot(x='total_bill', y='tip', data=tips)
plt.title('Scatter Plot of Total Bill vs. Tip')
plt.xlabel('Total Bill')
plt.ylabel('Tip')
plt.show()
```



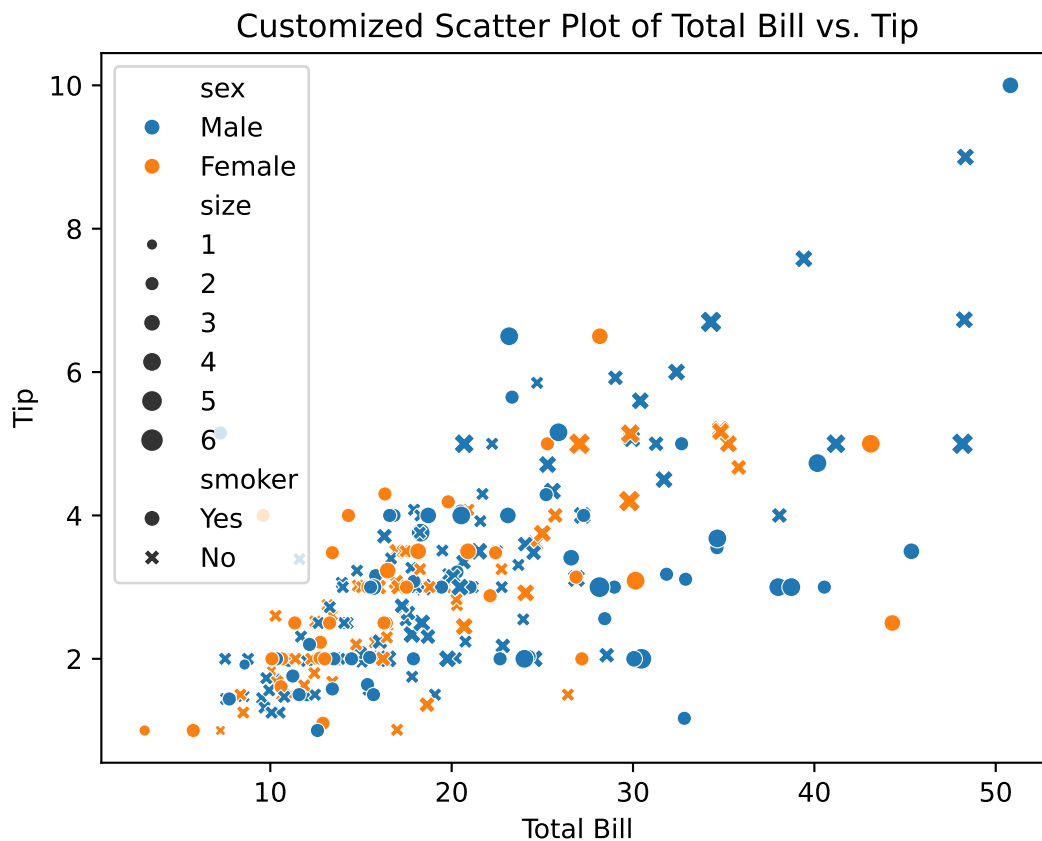
In this example, we use the `tips` dataset (which is built into Seaborn) to create a scatter plot that shows the relationship between the total bill and the tip amount.

15.2.3 Customizing Visualizations with Seaborn

Seaborn allows extensive customization of plot aesthetics to make visualizations more informative and appealing. You can easily modify colors, markers, and styles.

Example: Customizing a Scatter Plot

```
# Custom scatter plot with color and marker variations
sns.scatterplot(x='total_bill', y='tip', hue='sex', style='smoker', size='size', data=tips)
plt.title('Customized Scatter Plot of Total Bill vs. Tip')
plt.xlabel('Total Bill')
plt.ylabel('Tip')
plt.show()
```




Here, the scatter plot is customized to differentiate data points based on additional variables using color (**hue**), marker style (**style**), and size (**size**). This makes it easier to visualize relationships between multiple variables in a single plot.

15.2.4 Visualizing Distributions

Understanding the distribution of data is crucial in statistical analysis. Seaborn offers several functions specifically for visualizing distributions, such as histograms, KDE plots, and box plots.

Example: Creating a Histogram and KDE Plot

```
# Histogram and KDE plot for total bill
sns.histplot(tips['total_bill'], kde=True, color='blue')
plt.title('Distribution of Total Bill with Histogram and KDE')
plt.xlabel('Total Bill')
plt.ylabel('Frequency')
plt.show()
```



In this example, we create a histogram of the `total_bill` variable with an overlaid Kernel Density Estimate (KDE), which provides a smoothed curve representing the distribution.

15.2.5 Creating Categorical Plots

Seaborn excels in creating categorical plots that help analyze the relationship between categorical data and numerical data. The most common types include bar plots, box plots, and violin plots.

Example: Bar Plot

```
# Bar plot showing the average tip amount by day
sns.barplot(x='day', y='tip', data=tips, ci='sd', palette='viridis')
plt.title('Average Tip by Day')
plt.xlabel('Day')
plt.ylabel('Average Tip')
plt.show()
```

C:\Users\Joshua_Patrick\AppData\Local\Temp\ipykernel_42632\690395935.py:2: FutureWarning:

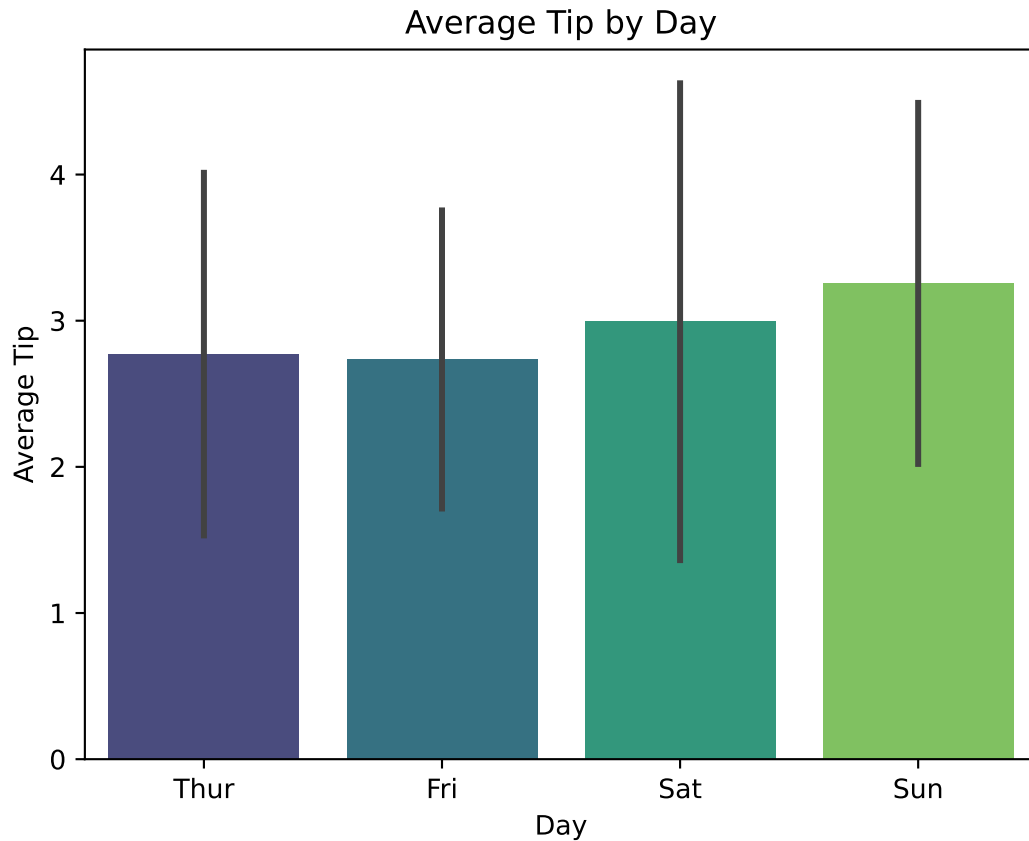
The `ci` parameter is deprecated. Use `errorbar='sd'` for the same effect.

```
sns.barplot(x='day', y='tip', data=tips, ci='sd', palette='viridis')
```

C:\Users\Joshua_Patrick\AppData\Local\Temp\ipykernel_42632\690395935.py:2: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign


```
sns.barplot(x='day', y='tip', data=tips, ci='sd', palette='viridis')
```



This bar plot displays the average tip amount for each day of the week, using error bars to indicate the standard deviation. Seaborn's color palettes, like `viridis`, enhance the plot's visual appeal.

Example: Box Plot

```
# Box plot showing the distribution of total bill by day
sns.boxplot(x='day', y='total_bill', hue='smoker', data=tips, palette='Set2')
plt.title('Box Plot of Total Bill by Day')
plt.xlabel('Day')
plt.ylabel('Total Bill')
plt.show()
```



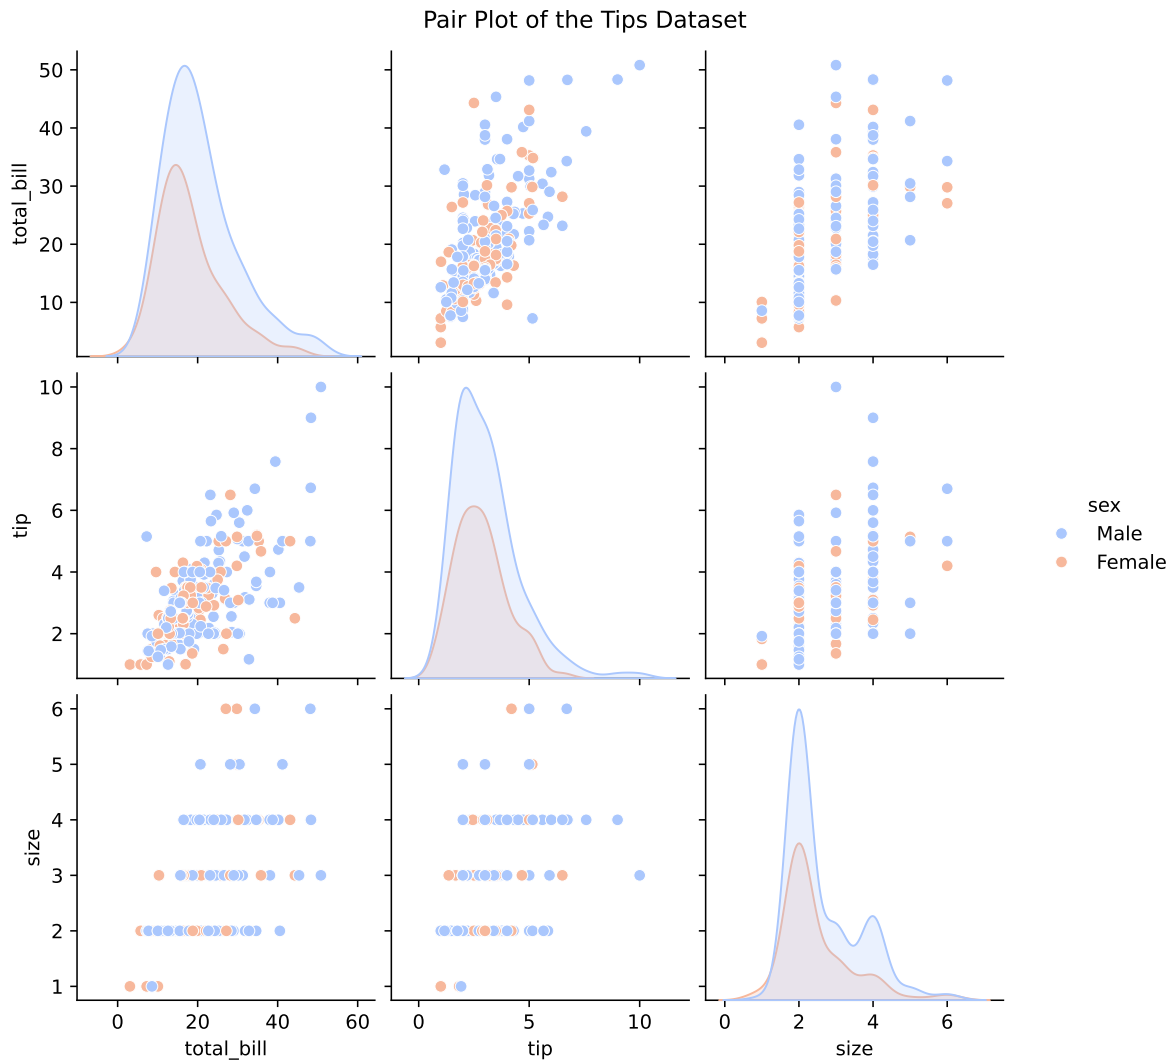
The box plot provides insights into the distribution of the `total_bill` by day, highlighting the median, interquartile range (IQR), and potential outliers. The `hue` parameter adds another layer to the plot, differentiating between smokers and non-smokers.

15.2.6 Pair Plots for Multivariate Analysis

Pair plots are useful for visualizing relationships between multiple variables at once. They create a grid of scatter plots for each pair of variables and display histograms for individual variable distributions.

Example: Pair Plot

```
# Creating a pair plot of the tips dataset
sns.pairplot(tips, hue='sex', palette='coolwarm')
plt.suptitle('Pair Plot of the Tips Dataset', y=1.02)
plt.show()
```




In this example, the pair plot provides a comprehensive view of pairwise relationships between all numerical variables in the dataset, with color-coded distinctions based on the `sex` variable.

15.2.7 Heatmaps for Correlation Analysis

Heatmaps are an effective way to visualize the correlation matrix of numerical data, showing how strongly pairs of variables are related.

Example: Creating a Heatmap

```
# Correlation heatmap for the tips dataset
corr = tips.corr(numeric_only = True)
sns.heatmap(corr, annot=True, cmap='coolwarm', linewidths=0.5)
plt.title('Correlation Heatmap of Tips Dataset')
plt.show()
```



The heatmap in this example displays the correlation coefficients between variables in the `tips` dataset, with color intensity indicating the strength of the correlation. The `annot=True` parameter ensures that the correlation values are displayed on the heatmap.

15.2.8 Customizing Seaborn Themes

Seaborn comes with several built-in themes to enhance the appearance of plots, making it easy to change the overall look with minimal code.

Example: Applying a Different Theme



```
# Set a Seaborn style for all plots
sns.set_style('whitegrid')

# Creating a box plot with the new style
sns.boxplot(x='day', y='total_bill', data=tips, palette='pastel')
plt.title('Box Plot with Whitegrid Theme')
plt.xlabel('Day')
plt.ylabel('Total Bill')
plt.show()
```

C:\Users\Joshua_Patrick\AppData\Local\Temp\ipykernel_42632\3431770401.py:5: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assigning `hue` to the variable will resolve this warning.

```
sns.boxplot(x='day', y='total_bill', data=tips, palette='pastel')
```



In this example, the `whitegrid` style adds grid lines to the plot background, which can help in analyzing data points more effectively.

15.2.9 Combining Seaborn with Matplotlib

Although Seaborn provides a high-level interface for creating plots, it is often beneficial to use Matplotlib functions for further customization.

Example: Combining Seaborn and Matplotlib for Plot Customization


```
# Creating a violin plot with Seaborn
sns.violinplot(x='day', y='total_bill', data=tips, inner='quartile', palette='Set3')

# Customizing with Matplotlib
plt.axhline(y=20, color='red', linestyle='--') # Add a reference line
plt.title('Violin Plot with Custom Reference Line')
plt.xlabel('Day')
plt.ylabel('Total Bill')
plt.show()
```

C:\Users\Joshua_Patrick\AppData\Local\Temp\ipykernel_42632\122318432.py:2: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign

```
sns.violinplot(x='day', y='total_bill', data=tips, inner='quartile', palette='Set3')
```



In this example, a Seaborn violin plot is enhanced with a Matplotlib horizontal line, illustrating the synergy between the two libraries for creating complex visualizations.

15.2.10 Advantages of Using Seaborn

- **Simplified Syntax:** Seaborn's API is intuitive and reduces the code complexity for creating statistical graphics.
- **Built-in Themes:** It offers aesthetically pleasing color palettes and themes that enhance the visual appeal.
- **Integration with Pandas:** Seamlessly works with Pandas DataFrames, making data manipulation and visualization straightforward.
- **Statistical Visualization:** Designed specifically for statistical plotting, providing functions that directly interpret and visualize statistical relationships.

15.2.11 Limitations and Considerations

- **Customization Limitations:** Although Seaborn is powerful, certain highly specific customizations may require falling back on Matplotlib.
- **Learning Curve:** Requires familiarity with data structures like DataFrames and understanding of statistical concepts for full utilization.

15.3 Introduction to ggplot (plotnine)

The `ggplot` library in Python, made accessible through the `plotnine` package, is inspired by the grammar of graphics principles introduced by Leland Wilkinson. This approach allows data visualizations to be built by combining independent layers, where each layer adds a new element to the plot. This methodology enables the creation of highly customizable and aesthetically pleasing plots with a structured syntax that emphasizes clarity and reproducibility.

15.3.1 The Grammar of Graphics

The grammar of graphics is a theoretical framework that breaks down the process of constructing data visualizations into independent components:

- **Data:** The dataset being visualized.
- **Aesthetics (aes):** The mapping of data variables to visual properties, such as position, color, and size.
- **Geometries (geoms):** The visual elements that represent data points, such as lines, bars, points, or histograms.
- **Facets:** The ability to split data into subplots for comparison.
- **Statistics:** Transformation or summary of data to highlight patterns.
- **Scales:** Controls the mapping between data values and their visual representation.
- **Coordinate systems:** Defines how data points are placed in a plot.

Understanding this layered approach helps in constructing meaningful and well-organized visualizations.

15.3.2 Creating a Basic Plot with ggplot

To use `ggplot` in Python, you will need to install the `plotnine` package if you haven't already:

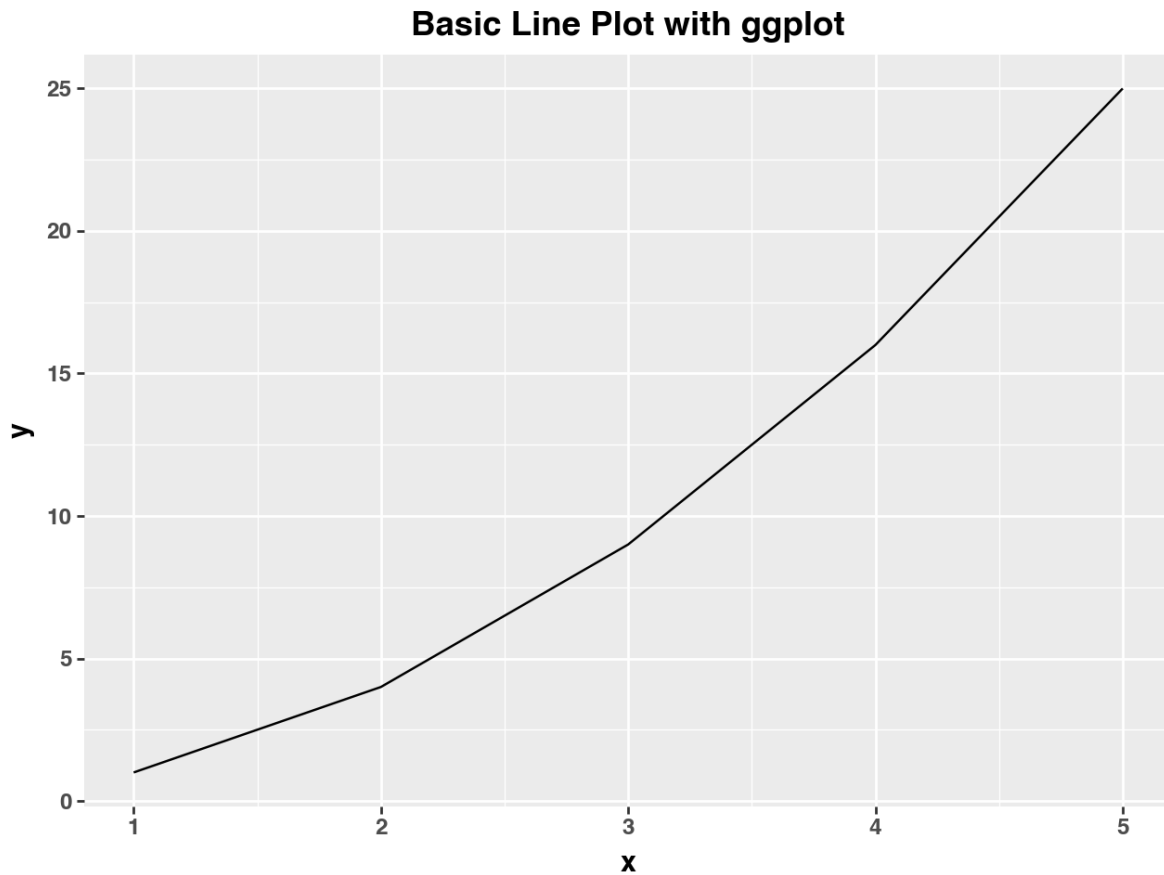
```
!pip install plotnine
```

Now, let's begin by creating a basic line plot using `plotnine`. We will use a simple dataset to demonstrate how to create a visualization using the `ggplot` syntax.

```
import pandas as pd
from plotnine import ggplot, aes, geom_line, ggtitle

# Creating a dataset
data = pd.DataFrame({
    'x': [1, 2, 3, 4, 5],
    'y': [1, 4, 9, 16, 25]
})

# Creating a basic line plot
(ggplot(data, aes(x='x', y='y')) +
 geom_line() +
 ggtitle('Basic Line Plot with ggplot'))
```



In this example, we define the dataset and specify the aesthetics using `aes()`, which maps the `x` and `y` variables to the axes of the plot. The `geom_line()` function adds a line graph layer to the plot. The title is set using `ggtitle()`.

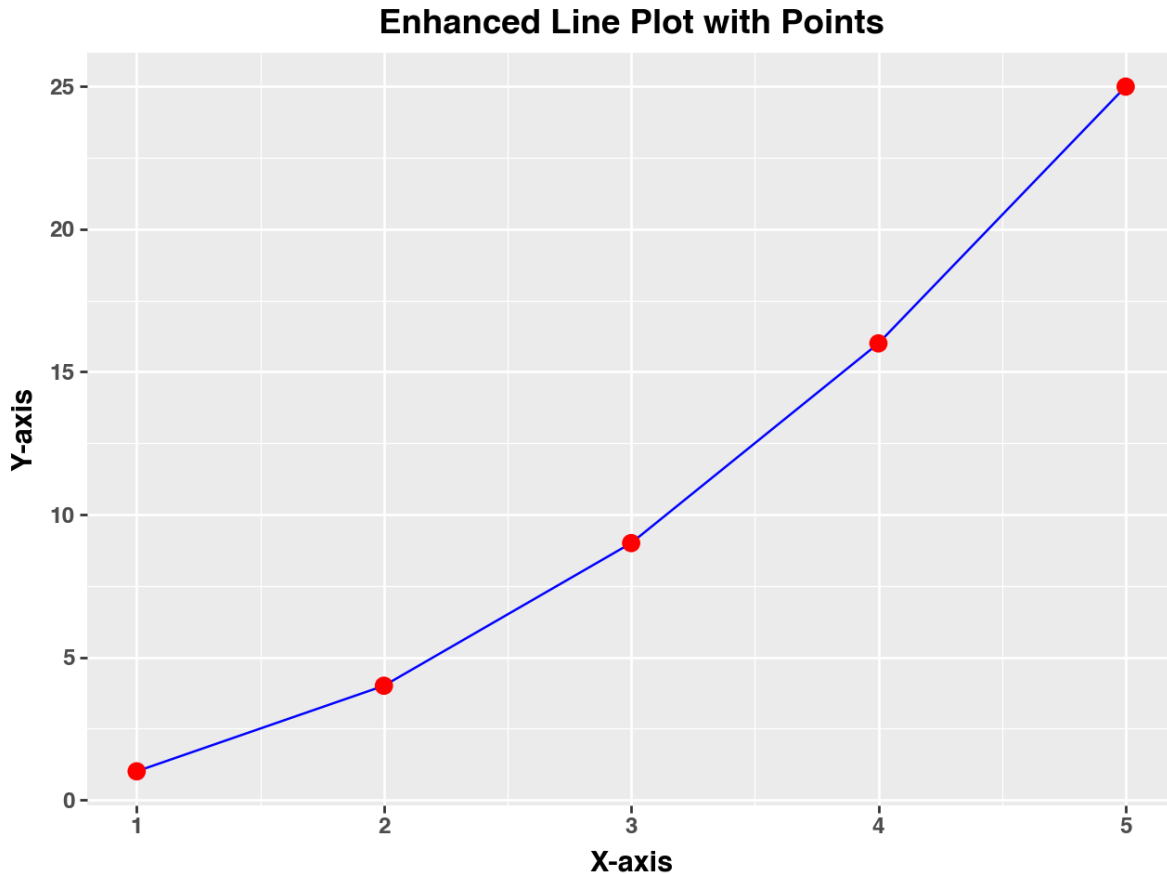
15.3.3 Enhancing Visualizations with Layers

One of the strengths of `ggplot` is the ability to add multiple layers to a plot, allowing you to build more informative visualizations.

Example: Adding Points to a Line Plot

```
from plotnine import geom_point, xlab, ylab

# Adding a layer of points to the line plot
(ggplot(data, aes(x='x', y='y')) +
 geom_line(color='blue') +
 geom_point(color='red', size=3) +
 ggtitle('Enhanced Line Plot with Points') +
 xlab('X-axis') +
 ylab('Y-axis'))
```



Here, we use the `geom_point()` function to overlay points on the existing line plot. This combination helps highlight individual data points while still showing the overall trend.

15.3.4 Common Plot Types in ggplot

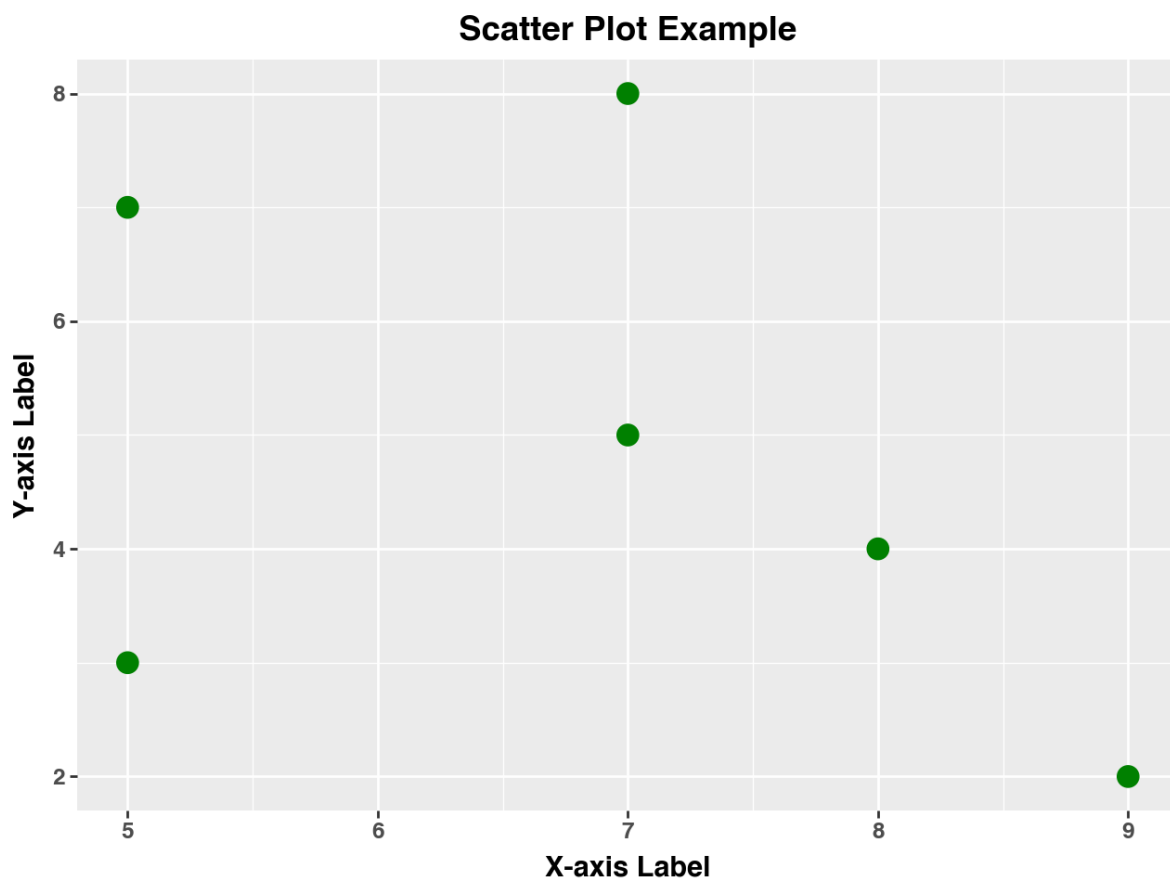
ggplot supports a wide variety of plot types, each suited for different types of data analysis:

1. **Scatter Plots:** Used to visualize relationships between two continuous variables.
2. **Bar Charts:** Ideal for comparing categorical data.
3. **Histograms:** Useful for visualizing the distribution of a single variable.
4. **Box Plots:** Helps to visualize the spread, central tendency, and outliers in data.

Example: Creating a Scatter Plot

```
# Creating a dataset for the scatter plot
scatter_data = pd.DataFrame({
    'x': [5, 7, 8, 5, 9, 7],
    'y': [3, 8, 4, 7, 2, 5]
})

# Scatter plot with ggplot
(ggplot(scatter_data, aes(x='x', y='y')) +
 geom_point(color='green', size=4) +
 ggtitle('Scatter Plot Example') +
 xlab('X-axis Label') +
 ylab('Y-axis Label'))
```



In this example, the scatter plot is created using `geom_point()`, with customizations for point color and size.

15.3.5 Faceting for Multi-Plot Layouts

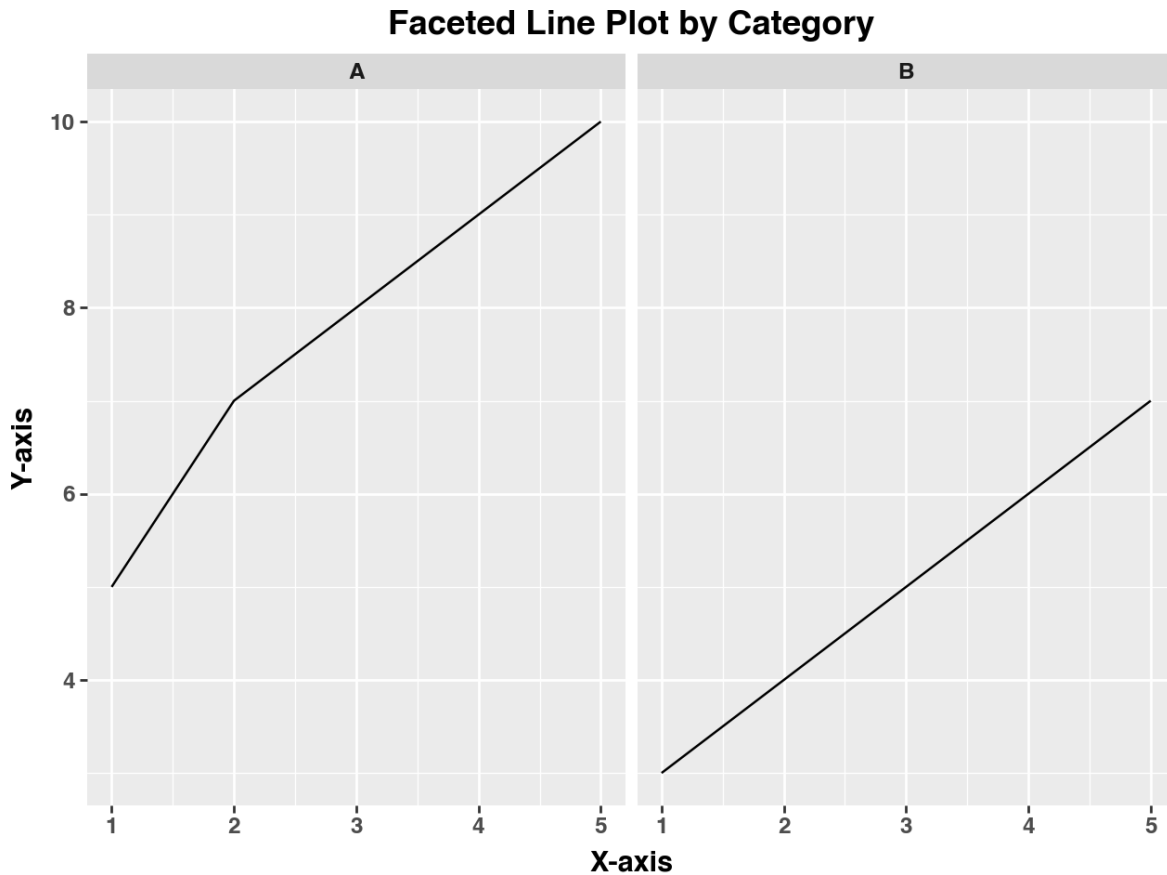
Faceting in `ggplot` allows you to create multiple subplots within a single plot, enabling easy comparison of different subsets of the data. Faceting can be done by rows, columns, or both.

Example: Faceting with a Categorical Variable

```
from plotnine import facet_wrap

# Creating a dataset for faceting
facet_data = pd.DataFrame({
    'x': [1, 2, 3, 4, 5, 1, 2, 3, 4, 5],
    'y': [5, 7, 8, 9, 10, 3, 4, 5, 6, 7],
    'category': ['A', 'A', 'A', 'A', 'A', 'B', 'B', 'B', 'B', 'B']
})

# Facet plot with ggplot
(ggplot(facet_data, aes(x='x', y='y')) +
 geom_line() +
 facet_wrap('~category') +
 ggtitle('Faceted Line Plot by Category') +
 xlab('X-axis') +
 ylab('Y-axis'))
```



This example uses `facet_wrap()` to create separate plots for each category in the dataset, allowing for a clear visual comparison between the groups.

15.3.6 Customizing Aesthetics and Themes

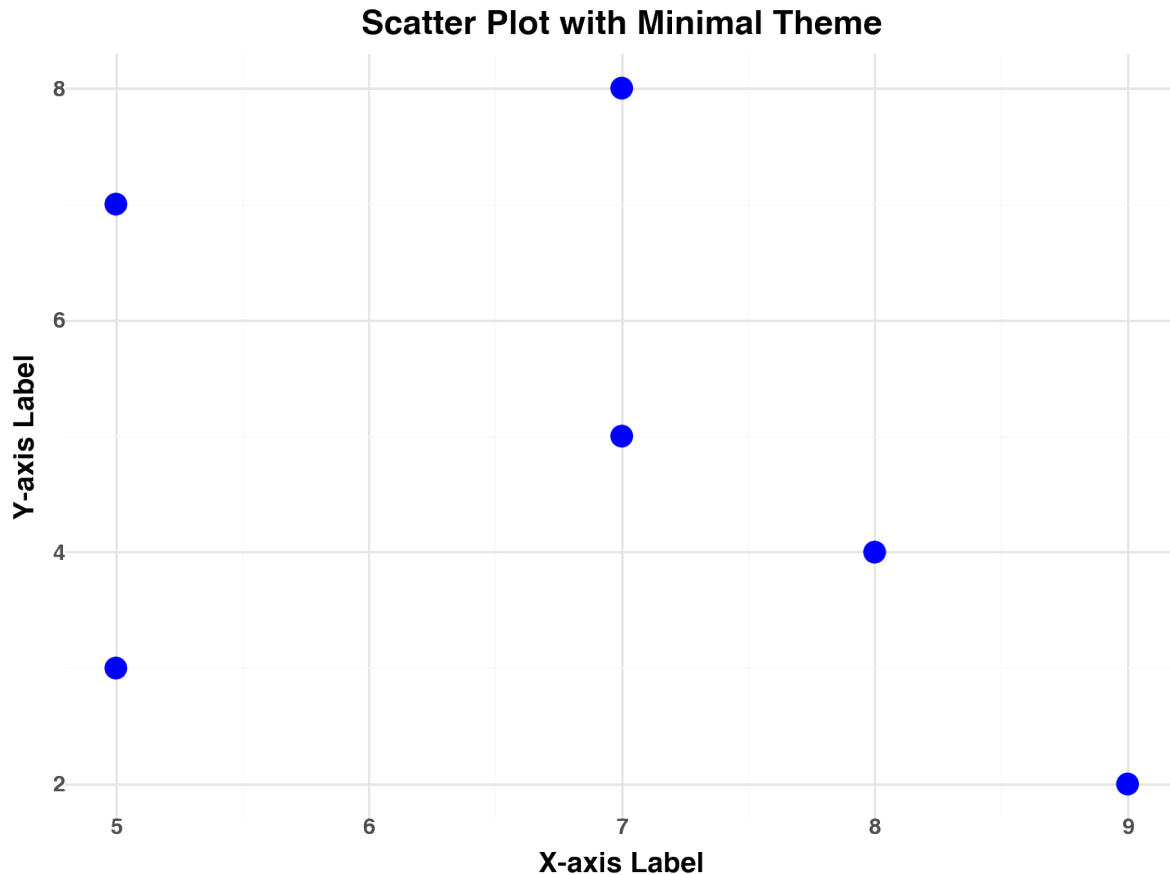
Customization in `ggplot` extends beyond simple color and size adjustments. The library allows for fine-tuning the plot's appearance through themes, which can dramatically change the look and feel of your visualizations.

Example: Using a Different Theme

```
from plotnine import theme_minimal, theme_bw

# Applying a minimal theme to a scatter plot
(ggplot(scatter_data, aes(x='x', y='y')) +
 geom_point(color='blue', size=4) +
```

```
ggtitle('Scatter Plot with Minimal Theme') +  
xlab('X-axis Label') +  
ylab('Y-axis Label') +  
theme_minimal()
```



In this example, the `theme_minimal()` function is applied to give the plot a clean and modern look, with minimal grid lines and axis labels.

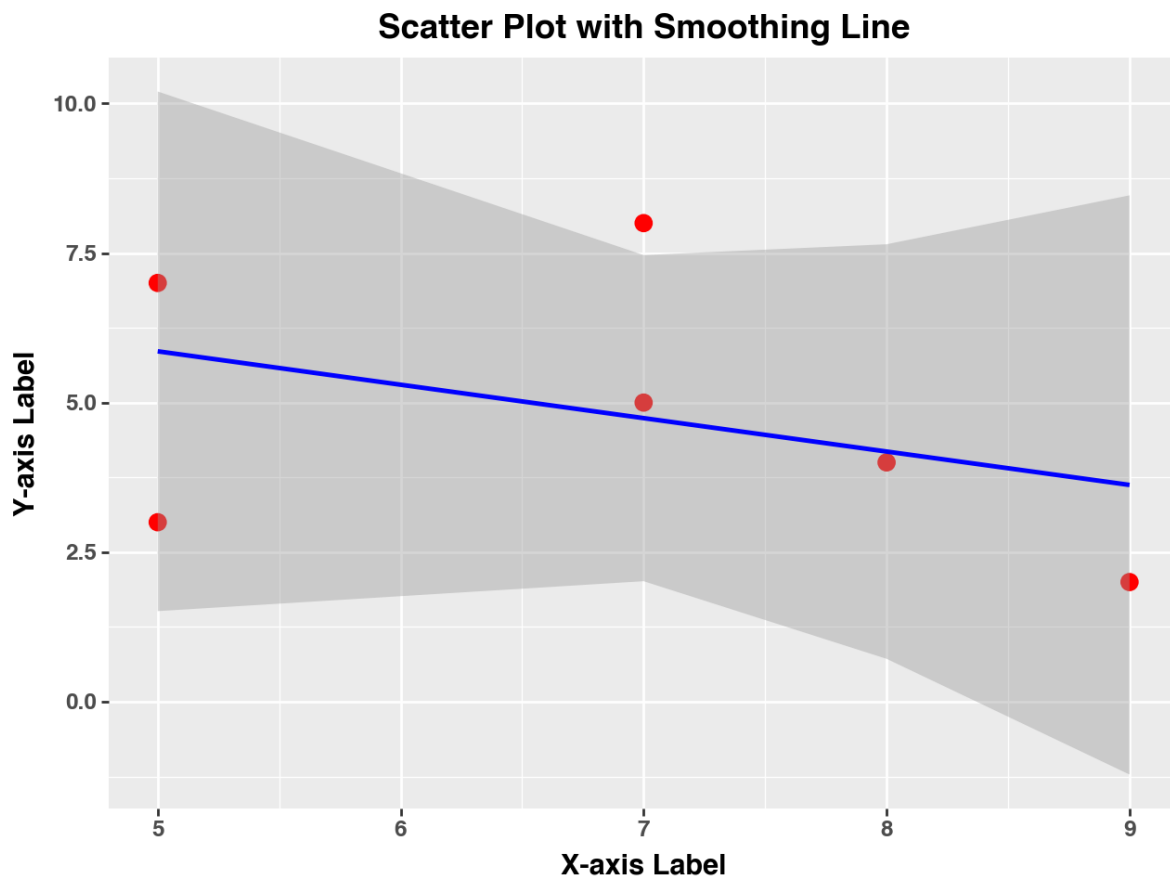
15.3.7 Statistical Transformations

`ggplot` in Python supports built-in statistical transformations to highlight trends and patterns within the data. These transformations can automatically compute summaries like smoothing lines or histograms.

Example: Adding a Smoothing Line to a Scatter Plot

```
from plotnine import geom_smooth

# Scatter plot with a smoothing line
(ggplot(scatter_data, aes(x='x', y='y')) +
 geom_point(color='red', size=3) +
 geom_smooth(method='lm', color='blue') +
 ggtitle('Scatter Plot with Smoothing Line') +
 xlab('X-axis Label') +
 ylab('Y-axis Label'))
```



Here, `geom_smooth()` adds a linear regression line to the scatter plot, showing the trend in the data points.

15.3.8 Advantages of Using ggplot

- **Consistency:** The grammar of graphics approach ensures that the process of building plots is systematic and repeatable.
- **Customization:** Every aspect of the plot can be adjusted to suit specific needs or aesthetic preferences.
- **Layered Design:** Allows for easy addition and modification of plot components without altering the underlying structure.
- **Built-In Statistical Tools:** Supports automatic statistical transformations to help identify patterns in the data.

15.3.9 Limitations and Considerations

While `ggplot` offers great flexibility and aesthetic appeal, there are some considerations to keep in mind:

- **Performance:** Rendering complex visualizations with large datasets can be slower compared to other libraries.
- **Learning Curve:** Requires a solid understanding of the grammar of graphics concepts, which might be challenging for beginners.
- **Dependencies:** As a port of `ggplot2` from R, some features might differ or be less developed than in the R version.

15.4 Exercises

15.4.0.1 Exercise 1: Working with Matplotlib - Creating Basic Line Plots

- Using Matplotlib, create a line plot of the function $f(x) = x^2$ for x values ranging from -10 to 10. Customize the plot by adding labels to the axes, a title, and a grid.
- Modify your plot to change the line color to red and use a dashed line style. Add circular markers to each data point.

15.4.0.2 Exercise 2: Working with Matplotlib - Scatter Plot Customization

- Generate a scatter plot using the following data:
 - $x = [2, 4, 6, 8, 10]$
 - $y = [1, 4, 9, 16, 25]$
- Customize the scatter plot by using triangle markers, coloring the points green, and increasing the marker size. Label the axes and add a title to the plot.

15.4.0.3 Exercise 3: Working with Matplotlib - Creating Subplots

- Create a figure with two subplots arranged vertically:
 - The first subplot should be a line plot of $f(x) = x^3$ for x values from -5 to 5.
 - The second subplot should be a scatter plot of the points $(-3, -27)$, $(-2, -8)$, $(0, 0)$, $(2, 8)$, and $(3, 27)$.
- Ensure that each subplot has a title and labeled axes, and adjust the spacing between the plots for better readability.

15.4.0.4 Exercise 4: Advanced Data Visualization with Seaborn - Analyzing Data Distributions

- Load the built-in `tips` dataset from Seaborn. Create a histogram of the `total_bill` variable with a KDE (Kernel Density Estimate) overlay. Customize the color of the histogram and KDE line.
- Interpret the resulting plot, describing any noticeable patterns in the distribution of the `total_bill` values.

15.4.0.5 Exercise 5: Advanced Data Visualization with Seaborn - Comparing Categorical Data

- Using the `tips` dataset, create a box plot to visualize the distribution of tips received by day of the week. Differentiate between smokers and non-smokers using the `hue` parameter.
- Analyze the box plot to determine on which day the highest median tip amount is given and whether smokers tend to tip more than non-smokers.

15.4.0.6 Exercise 6: Advanced Data Visualization with Seaborn - Correlation Analysis with Heatmaps

- Generate a heatmap showing the correlation matrix of the numerical variables in the `tips` dataset. Use the `coolwarm` color palette and display the correlation values on the heatmap.
- Based on the heatmap, identify which two variables have the strongest correlation and describe their relationship.

15.4.0.7 Exercise 7: Exploring ggplot with plotnine - Basic Line Plot with ggplot

- Using the `plotnine` package, create a basic line plot of the function $g(x) = \sin(x)$ for x values ranging from 0 to 2π . Add appropriate labels to the axes and a title to the plot.
- Enhance the plot by adding points at each integer value of x with different color markers.

15.4.0.8 Exercise 8: Exploring ggplot with plotnine - Creating Faceted Plots

- Create a dataset with two groups, “A” and “B”, and plot a faceted line plot using `ggplot`. Plot the following data points:
 - Group A: (1, 2), (2, 4), (3, 8), (4, 16)
 - Group B: (1, 1), (2, 2), (3, 3), (4, 4)
- Use facetting to display each group’s plot in a separate subplot and ensure that both plots share the same x-axis and y-axis labels.

15.4.0.9 Exercise 9: Exploring ggplot with plotnine - Applying Themes and Aesthetic Modifications

- Create a scatter plot using `ggplot` with a dataset of your choice. Apply the `theme_minimal` style and modify the aesthetics by changing the point color, size, and adding a smooth line to represent the trend in the data.
- Save your plot as a PNG file with a resolution of 300 dpi.

16 Object-Oriented Programming in Python

16.1 Introduction to Object-Oriented Programming

Object-Oriented Programming (OOP) is a programming paradigm that models real-world entities using objects and classes. While procedural programming focuses on functions and the sequence of operations, OOP centers around organizing code into objects, which encapsulate both data and behavior. This approach makes it easier to manage, extend, and modify large and complex programs, while also allowing for better reuse of code.

To understand OOP at a conceptual level, it can be helpful to look at examples from everyday life. These examples will illustrate the key concepts of OOP: **classes**, **objects**, **attributes**, and **methods**. Later in this chapter, we will see how these ideas map onto programming.

16.1.1 Objects and Classes in Everyday Life

Imagine you're organizing a car dealership. Every car on the lot can be seen as an **object**—a specific instance of a general concept, such as a car. The general concept is what we refer to as a **class** in OOP. A class acts like a blueprint or template for creating objects. Each car on the lot is unique in some way, but they all share certain common traits and behaviors, just like objects created from a class.

Let's break this down:

- **Class:** A car can be described by a blueprint, which details its general characteristics (such as make, model, engine type, etc.) and behaviors (like starting the engine, honking the horn). In OOP, this blueprint is the **class**.
- **Object:** Each specific car you have on the lot—a 2021 Toyota Camry, a 2020 Honda Civic, etc.—is an **object** created from the class blueprint. These cars have specific values for their characteristics (for example, one car is blue, another has a sunroof).

This distinction between class and object is fundamental to OOP. A **class** provides the general description, while an **object** is a specific instance created from that description.

16.1.2 Attributes and Methods

Objects have **attributes** and **methods**. Attributes are the data that describe the state of an object, while methods are the actions or behaviors an object can perform.

Continuing with our car dealership analogy:

- **Attributes:** Each car has specific properties like color, make, model, year, and mileage. These properties are the **attributes** of the car object.
- **Methods:** The behaviors of a car, such as starting the engine, honking the horn, or accelerating, are the **methods**. In the class blueprint, methods describe what a car can do, and in OOP terms, they are the actions or functions that can be performed on or by an object.

Just as a car's attributes are set when the car is manufactured, an object's attributes are defined when the object is created from the class.

16.1.2.1 Example: OOP in a Library System

Consider a library system, which can also be modeled using OOP principles. In this case, **books** are objects, and **LibraryBook** might be the class that defines their shared characteristics and behaviors.

- **Class (LibraryBook):** The class defines the general properties of a book, such as its title, author, publication date, and ISBN. It also defines behaviors, such as checking the book out or returning it.
- **Objects:** Each specific book in the library—like “To Kill a Mockingbird” or “Pride and Prejudice”—is an object created from the **LibraryBook** class. Each object has specific values for its attributes (e.g., title, author) but shares the same behaviors (e.g., can be checked out or returned).

When you borrow a book from the library, the system doesn't just perform random tasks; it interacts with the specific book object, checking if it is available or already checked out. This interaction between the object and its methods is one of the strengths of OOP—it provides a structured and organized way to manage complex systems.

16.1.3 Encapsulation: Keeping Attributes and Methods Together

One of the key advantages of OOP is **encapsulation**—the concept of bundling data (attributes) and behaviors (methods) together in one entity, the object. This allows for better organization and modularity.

Let's return to our car example: You can think of each car as a self-contained object. If you want to interact with the car, you don't need to know every intricate detail about its inner workings; you just need to use the car's methods. For instance, to start the car, you don't have to understand how the engine works—you just turn the key or press a button, and the car performs the appropriate behavior.

In programming, encapsulation works in a similar way. When a class is designed, its internal workings (the code and data) are hidden from other parts of the program. Instead, the class provides a clear interface (methods) that can be used to interact with its objects. This hiding of internal details makes the system easier to maintain and modify, because changes to the internal structure of a class won't affect the rest of the program as long as the interface remains the same.

16.1.3.1 Real-World Example: A Bank Account System

Consider a **BankAccount** class in an online banking system. This class defines the shared characteristics of all bank accounts (e.g., account number, balance) and the behaviors that all accounts can perform (e.g., deposit money, withdraw money, check balance).

- **Class (BankAccount):** This acts as the blueprint for creating specific bank accounts. It defines that each account will have an account number, a balance, and methods to deposit or withdraw funds.
- **Objects:** Each customer's bank account is an object created from the BankAccount class. For example, one object might represent Alice's bank account, while another represents Bob's bank account.

The concept of encapsulation is crucial here. You, as the user of the bank account, don't need to know how the bank stores your balance or how it processes withdrawals internally. All you need to know is that when you call the `deposit()` method, money is added to your account, and when you call the `withdraw()` method, money is subtracted. The internal details (data storage, security checks) are hidden from you.

This separation of concerns allows bank developers to modify the internal workings of the bank system without affecting the customers' interactions. For example, the bank might switch to a more efficient algorithm for handling transactions, but as long as the deposit and withdraw methods work the same way, the change is invisible to the user.

16.1.4 Why Use OOP?

Now that we've seen how OOP relates to real-world examples, let's summarize why this approach is so beneficial in programming:

1. **Modularity and Reusability:** By encapsulating data and behavior within classes, OOP promotes modular code that can be reused across different programs. For instance, a class written to represent a “student” in a school system can be reused in other systems (e.g., a library system).
2. **Ease of Maintenance:** Because objects are self-contained and interact with the rest of the system through defined methods, changes to one part of the system are less likely to break the entire program. This makes OOP programs easier to maintain and scale.
3. **Flexibility Through Inheritance:** OOP allows classes to be extended and specialized through inheritance, a powerful feature that we will discuss in detail later. This allows for the creation of complex systems that can be easily modified and extended.
4. **Real-World Modeling:** OOP allows programmers to model real-world systems in a way that’s intuitive and easy to understand. Objects in OOP can directly correspond to real-world entities, making the design and development process more natural.

In the next section, we will explore how to define a class in Python, bringing these concepts into the programming world and showing how to create and work with objects.

16.2 Defining a Class

Now that we’ve explored the key concepts behind Object-Oriented Programming (OOP) and how they relate to everyday examples, it’s time to move into the practical side: defining a class in Python. A class is the fundamental building block of OOP, serving as the blueprint for creating objects. Each object created from a class shares the structure and behavior defined by that class, but with unique values for its attributes.

In this section, we will dive into the syntax of class definition in Python, demonstrate how to create classes, and explain how attributes and methods are defined within a class. As you go through this section, keep in mind the real-world analogies from the previous section to understand how Python classes mimic real-world entities.

16.2.1 Basic Class Syntax

In Python, defining a class is straightforward. The keyword `class` is used, followed by the class name and a colon. The body of the class contains the attributes (data) and methods (functions) that define the behavior of objects created from the class.

Here’s a simple example:

```
class Car:
    pass
```

This is the simplest possible class definition: a class named `Car` with no attributes or methods. The keyword `pass` is used as a placeholder, indicating that there is no functionality yet. We will build on this example to make it more meaningful.

16.2.2 The `__init__` Method: Class Constructor

In most cases, you'll want your class to initialize some values when an object is created. This is where the `__init__` method comes into play. The `__init__` method is known as the **constructor**, and it is automatically called when an object is instantiated from the class. It is used to set up initial values for the object's attributes.

Let's expand on our `Car` class:

```
class Car:
    def __init__(self, make, model, year):
        self.make = make
        self.model = model
        self.year = year
```

Here's what's happening:

- The `__init__` method takes several arguments: `self`, `make`, `model`, and `year`.
- `self` refers to the specific object being created. It allows you to assign the values to the object's attributes.
- Inside the method, `self.make`, `self.model`, and `self.year` refer to the attributes of the object. We assign the values passed in (`make`, `model`, and `year`) to those attributes.

This way, each object (or instance) of the `Car` class will have its own `make`, `model`, and `year` attributes, which are set when the object is created.

16.2.2.1 Example: Creating Objects from a Class

Let's create some `Car` objects using this class:

```
car1 = Car("Toyota", "Camry", 2021)
car2 = Car("Honda", "Civic", 2020)

print(car1.make)
print(car2.year)
```

```
Toyota
2020
```

Here, `car1` and `car2` are objects created from the `Car` class. Each object has its own set of attributes: `car1` is a Toyota Camry from 2021, and `car2` is a Honda Civic from 2020. We can access these attributes using dot notation (`car1.make`, `car2.year`).

This structure mirrors real-world examples: just as each car in a dealership has specific properties, each object in OOP has specific values for its attributes, even though they are all created from the same class blueprint.

16.3 Methods and Attributes in a Class

In addition to attributes, a class can define **methods**—functions that describe the behaviors or actions the objects can perform. Methods in a class always take `self` as the first argument, which refers to the instance on which the method is being called.

Let's add a method to the `Car` class that allows the car to display its full description:

```
class Car:
    def __init__(self, make, model, year):
        self.make = make
        self.model = model
        self.year = year

    def description(self):
        return f"{self.year} {self.make} {self.model}"
```

In this example, the method `description` returns a formatted string with the car's year, make, and model.

```
car1 = Car("Toyota", "Camry", 2021)
print(car1.description())
```

2021 Toyota Camry

Here, we define the method `description` inside the class. When we call `car1.description()`, the car provides its full description. This method operates on the instance of the class (e.g., `car1`), and it has access to the attributes via `self`.

16.3.1 Instance Attributes vs. Class Attributes

In Python, attributes can be classified into two types: **instance attributes** and **class attributes**.

- **Instance Attributes:** These are specific to each object and are defined inside the `__init__` method. For example, in the `Car` class, `make`, `model`, and `year` are instance attributes. Each car object has its own unique values for these attributes.
- **Class Attributes:** These are shared across all instances of a class. They are defined directly within the class but outside of any method. All objects of the class will share the same value for a class attribute, unless it is explicitly overridden in an instance.

Let's modify the `Car` class to include a class attribute:

```
class Car:
    wheels = 4 # Class attribute

    def __init__(self, make, model, year):
        self.make = make
        self.model = model
        self.year = year

    def description(self):
        return f"{self.year} {self.make} {self.model}"
```

In this case, `wheels` is a class attribute, meaning that all cars created from the `Car` class will have 4 wheels. This attribute is shared among all instances of the class:

```
car1 = Car("Toyota", "Camry", 2021)
car2 = Car("Honda", "Civic", 2020)

print(car1.wheels)
print(car2.wheels)
```

4
4

Class attributes can be useful for defining properties that are the same across all instances, such as the number of wheels for cars.

16.3.2 Modifying Attributes

It is possible to modify both instance and class attributes after an object has been created. Let's see how this works.

16.3.2.1 Modifying Instance Attributes

```
car1 = Car("Toyota", "Camry", 2021)
print(car1.description())

# Modifying the year
car1.year = 2022
print(car1.description())
```

2021 Toyota Camry

2022 Toyota Camry

Here, we update the `year` attribute for `car1`, and the object reflects the new value when we call the `description` method again.

16.3.2.2 Modifying Class Attributes

Class attributes can be modified directly using the class name. Any change to a class attribute will be reflected across all instances of the class:

```
Car.wheels = 3
print(car1.wheels)
```

3

By modifying the `wheels` class attribute, we've updated the number of wheels for all `Car` objects.

16.4 Inheritance

Inheritance is one of the most powerful features of Object-Oriented Programming (OOP). It allows a new class to inherit the properties and behaviors (attributes and methods) of an existing class. The new class, known as the **subclass** or **child class**, can also have additional attributes and methods or override existing ones from the parent class. This promotes code reuse and makes it easier to build complex systems by extending functionality rather than starting from scratch.

To understand inheritance conceptually, let's once again consider some real-world examples.

16.4.1 Inheritance in Everyday Life

Imagine a university system where you have **students** and **graduate students**. Both graduate and undergraduate students share many characteristics: they have a name, student ID, and GPA. However, graduate students also have some unique attributes, such as a thesis title and an advisor. It would be redundant to create separate classes for students and graduate students from scratch because many of their attributes overlap.

Instead, we can create a general **Student** class and a more specialized **GraduateStudent** class. The GraduateStudent class can inherit the common attributes and behaviors from the Student class, while adding its own unique features. This hierarchical structure is a natural fit for OOP and is made possible through inheritance.

16.4.2 Defining Inheritance in Python

In Python, inheritance is defined by specifying the parent class in parentheses when defining the child class. The child class inherits all the attributes and methods from the parent class.

Let's define a basic example:

```
class Student:
    def __init__(self, name, student_id, gpa):
        self.name = name
        self.student_id = student_id
        self.gpa = gpa

    def display_info(self):
        return f"Student: {self.name}, ID: {self.student_id}, GPA: {self.gpa}"
```


Here, the `Student` class defines three attributes (`name`, `student_id`, and `gpa`) and one method (`display_info`). Now, we will define a subclass, `GraduateStudent`, which will inherit from `Student`.

```
class GraduateStudent(Student):
    def __init__(self, name, student_id, gpa, thesis_title):
        super().__init__(name, student_id, gpa)
        self.thesis_title = thesis_title

    def display_thesis(self):
        return f"Thesis: {self.thesis_title}"
```

In this example:

- The `GraduateStudent` class inherits from `Student`. We can see this by the syntax `class GraduateStudent(Student)`.
- The `__init__` method of the `GraduateStudent` class uses `super().__init__()` to call the parent class's constructor, inheriting the initialization of `name`, `student_id`, and `gpa`.
- A new attribute, `thesis_title`, is added, along with a method `display_thesis()` to show the thesis title.

16.4.2.1 Example: Creating Instances of Subclasses

Now let's create instances of the `Student` and `GraduateStudent` classes to see how inheritance works in practice:

```
# Creating a Student object
student1 = Student("Alice", "S12345", 3.9)
print(student1.display_info())

# Creating a GraduateStudent object
grad_student1 = GraduateStudent("Bob", "G54321", 4.0, "Quantum Computing")
print(grad_student1.display_info())

print(grad_student1.display_thesis())
```

```
Student: Alice, ID: S12345, GPA: 3.9
Student: Bob, ID: G54321, GPA: 4.0
Thesis: Quantum Computing
```

As shown, `GraduateStudent` inherits the `display_info()` method from the `Student` class, and we can also use the `display_thesis()` method that is specific to `GraduateStudent`.

16.4.3 Overriding Methods

One of the most useful aspects of inheritance is the ability to override methods from the parent class in the child class. This allows the child class to modify or extend the behavior of a method without affecting the parent class.

Let's say we want the `GraduateStudent` class to display more detailed information about the student, including the thesis title, when calling the `display_info()` method. We can override this method in the `GraduateStudent` class:

```
class GraduateStudent(Student):
    def __init__(self, name, student_id, gpa, thesis_title):
        super().__init__(name, student_id, gpa)
        self.thesis_title = thesis_title

    # Overriding the display_info method
    def display_info(self):
        return f"Graduate Student: {self.name}, ID: {self.student_id}, GPA: {self.gpa}, Th
```

Now, when `display_info()` is called on a `GraduateStudent` object, it uses the overridden version of the method instead of the parent class's version:

```
grad_student1 = GraduateStudent("Bob", "G54321", 4.0, "Quantum Computing")
print(grad_student1.display_info())
```

Graduate Student: Bob, ID: G54321, GPA: 4.0, Thesis: Quantum Computing

In this case, we have enhanced the method from the parent class to include additional information about the student's thesis.

16.4.4 The `super()` Function

The `super()` function is crucial when working with inheritance. It allows you to call methods from the parent class, which is especially useful when overriding methods. As we saw in the example above, `super()` was used to call the parent class's `__init__()` method, so we didn't have to duplicate the code that initializes the common attributes.

Here's a breakdown of when to use `super()`:

- **Constructor Chaining:** When defining a child class, you often need to initialize the attributes of the parent class. Using `super()` ensures that the parent class's `__init__()` method is called automatically, making it easier to maintain your code.

- **Overriding Methods:** If a method in the child class overrides a method in the parent class but you still need access to the parent class's version of the method, `super()` allows you to call it.

```
class GraduateStudent(Student):
    def __init__(self, name, student_id, gpa, thesis_title):
        super().__init__(name, student_id, gpa)
        self.thesis_title = thesis_title

    def display_info(self):
        return super().display_info() + f", Thesis: {self.thesis_title}"
```

In this case, we use `super().display_info()` to call the parent class's method and then add additional information about the thesis. This keeps the code for displaying the basic student information centralized in the `Student` class, avoiding code duplication.

16.4.5 Multiple Inheritance

Python supports **multiple inheritance**, which allows a class to inherit from more than one parent class. While this can be a powerful tool, it should be used with caution because it can lead to more complex and harder-to-maintain code. Let's look at an example:

```
class Athlete:
    def __init__(self, sport):
        self.sport = sport

    def show_sport(self):
        return f"Sport: {self.sport}"

class StudentAthlete(Student, Athlete):
    def __init__(self, name, student_id, gpa, sport):
        Student.__init__(self, name, student_id, gpa)
        Athlete.__init__(self, sport)

    def display_info(self):
        return f"{self.name}, ID: {self.student_id}, GPA: {self.gpa}, Sport: {self.sport}"
```

In this example, the `StudentAthlete` class inherits from both `Student` and `Athlete`. This allows the student-athlete to have both academic information (name, ID, GPA) and athletic information (sport). However, notice that we have to explicitly call both `Student.__init__()` and `Athlete.__init__()` because both parent classes need to be initialized.

```
student_athlete1 = StudentAthlete("Charlie", "S98765", 3.8, "Basketball")
print(student_athlete1.display_info())
```

Charlie, ID: S98765, GPA: 3.8, Sport: Basketball

In Python, multiple inheritance is resolved using a method resolution order (MRO), which determines the order in which classes are checked when calling a method. Python uses a depth-first, left-to-right search through the inheritance chain, ensuring that the proper methods are called.

16.5 Encapsulation and Data Hiding

Encapsulation is a core concept in Object-Oriented Programming (OOP) that refers to bundling data (attributes) and methods (functions) that operate on that data into a single unit called an object. This not only helps with organization but also protects the internal state of an object from unintended interference or misuse. Encapsulation allows objects to expose only necessary information to the outside world while hiding their internal workings.

This idea is similar to how most real-world objects work: you interact with them in simple ways, without needing to know all the underlying mechanics. Let's explore this further with examples, and then move on to the specific programming techniques used to implement encapsulation in Python.

16.5.1 Data Hiding: The Foundation of Encapsulation

In programming, **data hiding** is the practice of restricting access to certain attributes and methods within an object. Data hiding is a key aspect of encapsulation, as it ensures that an object's internal state cannot be altered in unintended or unpredictable ways by external code. In Python, we achieve data hiding by using certain naming conventions.

- **Public attributes:** These attributes and methods are accessible from outside the class and can be used or modified freely. By default, all attributes in Python are public unless specified otherwise.
- **Private attributes:** Private attributes are intended to be hidden from outside the class. In Python, we make an attribute private by prefixing its name with a double underscore (`__`).

Let's see how this works in practice with a simple example.

16.5.1.1 Example: Bank Account with Private Balance

```
class BankAccount:
    def __init__(self, account_number, initial_balance):
        self.account_number = account_number # Public attribute
        self.__balance = initial_balance # Private attribute

    def deposit(self, amount):
        if amount > 0:
            self.__balance += amount
        else:
            raise ValueError("Deposit amount must be positive")

    def withdraw(self, amount):
        if 0 < amount <= self.__balance:
            self.__balance -= amount
        else:
            raise ValueError("Insufficient funds or invalid amount")

    def get_balance(self):
        return self.__balance
```

In this example: - `account_number` is a public attribute, meaning anyone can access and modify it. - `__balance` is a private attribute, meaning it cannot be accessed or modified directly from outside the class. - The class provides public methods like `deposit()`, `withdraw()`, and `get_balance()` to interact with the private balance in a controlled way.

Let's see how this works in action:

```
# Creating a bank account
account = BankAccount("12345", 1000)

# Accessing public attribute
print(account.account_number)
```

12345

```
# Accessing private attribute directly (raises AttributeError)
print(account.__balance) # Raises: AttributeError: 'BankAccount' object has no attribute
```

```

# Accessing balance through the public method
print(account.get_balance())

# Depositing money
account.deposit(500)
print(account.get_balance())

# Withdrawing money
account.withdraw(200)
print(account.get_balance())

```

```

1000
1500
1300

```

As you can see, attempting to directly access `__balance` raises an error, demonstrating that the attribute is hidden from outside access. Instead, users of the `BankAccount` class must interact with the balance through the `deposit()`, `withdraw()`, and `get_balance()` methods. This ensures that the balance is only modified in valid ways, preserving the integrity of the object.

16.5.2 Getter and Setter Methods

Encapsulation does not mean that the data inside an object is completely inaccessible. Often, we need a way to safely access or modify private attributes. In such cases, classes can provide **getter** and **setter** methods to expose specific attributes to the outside world, while still maintaining control over how they are accessed or modified.

Let's modify our `BankAccount` class to include a setter for the balance attribute:

```

class BankAccount:
    def __init__(self, account_number, initial_balance):
        self.account_number = account_number
        self.__balance = initial_balance

    def deposit(self, amount):
        if amount > 0:
            self.__balance += amount
        else:
            raise ValueError("Deposit amount must be positive")

```

```

def withdraw(self, amount):
    if 0 < amount <= self.__balance:
        self.__balance -= amount
    else:
        raise ValueError("Insufficient funds or invalid amount")

# Getter for balance
def get_balance(self):
    return self.__balance

# Setter for balance (with validation)
def set_balance(self, amount):
    if amount >= 0:
        self.__balance = amount
    else:
        raise ValueError("Balance cannot be negative")

```

Now we have added a `set_balance()` method to allow controlled modification of the private `__balance` attribute. The setter ensures that the balance cannot be set to a negative value, preserving the consistency of the account's data.

```

# Setting the balance directly using the setter method
account.set_balance(2000)

# Trying to set a negative balance (raises ValueError)
account.set_balance(-500) # Raises: ValueError: Balance cannot be negative

```

By using getter and setter methods, we provide a controlled interface for accessing and modifying private data. This helps prevent improper use of an object's data, enforcing rules and validation where necessary.

16.6 Polymorphism

Polymorphism is another fundamental concept in Object-Oriented Programming (OOP). The term “polymorphism” comes from the Greek words *poly* (meaning “many”) and *morph* (meaning “forms”). In programming, polymorphism allows objects of different classes to be treated as objects of a common parent class, while still preserving their individual behaviors. The primary benefit of polymorphism is that it enables the same operation to behave differently on different types of objects.

Polymorphism is essential in creating flexible and reusable code. By allowing methods to take many forms, polymorphism makes it possible to design systems where the exact behavior of an object can vary depending on the object's class, yet all objects share a common interface. This allows for a more generalized and powerful approach to programming.

16.6.1 Real-World Example: Payment Systems

Imagine a payment processing system for an online store. The store accepts multiple types of payment: credit cards, PayPal, and gift cards. Even though each payment method is different, they all share the same core behavior: they allow customers to make payments. In this scenario, you can create a common interface (or parent class) called `Payment`, and then create specific subclasses like `CreditCardPayment`, `PayPalPayment`, and `GiftCardPayment`.

Each subclass implements its own version of a `process_payment()` method, but all can be treated as instances of the `Payment` class. This is polymorphism in action: the same method name (`process_payment()`) works on different types of payment objects, each with its own implementation of the behavior.

16.6.2 Polymorphism Through Method Overriding

Polymorphism in OOP is often achieved through **method overriding**. When a subclass provides a specific implementation of a method that is already defined in its parent class, this is known as overriding. The method in the subclass replaces the method in the parent class, while still sharing the same interface.

Let's build an example based on our payment processing system.

```
class Payment:
    def process_payment(self, amount):
        raise NotImplementedError("Subclasses must implement this method")

class CreditCardPayment(Payment):
    def process_payment(self, amount):
        return f"Processing credit card payment of {amount}"

class PayPalPayment(Payment):
    def process_payment(self, amount):
        return f"Processing PayPal payment of {amount}"

class GiftCardPayment(Payment):
    def process_payment(self, amount):
        return f"Processing gift card payment of {amount}"
```


In this example: - `Payment` is the parent class, which defines the method `process_payment()` but leaves its implementation to the subclasses. - The subclasses (`CreditCardPayment`, `PayPalPayment`, and `GiftCardPayment`) each override the `process_payment()` method with their specific implementations.

Now, let's see polymorphism in action by creating different payment objects and processing payments through the same interface:

```
payments = [CreditCardPayment(), PayPalPayment(), GiftCardPayment()]

for payment in payments:
    print(payment.process_payment(100))
```

Processing credit card payment of 100

Processing PayPal payment of 100

Processing gift card payment of 100

Even though we have three different types of payment objects, they all share the same interface (`process_payment()`), and we can treat them as instances of the `Payment` class. This makes the code more flexible and easier to extend.

16.6.3 Polymorphism Through Method Overloading (Not Native to Python)

Some programming languages allow **method overloading**, where multiple methods in the same class share the same name but have different parameter lists. While Python doesn't support method overloading natively, polymorphism can still be achieved through other means, such as using default parameters or the `*args` and `**kwargs` syntax to handle variable numbers of arguments.

Here's an example using default parameters to simulate method overloading in Python:

```
class Calculator:
    def add(self, a, b=0, c=0):
        return a + b + c
```

In this example, the `add()` method can be called with one, two, or three arguments, providing flexibility in how the method is used.

```
calc = Calculator()
print(calc.add(5))
print(calc.add(5, 10))
print(calc.add(5, 10, 20))
```

5
15
35

Although Python doesn't support method overloading in the traditional sense, this technique allows you to achieve similar functionality.

16.6.4 Polymorphism in Python with Duck Typing

Python's approach to polymorphism relies heavily on a concept called **duck typing**. Duck typing is based on the principle that “if it looks like a duck and quacks like a duck, it's probably a duck.” In Python, an object's suitability for a given operation is determined by whether it has the necessary attributes and methods, rather than its class hierarchy.

Let's see how duck typing works with an example:

```
class Dog:
    def speak(self):
        return "Woof!"

class Cat:
    def speak(self):
        return "Meow!"

def make_animal_speak(animal):
    print(animal.speak())

# Using different objects with the same method
dog = Dog()
cat = Cat()

make_animal_speak(dog)
make_animal_speak(cat)
```

Woof!
Meow!

In this example, both `Dog` and `Cat` have a `speak()` method, and the function `make_animal_speak()` works with both objects without caring about their specific types. This is a form of polymorphism enabled by duck typing—Python doesn't require a strict type hierarchy as long as the necessary method (`speak()`) is implemented.

16.6.5 Advantages of Polymorphism

Polymorphism offers several key benefits that make it a cornerstone of OOP:

1. **Code Reusability:** Polymorphism allows you to write code that works with objects of different types but share a common interface. This reduces duplication and promotes reusability.
2. **Extensibility:** When new subclasses are added, they automatically work with the existing polymorphic methods without requiring changes to the code that uses the parent class.
3. **Flexibility:** Since polymorphism allows methods to adapt to different object types, you can write more general and flexible programs. This makes it easier to build and extend complex systems.
4. **Readability and Maintainability:** Polymorphism allows you to design cleaner and more understandable code by abstracting common behaviors into a shared interface.

16.7 Exercises

16.7.0.1 Exercise 1: Create a Book Class

Define a `Book` class with the following attributes: - `title` (str) - `author` (str) - `pages` (int) - `year_published` (int)

Implement a method `description()` that returns a string with the book's title, author, and the number of pages. Create a few `Book` objects and print their descriptions.

16.7.0.2 Exercise 2: Car Dealership Simulation

Using the `Car` class example from the chapter, create an inventory system for a car dealership. Define a list of cars with different makes, models, and years. Write a method that searches the inventory and returns all cars from a particular year.

16.7.0.3 Exercise 3: Create a Vehicle Class with Subclasses

Define a parent class `Vehicle` with attributes `make` and `model`. Then, define two subclasses `Truck` and `Motorcycle`, each with additional attributes (`Truck` has a `payload_capacity`, and `Motorcycle` has `cc` for engine capacity). Both subclasses should override a method `vehicle_type()` that prints the specific type of vehicle. Create objects for both subclasses and call their methods.

16.7.0.4 Exercise 4: University System

Expand on the university system example in the chapter by adding a `Professor` class that inherits from `Person`. In addition to the `Person` attributes (e.g., name, age), the `Professor` class should include a `department` and `courses_taught`. Implement a method that prints the professor's name and department, and a method to add a course to the professor's list of courses.

16.7.0.5 Exercise 5: Bank Account with Private Attributes

Create a `BankAccount` class where the balance is a private attribute. Write getter and setter methods to access and modify the balance, ensuring that the balance can never be negative. Write a method `withdraw()` that subtracts a specified amount from the balance and a method `deposit()` to add to the balance. Test these methods to verify the data is encapsulated properly.

16.7.0.6 Exercise 6: Employee Management System

Create a class `Employee` that has private attributes `name` and `salary`. Implement getter and setter methods for the salary that ensure a minimum salary threshold of \$30,000. Write a method to display the employee's name and salary, and test the setter to ensure no invalid salaries are set.

16.7.0.7 Exercise 7: Animal Sound System

Create a parent class `Animal` with a method `speak()`. Then, create subclasses `Dog`, `Cat`, and `Bird`, each overriding the `speak()` method to return a sound appropriate to the animal. Write a function that accepts any `Animal` object and prints the result of calling its `speak()` method.

16.7.0.8 Exercise 8: Shape Class with Polymorphism

Define a parent class `Shape` with a method `area()`. Create two subclasses, `Circle` and `Rectangle`, each overriding the `area()` method to calculate the respective areas. Write a function that takes a list of shapes and prints their areas using polymorphism.

17 Writing Clear and Effective Documentation and PEP 8

17.1 Introduction

Documentation is a vital part of software development, playing a role analogous to that of proofs or derivations in mathematics. It provides the necessary guidance to users and developers on how to understand, maintain, and effectively utilize the code. In the absence of clear documentation, even well-written code can be difficult to interpret, particularly as projects grow in size and complexity.

Well-documented code acts as a **communication tool** between the original developer, collaborators, and future maintainers. It explains not only what the code does but also provides insights into design decisions and trade-offs made during development. This context helps mitigate common issues in collaborative environments—such as misunderstandings, redundancy, and rework—by making expectations and intentions clear.

17.1.1 Types of Documentation

There are multiple levels of documentation that contribute to an effective software development process. These include:

1. **Inline Documentation (Comments):** Provide localized explanations of code sections that are not self-evident.
2. **Docstrings:** A form of structured documentation attached to functions, classes, and modules, serving as a reference for users.
3. **Project-level Documentation:** Includes README files, API references, and manuals, which give users an overview of the system and how to engage with it.

17.1.2 Characteristics of Good Documentation

Good documentation shares several key characteristics:

- **Clarity:** It must be easy to read and understand, without unnecessary jargon or technical complexity.

- **Conciseness:** The documentation should be thorough but not overwhelming—providing the right amount of detail without redundancy.
- **Accuracy:** The information provided must match the behavior of the code. Outdated or incorrect documentation can be more harmful than none at all.
- **Consistency:** Use a consistent structure and style throughout to ensure readability. Following a standard format (like Google or NumPy style for docstrings) makes it easier for developers to engage with the documentation.

17.1.3 Documentation vs. Self-Documenting Code

Although Python encourages readable code, the notion that code can entirely document itself is a myth. While writing “self-documenting code”—that is, code with descriptive names and minimal need for comments—is good practice, **it cannot replace documentation entirely**. Complex algorithms or critical design decisions need explicit explanations.

- **When to Write Comments:** Comments are especially useful when you need to explain *why* a particular approach was chosen or describe non-trivial logic that isn’t immediately apparent from the code.
- **When Not to Use Comments:** Avoid commenting on obvious code—such as `x = x + 1`—where the purpose is clear from context.

17.1.4 The Role of Documentation in Collaborative Projects

In collaborative environments, documentation plays a pivotal role in:

- **Onboarding New Team Members:** Documentation allows new contributors to quickly familiarize themselves with the codebase, tools, and workflows, minimizing onboarding time.
- **Version Control Integration:** Documentation updates should accompany code changes in version control systems (e.g., Git). It is essential to document any changes in behavior to keep users informed.
- **Knowledge Transfer:** In academia or industry, team members may rotate in and out of projects. Documentation ensures continuity by reducing reliance on individuals for specialized knowledge.

For example, consider the following scenario: a researcher working on a complex statistical model shares their code with a team. Without clear documentation of assumptions, data preprocessing steps, and expected outputs, other members may struggle to replicate results or extend the model. Proper documentation ensures the reproducibility and scalability of such collaborative efforts.

17.1.5 The Relationship Between Documentation and Code Quality

Documentation is a reflection of the quality of your code. Projects with clear, well-organized documentation are perceived as more professional and trustworthy. In academic settings, code accompanied by robust documentation fosters reproducibility, a key principle in scientific research. Likewise, industry projects with well-documented APIs and user guides enhance user satisfaction and reduce support overhead.

17.2 Best Practices for Writing Documentation

To write documentation that adds genuine value to your codebase, following established best practices is essential. This section outlines practical strategies to ensure your documentation is clear, concise, and useful to both developers and end-users. By following these guidelines, you can create documentation that evolves seamlessly with your project and remains relevant over time.

17.2.1 Write Meaningful Docstrings

Docstrings are an integral part of Python's documentation strategy. They should clearly explain the *what*, *how*, and *why* of your code. Python's convention is to place a docstring at the beginning of every module, class, and function definition.

Key elements of good docstrings include:

1. **Description:** What the function, class, or module does.
2. **Parameters:** List all input arguments with their expected types.
3. **Return Values:** Indicate what the function returns and the data type.
4. **Raises:** Describe exceptions that might be raised, if any.

Example using Google-style docstring:

```
def calculate_mean(data):  
    """Calculate the mean of a list of numbers.  
  
    Args:  
        data (list of float): A list of numerical values.  
  
    Returns:
```

```

float: The mean of the input list.

Raises:
    ValueError: If the input list is empty.
"""
if len(data) == 0:
    raise ValueError("The input list must not be empty.")
return sum(data) / len(data)

```

This docstring provides a clear overview of the function's behavior, making it easy to understand its purpose and usage.

17.2.2 Comment Strategically

While docstrings describe *what* a function or module does, inline comments provide detailed insights into specific code sections. However, excessive comments can clutter the code, so use them judiciously.

When to Use Comments:

- To explain *why* certain design decisions were made.
- To clarify non-obvious logic or algorithms.
- To highlight potential areas of concern or future changes.

Example of strategic commenting:

```

# Using list comprehension for performance
squared_numbers = [x**2 for x in range(100)]

```

When to Avoid Comments:

- Avoid restating obvious code logic:

```

x = x + 1 # Add 1 to x (unnecessary comment)

```

- Use meaningful variable names to reduce the need for trivial comments.

17.2.3 Keep Documentation Up-to-Date

Outdated documentation can mislead users and developers, creating confusion. Documentation should evolve alongside the code. Consider the following strategies to ensure that your documentation stays relevant:

- **Document changes alongside code updates:** Incorporate documentation updates as part of the development workflow, especially when new features are added or APIs change.
- **Use version control:** Track changes to documentation using Git or another version control system to ensure consistency and allow for rollbacks if needed. We will discuss Git in a later chapter.

17.2.4 Provide Usage Examples

Including usage examples in your documentation helps readers understand how to use your code in practical scenarios. Examples also demonstrate expected inputs and outputs, which aids in faster comprehension.

Example with a function usage guide:

```
# Example usage of calculate_mean()
numbers = [10, 20, 30, 40]
print(calculate_mean(numbers)) # Output: 25.0
```

Usage examples are especially useful in API documentation, where users need quick access to common use cases.

17.2.5 Use Consistent Formatting

Consistency enhances readability. Adopting a standard format, such as **Google** or **NumPy style** docstrings, ensures that your documentation looks uniform throughout the codebase.

Examples of two common formats:

1. Google-style:

```
def foo(a, b):
    """Add two numbers.

    Args:
        a (int): The first number.
        b (int): The second number.

    Returns:
        int: The sum of the two numbers.
    """
    return a + b
```

2. NumPy-style:

```
def foo(a, b):
    """
    Add two numbers.

    Parameters
    -----
    a : int
        The first number.
    b : int
        The second number.

    Returns
    -----
    int
        The sum of the two numbers.
    """
    return a + b
```

Choose a format and apply it consistently across your project to maintain uniformity and reduce confusion.

17.2.6 Automate Documentation Generation

Automation can reduce the effort required to keep documentation consistent and up-to-date. Python offers several tools for generating documentation:

- **Sphinx:** Generates HTML and PDF documentation from docstrings and reStructured-Text files.
- **MkDocs:** A fast, simple tool for generating static websites from Markdown files.
- **pydoc:** A built-in Python tool for generating text-based documentation.

17.2.6.1 Using pydoc for Python Documentation

pydoc is a built-in Python tool that generates documentation for Python modules, classes, functions, and methods directly from their docstrings. It provides a quick way to view documentation either in the terminal or through a simple web interface.

17.2.6.2 Viewing Documentation in the Terminal

You can use `pydoc` from the command line to display documentation about any installed module or function.

Usage Example:

```
pydoc math
```

This command will display the documentation for the `math` module directly in the terminal. You can also use it to look up specific functions:

```
pydoc math.sqrt
```

17.2.6.3 Generating HTML Documentation

To generate HTML documentation for a module or package, use the following command:

```
pydoc -w <module_name>
```

For example:

```
pydoc -w math
```

This will create an HTML file (`math.html`) containing the documentation for the `math` module.

You can also do this within a `.py` script with the command

```
import pydoc
pydoc.writedoc('math')
```

17.2.6.4 Searching for Modules

You can use `pydoc` to search for installed modules on your system:

```
pydoc modules
```

This will list all available modules, helping you discover built-in functionality and installed packages.

17.2.7 Include Error Handling Information

Documenting potential exceptions or error conditions ensures that users can handle unexpected situations effectively. In addition to listing exceptions, explain scenarios where the function might raise them.

Example:

```
def divide(a, b):
    """Divide two numbers.

    Args:
        a (float): Numerator.
        b (float): Denominator.

    Returns:
        float: The result of the division.

    Raises:
        ZeroDivisionError: If the denominator is zero.
    """
    if b == 0:
        raise ZeroDivisionError("Denominator must not be zero.")
    return a / b
```

17.3 Understanding PEP 8

PEP 8—the Python Enhancement Proposal 8—is the official style guide for Python code. It provides guidelines on code formatting to promote consistency, making code easier to read, maintain, and share across projects. Following PEP 8 ensures that your code adheres to widely accepted best practices, fostering collaboration and professionalism. Just as mathematical notation brings clarity to equations, PEP 8 ensures that Python code is both elegant and accessible.

17.3.1 Why PEP 8 Matters

Consistent style throughout a project enhances readability, reduces cognitive load, and minimizes friction in collaborative efforts. Adopting PEP 8 helps teams:

- **Avoid ambiguity** by enforcing clear, logical code structures.
- **Improve code reviews** by focusing on logic rather than formatting issues.

- **Increase maintainability** by ensuring that code written months later remains understandable.

PEP 8 is especially important for open-source projects, where contributors need to align their work with community standards.

17.3.2 Key PEP 8 Guidelines

17.3.2.1 Indentation

- **Use 4 spaces per indentation level.** Avoid using tabs, as mixing tabs and spaces can lead to errors and inconsistencies.

```
def example_function():  
    print("This is an example.")
```

- Tools like **PyCharm** or **VS Code** allow automatic enforcement of this rule.

17.3.2.2 Line Length

- **Limit lines to 79 characters.** For longer code lines, break them across multiple lines using parentheses or backslashes.

```
total = (first_number  
        + second_number  
        + third_number)
```

- For comments and docstrings, the recommended limit is 72 characters.

17.3.2.3 Blank Lines

- Use **two blank lines** between top-level functions or class definitions.
- Use **one blank line** between methods within a class or between sections in a function.

```
def func1():  
    pass  
  
def func2():  
    pass
```

17.3.2.4 Imports

- Place **all imports at the top** of the file.
- Group imports as follows:
 1. **Standard library imports** (e.g., `os`, `sys`).
 2. **Third-party imports** (e.g., `numpy`, `pandas`).
 3. **Local module imports** (e.g., `from my_module import helper`).

Example:

```
import os
import sys
import numpy as np
from my_module import helper_function
```

17.3.2.5 Naming Conventions

- **Variables and functions:** Use lowercase_with_underscores.

```
def calculate_mean(data):
    return sum(data) / len(data)
```

- **Classes:** Use CapWords.

```
class DataFrameHandler:
    pass
```

- **Constants:** Use UPPERCASE_WITH_UNDERSCORES.

```
MAX_CONNECTIONS = 10
```

17.3.2.6 Whitespace Usage

- **Avoid extraneous spaces** around operators, brackets, or commas:

```
x = a + b # Correct
y = (1, 2) # Correct
```

- **Incorrect:**

```
x = a + b
y = ( 1 , 2 )
```

17.3.2.7 Inline Comments and Block Comments

- Use **inline comments** sparingly and only when necessary.

```
x = x * 2 # Doubling the value of x
```

- Use **block comments** for more detailed explanations.

```
# This section of code handles
# file input and error checking.
with open('file.txt', 'r') as f:
    data = f.read()
```

17.3.2.8 Docstring Conventions

- Use **triple double quotes** for all docstrings, even for one-liners:

```
def func():
    """Do nothing."""
    pass
```

- Multi-line docstrings should have a summary line followed by more details:

```
def add(a, b):
    """
    Add two numbers.

    This function takes two integers and returns their sum.
    """
    return a + b
```

17.3.3 Common PEP 8 Pitfalls

1. **Inconsistent indentation:** Mixing tabs and spaces can break the code.
2. **Long lines:** Resist the urge to cram too much logic onto one line.

3. **Improper import ordering:** Be mindful to separate imports logically.
4. **Excessive comments:** Comment only when necessary—clear code is better than verbose comments.

17.3.4 Exceptions to PEP 8

While PEP 8 is a valuable guide, it's not an absolute rule. In certain cases—such as writing complex scientific code or working with third-party libraries—it may be necessary to deviate from PEP 8 for clarity or compatibility. Use discretion when making exceptions, ensuring that the code remains readable and maintainable.

18 Exception Handling and Debugging Techniques

18.1 Introduction to Exception Handling

In computer programming, errors are an inevitable part of development. These errors, known as *exceptions*, occur when a program encounters something unexpected during execution. For example, dividing by zero, accessing a file that doesn't exist, or trying to convert a string to a number can all cause exceptions. If not handled properly, these exceptions can cause a program to crash, resulting in a poor user experience and potential data loss.

Exception handling is the mechanism that allows programs to anticipate and gracefully respond to errors. Rather than terminating abruptly, a program can detect an exception, respond appropriately, and continue functioning. This approach is especially important in scientific computing, statistical modeling, and other data-intensive tasks where robustness and reliability are paramount.

18.1.1 Types of Errors in Python

There are three primary categories of errors that can occur in Python:

1. Syntax Errors

Syntax errors, also known as *parsing errors*, occur when Python cannot interpret the code because it violates the language's syntax rules. These are usually caught before the program runs.

Example:

```
print("Hello"
```

SyntaxError: incomplete input (3469545723.py, line 1)

The missing closing parenthesis will raise a `SyntaxError`.

2. Runtime Errors (Exceptions)

These errors occur during program execution, even if the code is syntactically correct. Runtime errors are what we typically handle using exception handling mechanisms.

Example:

```
x = 10 / 0
```

```
ZeroDivisionError: division by zero
```

3. Logical Errors

Logical errors do not cause the program to crash, but the output is incorrect due to flaws in the logic of the code. These are often the hardest to detect and require thorough testing to identify.

18.1.2 Why Handle Exceptions?

In large-scale applications—such as those used for statistical modeling, data collection, or machine learning—exceptions are common and can arise from user input errors, missing data, or network connectivity issues. Handling exceptions ensures that:

- **Programs Remain Functional:** Rather than crashing, the program can respond to the error and continue executing.
- **Users Receive Feedback:** Users get meaningful error messages, helping them understand and correct their input.
- **Data Integrity is Preserved:** By handling exceptions, the program can prevent incomplete operations that could lead to data corruption.
- **Code is Easier to Maintain:** Exception handling creates structured and predictable error management, making the codebase easier to debug and maintain.

18.1.3 The Life Cycle of an Exception

When an exception occurs, the following sequence of events takes place:

1. **Detection:** An exception is raised when Python encounters an unexpected situation.
2. **Propagation:** If the exception is not handled in the immediate scope, Python searches for a matching handler up the call stack.
3. **Handling:** If a matching handler is found, the program executes the code in the corresponding `except` block.
4. **Termination or Recovery:** Depending on the program's design, it may either terminate gracefully or recover and continue.

18.1.4 Common Exceptions in Python

Below is a list of some frequently encountered exceptions in Python, along with descriptions:

Exception	Cause
<code>ZeroDivisionError</code>	Attempt to divide by zero.
<code>ValueError</code>	Invalid argument passed to a function.
<code>TypeError</code>	Operation applied to an inappropriate type.
<code>FileNotFoundError</code>	File or directory does not exist.
<code>IndexError</code>	Index out of range for a list or tuple.
<code>KeyError</code>	Key not found in a dictionary.
<code>AttributeError</code>	Accessing an attribute that doesn't exist.
<code>IOError</code>	Input/output operation failed (e.g., file read error).
<code>AssertionError</code>	Raised when an assertion fails.

18.1.5 The Philosophy Behind Python's Exception Handling

Python's philosophy for exception handling emphasizes readability, simplicity, and robustness, in line with the language's overarching principle, expressed in the Zen of Python: *"Errors should never pass silently."* By making error handling explicit, Python encourages developers to anticipate and plan for potential problems.

Python also follows the principle of *EAFP* (*Easier to Ask for Forgiveness than Permission*). Rather than checking in advance if an operation will succeed, it is often more efficient to perform the operation and handle any exceptions if they arise. This approach simplifies code while ensuring errors are dealt with properly.

Example of EAFP in practice:

```
try:
    value = int(input("Enter a number: "))
except ValueError:
    print("Invalid input. Please enter a valid integer.")
```

Instead of preemptively validating input, the code attempts to perform the operation and handles any issues via exception handling.

18.1.6 Exceptions in Statistical and Mathematical Computing

In statistical programming, exception handling plays a critical role. Consider the following scenarios:

1. Handling Division by Zero in Calculations

When performing statistical computations, division by zero might occur due to missing or incomplete data. Proper exception handling ensures the program doesn't crash and provides meaningful feedback.

Example:

```
try:
    mean = total_sum / count
except ZeroDivisionError:
    print("Cannot compute mean: division by zero.")
```

2. Handling File Input Errors

Data processing often involves reading data from external files. Exception handling ensures the program gracefully handles missing files or incorrect paths.

Example:

```
try:
    with open("data.csv", "r") as file:
        data = file.read()
except FileNotFoundError:
    print("The data file is missing.")
```

3. Managing User Input in Interactive Applications

Many programs used for data analysis involve user interaction. Ensuring users provide valid input is critical to avoid runtime errors.

Example:

```
while True:
    try:
        n = int(input("Enter the sample size: "))
        break
    except ValueError:
        print("Invalid input. Please enter a number.")
```

Exception handling allows developers to anticipate, detect, and respond to errors effectively, ensuring the program runs smoothly under unexpected conditions. Whether in data analysis or mathematical modeling, properly handling exceptions is essential for developing reliable and user-friendly applications. The next sections will dive into the syntax and practical examples of handling exceptions in Python, providing tools to manage errors and improve program stability.

18.2 The try and except Blocks

The foundation of exception handling in Python lies in the use of `try` and `except` blocks. These blocks allow you to “try” a block of code that might raise an exception and catch that exception to handle it gracefully. The `try` block contains the code that may cause an error, while the `except` block specifies how to handle the error if it occurs.

18.2.1 Basic Structure of try and except

The basic syntax for a `try-except` block is as follows:

```
try:
    # Code that might raise an exception
    result = 10 / 0
except ZeroDivisionError:
    # Code to handle the exception
    print("Cannot divide by zero.")
```

Cannot divide by zero.

In the example above, Python attempts to execute the code in the `try` block. Since dividing by zero is not allowed, a `ZeroDivisionError` is raised, and control is passed to the `except` block, which prints an appropriate message.

18.2.2 Catching Specific Exceptions

You can specify multiple `except` blocks to catch different types of exceptions. This is useful when you want to handle different errors in different ways.

```
try:
    value = int(input("Enter a number: "))
    result = 10 / value
except ValueError:
    print("Invalid input. Please enter a valid number.")
except ZeroDivisionError:
    print("You can't divide by zero!")
```

This example ensures that invalid input and division by zero are handled separately, improving the user experience with specific error messages.

18.2.3 Catching Multiple Exceptions in One Block

When multiple exceptions are expected, you can handle them in a single `except` block by grouping them in a tuple.

```
try:
    value = int(input("Enter a number: "))
    result = 10 / value
except (ValueError, ZeroDivisionError) as e:
    print(f"Error occurred: {e}")
```

This approach simplifies code when the same response is appropriate for different exceptions.

18.2.4 Handling All Exceptions (Not Recommended)

If you do not know what exceptions might occur, you can use a bare `except` block. However, this is generally discouraged because it can catch unexpected exceptions, making debugging more difficult.

```
try:
    result = 10 / value
except:
    print("An error occurred.")
```

A better practice is to catch exceptions explicitly or use `Exception` to capture any standard error while leaving system-exit exceptions unhandled.

```
try:
    result = 10 / value
except Exception as e:
    print(f"An unexpected error occurred: {e}")
```

18.2.5 Using the else Block

Python provides an optional `else` block that runs only if no exceptions are raised in the `try` block. This ensures that the `else` block is executed only when everything runs smoothly.

```
try:
    value = int(input("Enter a number: "))
    result = 10 / value
```

```

except ZeroDivisionError:
    print("You can't divide by zero!")
except ValueError:
    print("Please enter a valid number.")
else:
    print(f"The result is {result}")

```

This example ensures that the message in the `else` block is printed only if the input is valid and no exception occurs.

18.2.6 The finally Block

A `finally` block is always executed, regardless of whether an exception occurs. It is typically used for cleanup operations, such as closing files or releasing resources.

```

try:
    file = open("data.txt", "r")
    data = file.read()
except FileNotFoundError:
    print("The file was not found.")
finally:
    file.close() # Ensures the file is closed

```

Even if an exception occurs, the `finally` block ensures the file is closed, preventing resource leaks.

18.2.7 Nesting try and except Blocks

In more complex scenarios, `try` and `except` blocks can be nested to handle multiple layers of potential exceptions.

```

try:
    file = open("data.txt", "r")
    try:
        data = file.read()
        value = int(data)
    except ValueError:
        print("Data is not a valid integer.")
finally:
    file.close()

```

Here, the inner `try` block handles issues with reading or processing data, while the outer block ensures the file is closed no matter what happens.

18.2.8 Raising Exceptions

Python allows developers to manually raise exceptions using the `raise` statement. This is helpful when you want to enforce certain conditions in your code.

```
def set_age(age):
    if age < 0:
        raise ValueError("Age cannot be negative.")
    print(f"Age is set to {age}")

try:
    set_age(-1)
except ValueError as e:
    print(f"Error: {e}")
```

Error: Age cannot be negative.

In this example, the function raises a `ValueError` if the age is negative. This exception is caught in the `try-except` block to provide meaningful feedback.

The `try` and `except` blocks provide a powerful way to manage exceptions gracefully. By anticipating potential errors and designing structured exception handling, you can build programs that are more robust, maintainable, and user-friendly. The next section will explore best practices for debugging to further enhance code reliability.

18.3 Debugging Techniques

Debugging is the process of identifying, analyzing, and resolving errors or bugs in a program. Even with well-structured code, errors can arise from various sources, such as logical flaws, incorrect assumptions, or misinterpretation of data. Therefore, mastering debugging techniques is essential for writing reliable and efficient programs, particularly in fields like statistical computing and data analysis, where accuracy is paramount.

18.3.1 Using Print Statements for Debugging

One of the simplest and most common debugging techniques is the use of `print()` statements. By printing intermediate values or messages, you can track the flow of your program and observe how variables change over time. This method is especially useful for tracking down logical errors.

Example:

```
def calculate_average(numbers):
    total = sum(numbers)
    print(f"Total: {total}") # Debugging statement
    average = total / len(numbers)
    print(f"Average: {average}") # Debugging statement
    return average

numbers = [10, 20, 30, 40]
calculate_average(numbers)
```

Total: 100

Average: 25.0

25.0

In this example, the `print()` statements help visualize how the total and average values are calculated. This technique is simple but effective for small programs or sections of code.

Limitations:

- It can clutter your code with `print()` statements.
- You must remove or comment out these statements once the bug is fixed.
- It can be inefficient when debugging larger, more complex programs.

18.3.2 Using Python's Built-In Debugger (pdb)

Python provides a built-in debugger, `pdb`, which offers more powerful debugging capabilities. It allows you to set breakpoints, step through your code, and inspect variables at runtime.

To start using `pdb`, you can insert the following line in your code where you want the execution to pause:

```
import pdb; pdb.set_trace()
```

Example:

```
def calculate_average(numbers):  
    total = sum(numbers)  
    import pdb; pdb.set_trace() # Sets a breakpoint here  
    average = total / len(numbers)  
    return average  
  
numbers = [10, 20, 30, 40]  
calculate_average(numbers)
```

When you run the program, it will pause at the breakpoint. You can then use the following commands to debug:

- **n** (next): Execute the next line of code.
- **c** (continue): Continue execution until the next breakpoint.
- **p** (print): Print the value of a variable.
- **q** (quit): Exit the debugger.

Using `pdb`, you can navigate through your program interactively and inspect the state of variables at specific points, making it easier to identify the cause of an issue.

18.3.3 IDE Debugging Tools

Most modern Integrated Development Environments (IDEs), such as PyCharm, VSCode, and RStudio, offer built-in debugging tools that provide a graphical interface for setting breakpoints, stepping through code, and inspecting variables. These tools enhance productivity and streamline the debugging process by making it more intuitive.

In PyCharm, for example, you can:

1. Set breakpoints by clicking next to the line number in the editor.
2. Run the program in debug mode by clicking the “bug” icon.
3. Inspect the value of variables in the “Variables” pane.
4. Step through code using the “Step Into,” “Step Over,” and “Step Out” buttons.

Advantages of IDE Debuggers:

- No need to modify your code with debugging statements like `print()` or `pdb.set_trace()`.
- Easier navigation through complex codebases.
- Visual inspection of variables, call stacks, and execution flow.

18.3.4 Testing and Assertions

Testing is a proactive debugging technique that ensures your code behaves as expected. Writing tests allows you to check the correctness of your functions and detect issues early. Testing frameworks like `unittest` and `pytest` help automate the process, ensuring that any future changes don't introduce new bugs.

Example Using `unittest`:

```
import unittest

def calculate_average(numbers):
    return sum(numbers) / len(numbers)

class TestAverage(unittest.TestCase):
    def test_average(self):
        self.assertEqual(calculate_average([10, 20, 30]), 20)
        self.assertEqual(calculate_average([5, 5, 5]), 5)

if __name__ == "__main__":
    unittest.main()
```

In this example, the `unittest` framework is used to check if the `calculate_average` function returns the expected values. Tests can be run repeatedly to ensure code correctness.

The Purpose of `if __name__ == "__main__":`

In Python, the `if __name__ == "__main__":` construct is used to control the execution of code, ensuring that certain blocks run only when the script is executed directly, and not when it is imported as a module. This is a best practice for writing reusable code and organizing scripts that might serve as both standalone programs and libraries.

18.3.4.1 How It Works

When a Python file is executed, the interpreter sets a special built-in variable, `__name__`. If the script is being run directly, `__name__` is set to `"__main__"`. However, if the script is imported into another module, `__name__` takes the name of the module instead.

18.3.4.2 Example:

```
def greet():
    print("Hello, world!")

if __name__ == "__main__":
    greet()
```

Explanation:

- If this script is run directly (e.g., `python script.py`), the `greet()` function will be called, printing `"Hello, world!"`.
- If the same script is imported into another module, the `greet()` function won't run automatically because the code inside the `if __name__ == "__main__":` block will be skipped.

18.3.4.3 Why It's Useful

1. **Prevents Unintended Execution:** Ensures that code meant for direct execution does not run when the module is imported elsewhere.
2. **Facilitates Code Reusability:** Allows you to reuse functions, classes, or variables in other scripts without triggering the script's main logic.
3. **Supports Testing:** Makes it easy to write code that behaves differently during development, testing, and production.

This construct is an essential part of Python programming, especially when writing scripts that may also serve as importable modules.

18.3.4.4 Assertions

Assertions are a built-in mechanism in Python to enforce assumptions in your code. If an assertion fails, it raises an `AssertionError` and stops the program, allowing you to catch logic errors during development.

Example:

```
def calculate_average(numbers):
    assert len(numbers) > 0, "List of numbers cannot be empty"
    return sum(numbers) / len(numbers)
```

```
# Raises AssertionError if an empty list is passed
```

```
calculate_average([])
```

AssertionError: List of numbers cannot be empty

Assertions provide a simple way to ensure that certain conditions hold true during execution, helping you catch errors early.

18.3.5 Logging for Debugging

Logging is a more advanced form of debugging, particularly useful for larger applications. While `print()` statements are great for quick checks, logging provides a more robust and configurable way to record what is happening in your program.

Python's `logging` module allows you to log messages at different levels of severity, such as `DEBUG`, `INFO`, `WARNING`, `ERROR`, and `CRITICAL`. This makes it easy to record important information about the state of your program while it runs.

Example Using logging:

```
import logging

logging.basicConfig(level=logging.DEBUG)

def calculate_average(numbers):
    logging.debug(f"Calculating average for: {numbers}")
    total = sum(numbers)
    logging.debug(f"Total: {total}")
    average = total / len(numbers)
    logging.debug(f"Average: {average}")
    return average

numbers = [10, 20, 30]
calculate_average(numbers)
```

In this example, `logging.debug()` statements provide a running log of the calculations performed by the program. Unlike `print()` statements, you can easily disable or change the logging level without modifying the code.

Benefits of Logging:

- You can control the amount of output by setting different logging levels.
- Log files can be created to store output for later analysis.
- It helps in debugging production code without interrupting the program's flow.

18.3.6 Writing Clean Code to Minimize Bugs

The best way to reduce debugging time is to write clean, well-structured code from the beginning. Adopting best practices such as following PEP 8 (Python’s style guide), writing meaningful variable names, and breaking complex tasks into smaller functions can drastically reduce the likelihood of bugs.

Some tips include:

- Use descriptive names for variables and functions.
- Write functions that do one thing and do it well.
- Keep your code DRY (Don’t Repeat Yourself).
- Write unit tests to ensure each part of your code works as expected.

Debugging is an essential part of the development process, especially in fields like statistics and data science, where accuracy is critical. From simple print statements to advanced logging techniques, there are various tools available to help identify and resolve issues. By combining these strategies with careful coding practices, you can reduce the number of bugs in your programs and make the debugging process more efficient.

18.4 Exercises

18.4.0.1 Exercise 1: Handling Division by Zero

Write a program that asks the user for two numbers and prints the result of dividing the first number by the second. Use exception handling to manage cases where the second number is zero. Ensure that the program does not crash and provides a helpful error message.

Example Output:

```
Enter the numerator: 10
Enter the denominator: 0
Error: Cannot divide by zero.
```

18.4.0.2 Exercise 2: File Reading with Exception Handling

Write a program that attempts to open a file and read its contents. If the file does not exist, handle the `FileNotFoundError` exception by printing an appropriate message. Add a `finally` block to ensure the file is closed after reading, even if an error occurs.

Example Output:

```
Enter the file name: missing_file.txt
Error: The file was not found.
```

18.4.0.3 Exercise 3: Validating User Input

Create a function that prompts the user to enter a positive integer. Use exception handling to ensure the input is a valid integer and greater than zero. If the user enters invalid input, prompt them to try again until they enter a valid number.

Example Output:

```
Enter a positive integer: -5
Invalid input. Please enter a positive integer.
Enter a positive integer: hello
Invalid input. Please enter a positive integer.
Enter a positive integer: 10
You entered: 10
```

18.4.0.4 Exercise 4: Testing with Assertions

Write a function that calculates the square root of a number, but only for non-negative inputs. Use an assertion to ensure that the input is non-negative, and if the input is negative, raise an `AssertionError`. Write test cases to verify the behavior of the function.

18.4.0.5 Exercise 5: Using pdb for Debugging

Insert a breakpoint in the following code using the pdb debugger. Run the program and step through it to find the error.

```
def average(numbers):
    total = 0
    for num in numbers:
        total -= num
    return total/(len(numbers) - 1)

numbers = [1, 2, 3, 4, 5]
print(average(numbers))
```

Question: What steps did you use in the debugger to confirm the program works?

18.4.0.6 Exercise 6: Logging and Error Messages

Write a program that logs the steps of a basic math operation (addition, subtraction, multiplication, and division). Use the `logging` module to track each operation, and add appropriate log levels for normal execution (`INFO`) and errors (`ERROR`).

18.4.0.7 Exercise 7: Nesting `try` and `except` Blocks

Create a program that reads a number from the user, writes it to a file, and then reads the file back. Use nested `try` and `except` blocks to handle potential errors, such as invalid input, file writing errors, or file reading errors.

Example Output:

```
Enter a number: 42
Successfully wrote the number to file.
Error: Unable to read the file.
```


19 Regular Expressions

In programming, we often encounter tasks that require searching, extracting, or validating specific patterns within strings. *Regular Expressions* (regex or regexp) provide a powerful tool for these tasks by allowing us to define search patterns with precision. Regular expressions are particularly useful in text processing, making it easier to handle tasks such as finding email addresses, extracting numerical values, or validating formatted inputs.

19.1 The Concept of Pattern Matching

At its core, a regular expression is a sequence of characters that form a **pattern**. This pattern describes the structure or sequence you are interested in finding within a larger text. Unlike simple string matching, which relies on an exact match (e.g., searching for the word “data” would only find “data” and not “dataset” or “database”), regular expressions allow for more **flexible** and **dynamic** matching, capturing partial matches or patterns that meet certain criteria.

For instance, the pattern `data\w*` will match:

- “data” (the word itself)
- “dataset”
- “database”

This flexibility is what makes regular expressions so powerful for text processing.

19.1.1 Anatomy of a Regular Expression

To build effective regular expressions, it’s helpful to understand the components that make up a pattern. Each character in a regular expression can have a special meaning, enabling more complex pattern definitions.

1. **Literal Characters:** Most characters (like letters or numbers) match themselves. For instance, the pattern `data` will only match occurrences of the word “data”.
2. **Metacharacters:** These are special characters that have specific functions in regex syntax. Some commonly used metacharacters are:

- **.** (**dot**): Matches any single character except a newline.
 - **^** (**caret**): Indicates the start of a string.
 - **\$** (**dollar**): Indicates the end of a string.
 - *****, **+**, **?**: Define the frequency or repetition of patterns.
3. **Character Classes:** Defined by brackets `[]`, character classes allow you to match any one of a set of characters.
 - For example, `[abc]` will match either “a”, “b”, or “c”.
 - Ranges can also be used within brackets, like `[0-9]` to match any digit.
 4. **Escape Sequences:** Some characters (such as `.` or `*`) have special meanings, so you may need to **escape** them with a backslash (`\`) if you want to match them literally. For instance, `\.` will match a period rather than any character.
 5. **Quantifiers:** These symbols define how many times a character or pattern must appear.
 - `*` matches zero or more occurrences.
 - `+` matches one or more occurrences.
 - `?` matches zero or one occurrence.

19.1.2 Why Use Regular Expressions?

Regular expressions are particularly useful when handling unstructured data, such as:

- **Data Extraction:** Pulling specific data, like email addresses, phone numbers, or dates, from large text files.
- **Data Validation:** Ensuring that an input string follows a required format, such as validating email addresses or passwords.
- **Text Replacement:** Substituting parts of a string with new values, which can help anonymize data or clean up formatting.

Regular expressions are versatile enough to support a wide variety of tasks in text processing, making them invaluable in fields like data science, web development, and natural language processing.

19.2 Using Regular Expressions in Python

Python’s `re` library is designed to work with regular expressions, providing several powerful functions for pattern matching and text manipulation. Below, we explore some of the key functions in the `re` library, each accompanied by examples to demonstrate their practical applications.

19.2.1 Key Functions in Python's re Library

1. **re.search()**: This function searches for the first occurrence of a pattern within a string. If it finds a match, it returns a match object; otherwise, it returns **None**.
2. **re.findall()**: This function finds all non-overlapping matches of a pattern in a string and returns them as a list.
3. **re.match()**: This function checks for a match only at the beginning of a string.
4. **re.sub()**: This function replaces occurrences of a pattern with a specified replacement string.
5. **re.split()**: This function splits a string by occurrences of a pattern, returning a list of substrings.

Each function serves a unique purpose and can be applied to various text-processing tasks.

19.2.2 Using re.search()

The **re.search()** function finds the first match of a pattern within a string. This is useful when you only need to confirm the existence of a pattern or retrieve the first match.

Example 1: Searching for a word that starts with “data”

```
import re

text = "The dataset is being updated in the database."
pattern = r"data\w*"

# Search for the first occurrence of the pattern
result = re.search(pattern, text)
if result:
    print("Found:", result.group())
```

Found: dataset

In this example: - The pattern **data\w*** matches “data” followed by zero or more word characters (**\w***). - The **result.group()** method returns the matched string.

Example 2: Checking for the presence of a phone number pattern

```
text = "Contact us at 123-867-5309."
pattern = r"\d{3}-\d{3}-\d{4}"
```

```
result = re.search(pattern, text)
if result:
    print("Phone number found:", result.group())
```

Phone number found: 123-867-5309

The `\d` metacharacter matches any digit and is equivalent to `[0-9]`. When followed by `{3}`, then we are looking for three consecutive digits.

19.2.3 Using `re.findall()`

The `re.findall()` function returns all matches of a pattern as a list, making it ideal for finding multiple occurrences of a pattern within a string.

Example 1: Extracting all email addresses

```
text = "Emails: alice@example.com, bob@example.org, charlie@test.com"
pattern = r"\b\w+@\w+\.\w+\b"

emails = re.findall(pattern, text)
print("Emails found:", emails)
```

Emails found: ['alice@example.com', 'bob@example.org', 'charlie@test.com']

In this example:

- The pattern `\b\w+@\w+\.\w+\b` matches email addresses by looking for word characters around the “@” and “.” symbols.
- `re.findall()` captures all email addresses in the text.

Example 2: Finding all words starting with “stat”

```
text = "Statistics is a fascinating field, with stats and statistical methods widely appli"
pattern = r"\bstat\w*"

matches = re.findall(pattern, text)
print("Words found:", matches)
```

Words found: ['stats', 'statistical']

19.2.4 re.match()

The `re.match()` function checks if a pattern matches only at the beginning of a string, returning `None` if the pattern appears elsewhere.

Example 1: Verifying the format of a postal code at the beginning of a string

```
text = "12345-6789 is the postal code."
pattern = r"^\d{5}-\d{4}"

if re.match(pattern, text):
    print("Valid postal code format.")
else:
    print("Invalid format.")
```

Valid postal code format.

Here:

- The pattern `^\d{5}-\d{4}` matches a 5-digit number, a hyphen, and a 4-digit number, specifically at the start of the string.

19.2.5 Using re.sub()

The `re.sub()` function replaces all occurrences of a pattern with a specified replacement string, making it useful for sanitizing or formatting data.

Example 1: Replacing phone numbers with “[PHONE]”

```
text = "Reach us at 123-456-7890 or 987-654-3210."
pattern = r"\d{3}-\d{3}-\d{4}"

new_text = re.sub(pattern, "[PHONE]", text)
print(new_text)
```

Reach us at [PHONE] or [PHONE].

Example 2: Standardizing date formats

```
text = "Event on 2023-01-25, follow-up on 01/26/2023."
pattern = r"(\d{4})-(\d{2})-(\d{2})"
```

```
# Reformat dates to MM/DD/YYYY
new_text = re.sub(pattern, r"\2/\3/\1", text)
print(new_text)
```

Event on 01/25/2023, follow-up on 01/26/2023.

In this example:

- The pattern `(\d{4})-(\d{2})-(\d{2})` captures the year, month, and day.
- The replacement pattern `\2/\3/\1` reorders the components to MM/DD/YYYY.

19.2.6 Using `re.split()`

The `re.split()` function splits a string based on a regular expression, which is useful when splitting by complex delimiters.

Example 1: Splitting text by multiple delimiters

```
text = "apple; orange, banana|grape"
pattern = r"[;,|]"

fruits = re.split(pattern, text)
print(fruits)
```

```
['apple', ' orange', ' banana', 'grape']
```

Here, the pattern `[;,|]` matches any one of `;`, `,`, or `|` as delimiters.

19.3 Common Metacharacters

Here's a list of commonly used special characters (also called *escape sequences* or *metacharacters*) in regular expressions, along with their functions:

19.3.0.1 Anchors

- `^`: Matches the start of a string.
- `$`: Matches the end of a string.
- `\b`: Matches a word boundary (position between a word and a non-word character).
- `\B`: Matches a position that is not a word boundary.

19.3.0.2 Character Classes

- `.`: Matches any character except a newline.
- `\d`: Matches any digit, equivalent to `[0-9]`.
- `\D`: Matches any non-digit character, equivalent to `[^0-9]`.
- `\w`: Matches any word character (alphanumeric or underscore), equivalent to `[a-zA-Z0-9_]`.
- `\W`: Matches any non-word character, equivalent to `[^a-zA-Z0-9_]`.
- `\s`: Matches any whitespace character (spaces, tabs, newlines).
- `\S`: Matches any non-whitespace character.

19.3.0.3 Quantifiers

- `*`: Matches 0 or more occurrences of the preceding element.
- `+`: Matches 1 or more occurrences of the preceding element.
- `?`: Matches 0 or 1 occurrence of the preceding element.
- `{n}`: Matches exactly `n` occurrences of the preceding element.
- `{n,}`: Matches `n` or more occurrences.
- `{n,m}`: Matches between `n` and `m` occurrences.

19.3.0.4 Groups and References

- `()`: Groups a pattern together, allowing you to apply quantifiers to the entire group or capture matches.
- `|`: Alternation operator, meaning “or” (e.g., `data|info` matches “data” or “info”).
- `\`: Escapes special characters, allowing them to be used as literal characters (e.g., `\.` matches a period).
- `\1`, `\2`, etc.: Refers to matched groups in the pattern, useful for back-references.

19.3.0.5 Lookaheads and Lookbehinds

- `(?=...)`: Positive lookahead, ensures that what follows matches
- `(?!...)`: Negative lookahead, ensures that what follows does not match
- `(?<=...)`: Positive lookbehind, ensures that what precedes matches
- `(?<!=...)`: Negative lookbehind, ensures that what precedes does not match

19.3.0.6 Escaped Characters for Specific Needs

- `\t`: Matches a tab character.
- `\n`: Matches a newline character.
- `\r`: Matches a carriage return character.
- `\f`: Matches a form feed character.
- `\v`: Matches a vertical tab character.
- `\0`: Matches the null character.

19.4 Building Complex Patterns

Regular expressions become particularly powerful when we combine multiple elements to create complex patterns. This section explores some advanced techniques to build sophisticated expressions, allowing for precise control over pattern matching.

19.4.1 Using Grouping and Capturing

Grouping is achieved by enclosing parts of a pattern within parentheses `()`. Groups allow you to:

1. Apply quantifiers to an entire section of a pattern.
2. Capture parts of a match, enabling you to reference them later (known as **capturing groups**).

Example: Capturing parts of a date

```
import re

text = "Today's date is 2024-11-03."
pattern = r"(\d{4})-(\d{2})-(\d{2})"

match = re.search(pattern, text)
if match:
    year, month, day = match.groups()
    print(f"Year: {year}, Month: {month}, Day: {day}")
```

Year: 2024, Month: 11, Day: 03

In this example:

- The pattern `(\d{4})-(\d{2})-(\d{2})` captures the year, month, and day as separate groups.
- The `match.groups()` method returns a tuple with the captured parts.

19.4.2 Lookahead and Lookbehind Assertions

Lookahead and **lookbehind** assertions allow you to match a pattern based on what follows or precedes it, without including that part in the match.

- **Positive Lookahead** `(?=...)`: Asserts that a match is followed by a specific pattern.
- **Negative Lookahead** `(?!...)`: Asserts that a match is **not** followed by a specific pattern.
- **Positive Lookbehind** `(?<=...)`: Asserts that a match is preceded by a specific pattern.
- **Negative Lookbehind** `(?<!=...)`: Asserts that a match is **not** preceded by a specific pattern.

Example 1: Finding words followed by “ing” (positive lookahead)

```
text = "The following items are walking, talking, and reading."
pattern = r"\b\w+(?=ing\b)"

matches = re.findall(pattern, text)
print(matches)
```

```
['follow', 'walk', 'talk', 'read']
```

Example 2: Finding numbers not preceded by a dollar sign (negative lookbehind)

```
text = "Price is $100 but I paid 150."
pattern = r"(?

```

```
['150']
```

19.4.3 Alternation with the | Operator

The | operator allows you to specify multiple patterns, matching if any of the alternatives are found.

Example: Matching multiple file extensions

```
text = "Files: report.pdf, image.jpg, document.docx"
pattern = r"\b\w+\.(pdf|jpg|docx)\b"

matches = re.findall(pattern, text)
print(matches)
```

```
['pdf', 'jpg', 'docx']
```

Here, the pattern matches any word followed by a file extension (either “pdf”, “jpg”, or “docx”).

19.4.4 Working with Character Classes and Ranges

Character classes [] allow for more flexibility by matching any one character within the brackets. You can also define ranges within classes to match multiple characters more succinctly.

- **Ranges:** For example, [a-z] matches any lowercase letter, while [0-9] matches any digit.
- **Negated Classes:** Using [^...] matches any character **not** in the brackets.

Example 1: Matching only vowels

```
text = "Regular expressions are useful."
pattern = r"[aeiou]"

vowels = re.findall(pattern, text)
print(vowels)
```

```
['e', 'u', 'a', 'e', 'e', 'i', 'o', 'a', 'e', 'u', 'e', 'u']
```

Example 2: Matching any character except vowels

```

pattern = r"[^aeiou]"

non_vowels = re.findall(pattern, text)
print(non_vowels)

```

```

['R', 'g', 'l', 'r', ' ', 'x', 'p', 'r', 's', 's', 'n', 's', ' ', 'r', ' ', 's', 'f', 'l', ' '

```

19.4.5 Using Backreferences

Backreferences allow you to reuse a captured group within the same pattern. This is especially useful for finding repeated patterns.

Example: Matching repeated words

```

text = "The the test was a success."
pattern = r"\b(\w+)\s+\1\b"

matches = re.findall(pattern, text, re.IGNORECASE)
print(matches)

```

```

['The']

```

Here: - `(\w+)` captures a word, and `\1` references this captured group, matching any repeated word.

19.4.6 Conditional Matching with Quantifiers

Quantifiers, such as `*`, `+`, and `{n,m}`, allow you to control the frequency of matched elements, providing even more flexibility for complex patterns.

- `*`: Matches zero or more occurrences.
- `+`: Matches one or more occurrences.
- `?`: Matches zero or one occurrence.
- `{n}`: Matches exactly `n` occurrences.
- `{n,}`: Matches `n` or more occurrences.
- `{n,m}`: Matches between `n` and `m` occurrences.

Example 1: Matching sequences of digits with varying lengths

```

text = "Order numbers: 123, 4567, 89, 23456"
pattern = r"\b\d{2,4}\b"

matches = re.findall(pattern, text)
print(matches)

```

```
['123', '4567', '89']
```

This pattern matches numbers that are 2 to 4 digits long.

19.4.7 Combining Techniques for Complex Patterns

By combining groups, lookaheads, alternation, and quantifiers, you can construct highly specific patterns.

Example: Extracting full names in the format “Last, First M.”

```

text = "Attendees: Smith, John A.; Doe, Jane B."
pattern = r"\b([A-Z] [a-z]+),\s([A-Z] [a-z]+)\b"

names = re.findall(pattern, text)
print(names)

```

```
[('Smith', 'John'), ('Doe', 'Jane')]
```

Here:

- `[A-Z] [a-z]+` matches capitalized words.
- `\s` matches spaces.

19.5 Exercises

19.5.0.1 Exercise 1: Extracting Domain Names

- Write a regular expression to find all domain names in a text. Assume that domain names have the format `www.something.com`, `something.org`, or `something.edu`.
- **Sample Input:** "Visit us at `www.example.com` or go to `research.org` for more info."
- **Expected Output:** `['example.com', 'research.org']`

19.5.0.2 Exercise 2: Removing Punctuation

- Use `re.sub()` to remove all punctuation from a sentence, leaving only alphanumeric characters and spaces.
- **Sample Input:** "Hello, world! Let's test this: are you ready?"
- **Expected Output:** "Hello world Lets test this are you ready"

19.5.0.3 Exercise 3: Extracting Hashtags

- Write a pattern to find all hashtags in a given string.
- **Sample Input:** "Join the conversation with #Python, #DataScience, and #MachineLearning!"
- **Expected Output:** ['#Python', '#DataScience', '#MachineLearning']

19.5.0.4 Exercise 4: Finding Dates in Different Formats

- Write a regular expression to find dates in multiple formats, such as MM/DD/YYYY, YYYY-MM-DD, and DD.MM.YYYY.
- **Sample Input:** "We met on 2024-11-03, and we'll meet again on 11/03/2024."
- **Expected Output:** ['2024-11-03', '11/03/2024']

19.5.0.5 Exercise 5: Finding Words That Start with a Specific Letter

- Write a regular expression to find all words that start with the letter "s" (case-insensitive).
- **Sample Input:** "She sells seashells by the seashore."
- **Expected Output:** ['She', 'sells', 'seashells', 'seashore']

19.5.0.6 Exercise 6: Redacting Sensitive Information

- Use `re.sub()` to replace all Social Security numbers (formatted as ###-##-####) with [REDACTED].
- **Sample Input:** "Client's SSN is 123-45-6789."
- **Expected Output:** "Client's SSN is [REDACTED]."

19.5.0.7 Exercise 7: Validating Password Strength

- Write a regular expression to validate that a password is at least 8 characters long, contains at least one uppercase letter, one lowercase letter, one digit, and one special character (e.g., @, #, !, \$).
- **Sample Input:** Password123!
- **Expected Output:** Match found
- **Sample Input:** weakpass
- **Expected Output:** No match

19.5.0.8 Exercise 8: Extracting Full Names

- Write a regular expression to capture full names with the format "Last, First M.", where "M" is a middle initial.
- **Sample Input:** `"Dr. Smith, John A. and Dr. Brown, Lisa attended the meeting."`
- **Expected Output:** `[('Smith', 'John', 'A.'), ('Brown', 'Lisa', '')]`

19.5.0.9 Exercise 9: Finding Repeated Words

- Write a regular expression to find and return any repeated words in a sentence.
- **Sample Input:** `"This is a test test sentence."`
- **Expected Output:** `['test']`

19.5.0.10 Exercise 10: Splitting a Sentence by Words with Multiple Delimiters

- Write a regular expression to split a sentence by multiple delimiters, such as spaces, commas, semicolons, and dashes.
- **Sample Input:** `"Apples; oranges, and bananas are tasty."`
- **Expected Output:** `['Apples', 'oranges', 'and', 'bananas', 'are', 'tasty']`

19.5.0.11 Exercise 11: Extracting File Names with Specific Extensions

- Write a pattern to find all files with `.pdf`, `.docx`, or `.xlsx` extensions.
- **Sample Input:** `"Documents include report.pdf, summary.docx, and data.xlsx."`
- **Expected Output:** `['report.pdf', 'summary.docx', 'data.xlsx']`

19.5.0.12 Exercise 12: Extracting Sentences without Ending with a Question Mark

- Write a regular expression to capture all sentences that do not end with a question mark.
- **Sample Input:** `"What is your name? My name is Alice. Where are you from?"`
- **Expected Output:** `['My name is Alice']`